


01: Coding for Performance

Engr 315: Hardware / Software Codesign

Andrew Lukefahr

Indiana University



Some material taken from:

https://github.com/trekhleb/homemade-machine-learning/tree/master/homemade/neural_network

<http://cs231n.github.io/neural-networks-1/>

Course Website

engr315.github.io

Write that down!

Slack? - <https://engr-315.slack.com>

- *Thanks Joel*

I try to post all the code I use in class

The screenshot shows the GitHub interface for the repository 'Engr315 / lecture_slides'. The repository is public and has an 'Unwatch' button. The 'Code' tab is selected, showing the file structure and recent commits. The file structure includes a folder 'code_00' and a file '00_intro.pdf', both of which were added in a commit 35 seconds ago. The commit message is 'Adding 00_intro.pdf'.

Engr315 / lecture_slides Public Unwatch

<> Code Issues Pull requests Actions Projects Wiki Security


master 1 branch 0 tags Go to file Add file Code

Andrew Lukefahr Adding 00_intro.pdf 2d15b57 35 seconds ago 6 commits



code_00	Adding 00_intro.pdf	35 seconds ago
00_intro.pdf	Adding 00_intro.pdf	35 seconds ago

Remind me if (when) I forget.

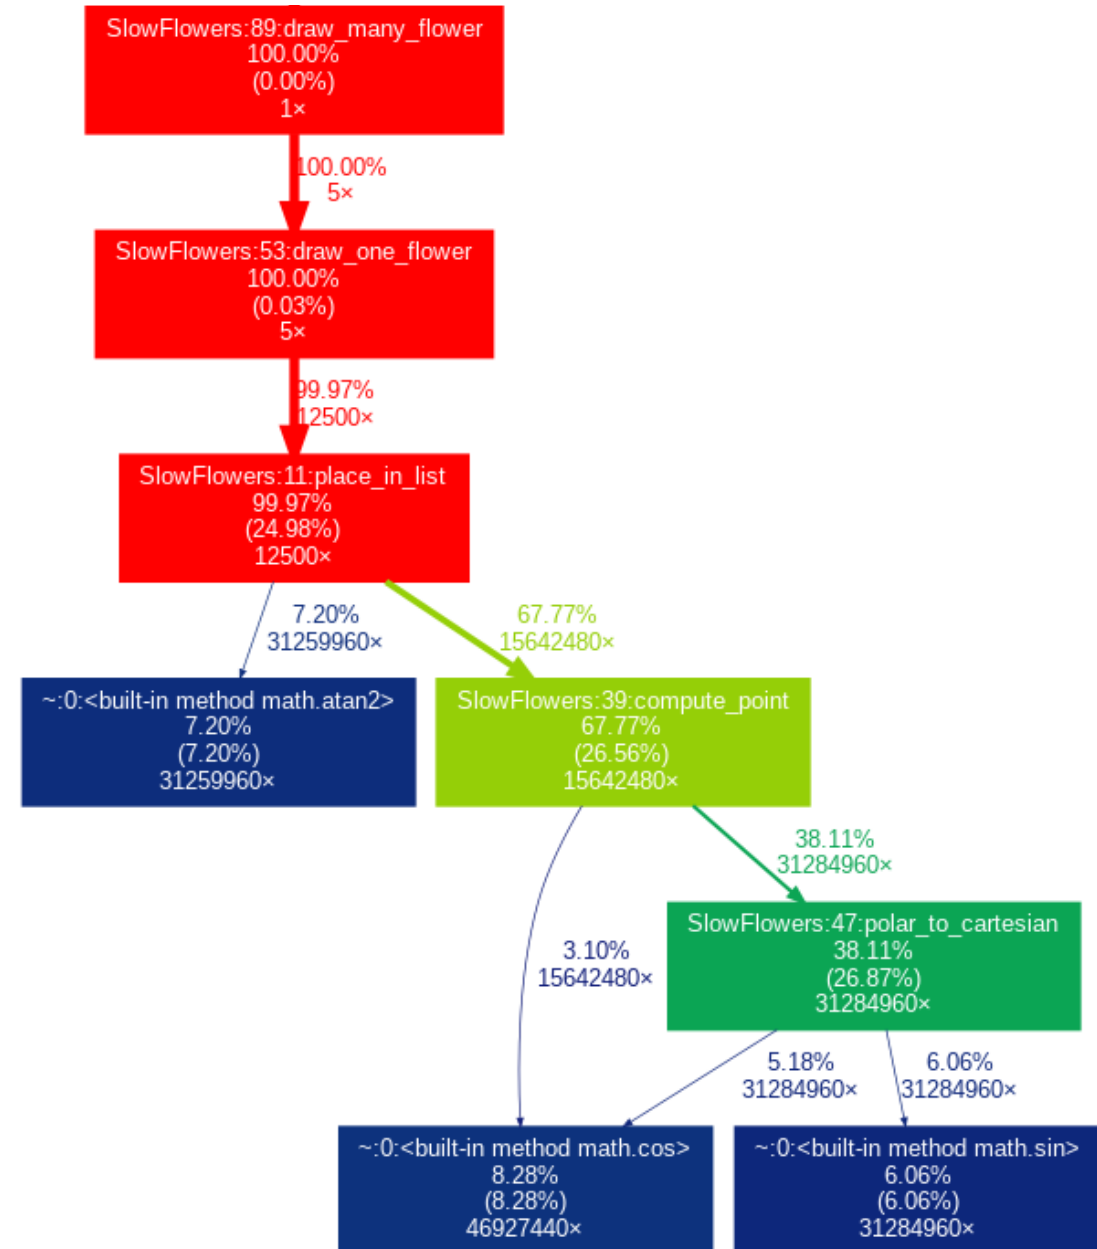
Project 1: Optimization



Project 1: Flowers



Project 1: Flowers



Project 1: Flowers

- This is optimized already:

```
oldTheta = atan2(point[0]-x_c, point[1]-y_c)
r = (a * cos(5 * (theta - (pi/2)))) + (a * 1.3)
```

- You aren't going to accelerate math functions. Don't try.
- Figure out how to call it less.

Project 1: “Bonus”

- Bonus: Gets you bonus points!
 - My old best time
- Better Bonus: Gets you even more bonus points!
 - (best time from student)

Good luck!

Code Profiling

- In software engineering, **profiling** ("program profiling", "software profiling") is a form of dynamic program analysis that **measures**, for example, the space (memory) or time complexity of a program, the usage of particular instructions, or **the frequency and duration of function calls**. Most commonly, profiling information serves to aid program optimization. [Wiki]

Code Profiling can measure

- Program Runtimes
- **Function Call Numbers/Runtimes**
- Memory Usage
- Instruction Usage
- Others

Profiling guide us on where to look to reduce runtime

```
1 def squares(n):
2     if n <= 1:
3         return [1]
4     else:
5         seq = squares(n-1)
6         seq.append(n*n)
7         return seq
```

40002 function calls (20003 primitive calls) in 0.021 seconds

Ordered by: standard name

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
20000/1	0.019	0.000	0.021	0.021	<ipython-input-8-50d13c5dd8df>:1(squares)
1	0.000	0.000	0.021	0.021	<string>:1(<module>)
1	0.000	0.000	0.021	0.021	{built-in method builtins.exec}
19999	0.002	0.000	0.002	0.000	{method 'append' of 'list' objects}
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}

Conclusion #1: Function calls are not free!

- Setup/Return overheads with function calls
 - Small
- Recursion: small overheads x many calls
 - Can add notable overheads
- Only use recursion if don't care about performance

Cutting recursion buys us $\sim 2x$

```
1 def squares(n):
2     if n <= 1:
3         return [1]
4     else:
5         seq = squares(n-1)
6         seq.append(n*n)
7         return seq
```

```
1 def squares2(n):
2     if n <= 1:
3         return [1]
4     else:
5         seq = []
6         for i in range(1,n):
7             seq.append(i*i)
8         return seq
```

```
1 import time
2 import sys
3 sys.setrecursionlimit(21000)
4
5 start_time = time.time()
6 squares(20000)
7 end_time = time.time()
8
9 # at the end of the program:
10 print("%f seconds" % (end_time - start_time))
```

0.009825 seconds

```
1 import time
2
3 start_time = time.time()
4 squares2(20000)
5 end_time = time.time()
6
7 # at the end of the program:
8 print("%f seconds" % (end_time - start_time))
```

0.004209 seconds

Can we make it go even faster?

```
1 def squares2(n):
2     if n <= 1:
3         return [1]
4     else:
5         seq = []
6         for i in range(1,n):
7             seq.append(i*i)
8         return seq
```

```
1 import time
2
3 start_time = time.time()
4 squares2(20000)
5 end_time = time.time()
6
7 # at the end of the program:
8 print("%f seconds" % (end_time - start_time))
```

0.004209 seconds

```
1 import numpy as np
2 def squares3(n):
3
4     seq = np.zeros(n, dtype=np.int)
5     for i in range(1, n+1):
6         seq[i-1] = i * i
7     return seq
```

```
1 import time
2
3 start_time = time.time()
4 squares3(20000)
5 end_time = time.time()
6
7 # at the end of the program:
8 print("%f seconds" % (end_time - start_time))
```

0.003960 seconds

... And I'm bested!


```
1  #Thanks Drason!
2  def squares4(n):
3      return [i * i for i in range(1, n+1)]
4
5  start_time = time.time()
6  squares4(20000)
7  end_time = time.time()
8
9  # at the end of the program:
10 print("%.f seconds" % (end_time - start_time))
```

0.003010 seconds


Conclusion #2: memory preallocation *should be faster*

- Numpy reallocates a large contiguous block 😊
 - But also zeros it out. 😞
- List.append() allocates new memory as needed
- Python's "List Comprehension" is weird.

np.empty



Array vs. Linked List: Which is faster?



- Randomly accessing a specific element? *array*
- Appending new values? \rightarrow *list*

Array vs. Linked List: Random Access

```
1 lst = collections.deque(nums)
2 arr = np.array(nums)
3 print (lst)
4 print (arr)
```

```
deque([5, 1, 9, 0, 3, 2, 6, 4, 8, 7])
[5 1 9 0 3 2 6 4 8 7]
```

```
1 def traverse( thing, times):
2     idx = 0
3     for i in range(times):
4         nidx = thing[idx]
5         print (i, ': ', idx, '->', nidx)
6         idx = nidx
```

```
1 trips = 10
2 traverse(lst, trips)
```

```
0 : 0 -> 5
1 : 5 -> 2
2 : 2 -> 9
3 : 9 -> 7
4 : 7 -> 4
5 : 4 -> 3
6 : 3 -> 0
7 : 0 -> 5
8 : 5 -> 2
9 : 2 -> 9
```

Array vs. Linked List: Random Access

```
1 start_time = time.time()
2 traverse(lst, trips)
3 end_time = time.time()
4
5 # at the end of the program:
6 print("True List: %f seconds" % (end_time - start_time))
7
8 start_time = time.time()
9 traverse(arr, trips)
10 end_time = time.time()
11
12 # at the end of the program:
13 print("Array: %f seconds" % (end_time - start_time))
```

```
0 : 0 -> 5
1 : 5 -> 2
2 : 2 -> 9
3 : 9 -> 7
4 : 7 -> 4
5 : 4 -> 3
6 : 3 -> 0
7 : 0 -> 5
8 : 5 -> 2
9 : 2 -> 9
True List: 0.001251 seconds
0 : 0 -> 5
1 : 5 -> 2
2 : 2 -> 9
3 : 9 -> 7
4 : 7 -> 4
5 : 4 -> 3
6 : 3 -> 0
7 : 0 -> 5
8 : 5 -> 2
9 : 2 -> 9
Array: 0.006385 seconds
```

Array vs. Linked List: Random Access

```
1 def traverse( thing, times):
2     idx = 0
3     for i in range(times):
4         idx = thing[idx]
5
6 random.seed(1)
7 sz = 1000000
8 nums = [x for x in range(sz)]
9 random.shuffle(nums)
10 random.shuffle(nums)
11 lst = collections.deque(nums)
12 arr = np.array(nums)
13 trips = 1000
14
15 start_time = time.time()
16 traverse(lst, trips)
17 end_time = time.time()
18 print("True List: %f seconds" % (end_time - start_time))
19
20 start_time = time.time()
21 traverse(arr, trips)
22 end_time = time.time()
23 print("Array: %f seconds" % (end_time - start_time))
24
25 start_time = time.time()
26 traverse(nums, trips)
27 end_time = time.time()
28 print("Python List: %f seconds" % (end_time - start_time))
```

True List: 0.037878 seconds

Array: 0.000312 seconds

Python List: 0.000410 seconds

Python's “List” isn't actually a “List”

- It's a list of arrays!

Array vs. Linked List: Sequential Insert

```
1 def insert(thing, idx, values):
2     print (thing)
3     for value in values:
4         thing.insert(idx, value)
5     print (thing)
6
7 random.seed(1)
8 sz = 10
9 nums = [x for x in range(sz)]
10 random.shuffle(nums)
11 random.shuffle(nums)
12 lst = collections.deque(nums)
13 arr = np.array(nums)
14
15 idxs = int(sz/2)
16 insert(nums, idxs, [-1,-2,-3,-4])
```

[5, 1, 9, 0, 3, 2, 6, 4, 8, 7]

[5, 1, 9, 0, 3, -4, -3, -2, -1, 2, 6, 4, 8, 7]

Array vs. Linked List: Sequential Insert of 1M elements

Insert at: 0

True List: 0.000085 seconds

Array: 0.335853 seconds

Python List: 0.115629 seconds

Insert at: 750000

True List: 0.054327 seconds

Array: 0.336377 seconds

Python List: 0.022257 seconds

Big O Complexity

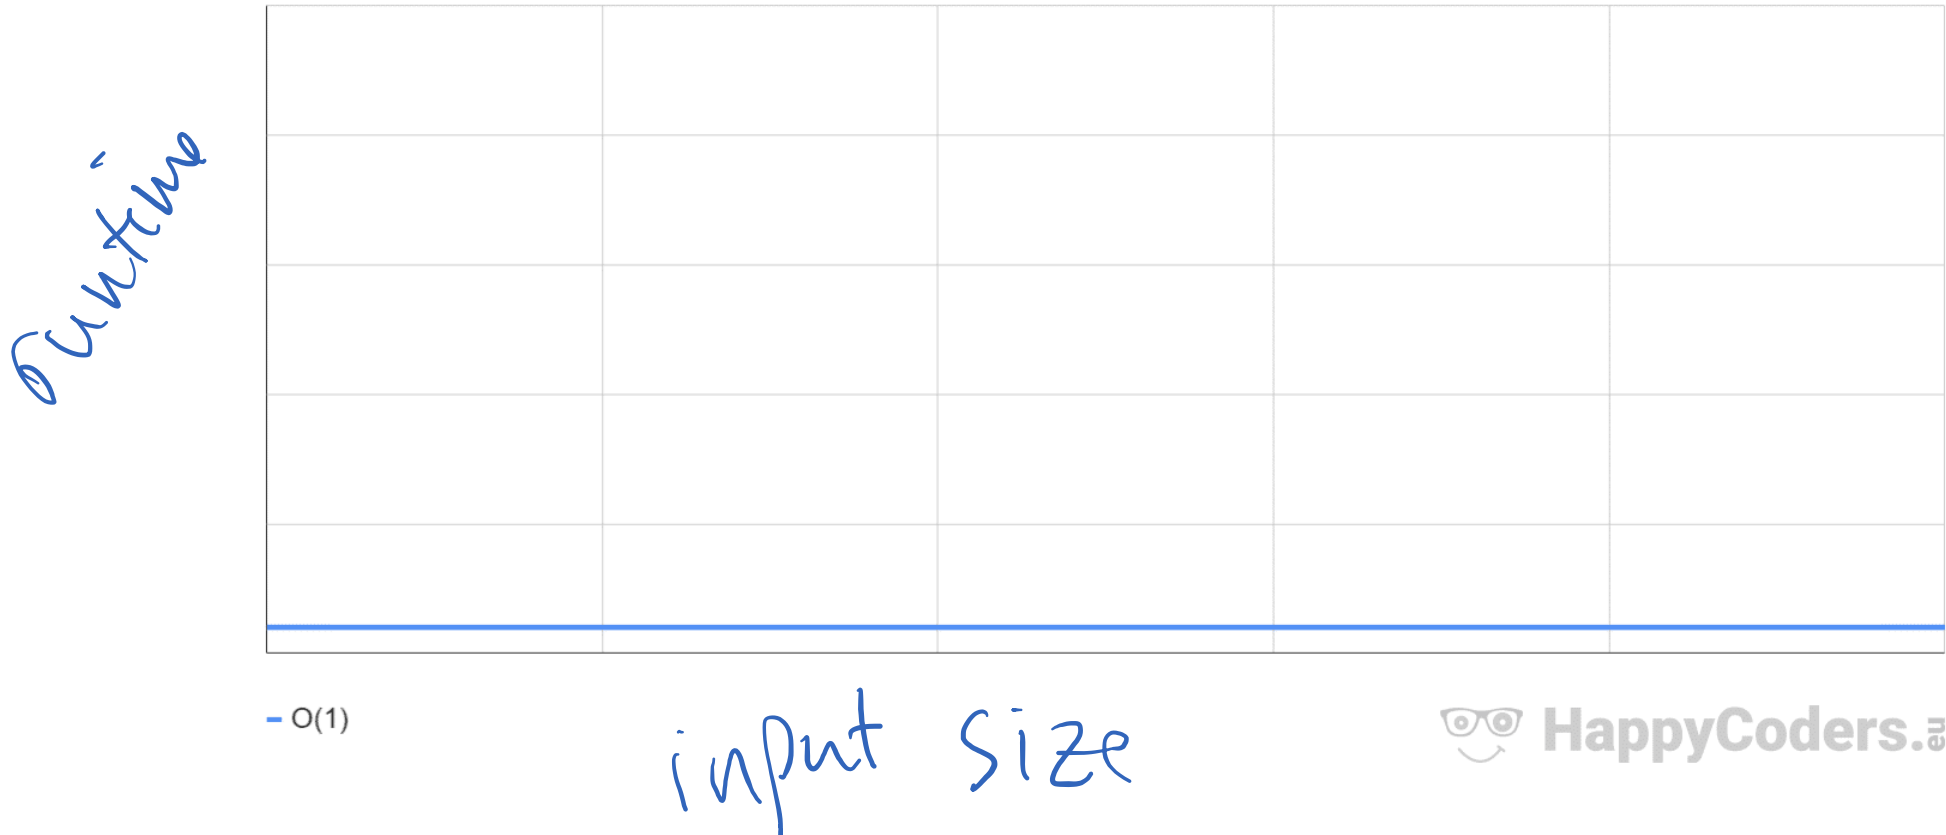
- Computational time complexity describes the change in the runtime of an algorithm, depending on the change in the input data's size.
- "How much does an algorithm's performance change when the amount of input data changes?"

$O(1)$ – Constant Time

- “big O of 1”
- Runtime is constant, regardless of input size
- Example: `x = array[n]`

$O(1)$ – Constant Time

Complexity class $O(1)$ – constant time




$O(n)$ – Linear Time

- “big O of n”
- Runtime grows linearly with input size
- Example: Linked Lists!

$O(n)$ – Linear Time

Complexity class $O(n)$ – linear time

more time



— $O(n)$

bigger

size




HappyCoders.eu

$O(n^2)$ – Quadratic Time

- “big O of n squared”
- Runtime grows linearly with square of the input size
- Example: Bubble Sort
 - `haystack.sort(low->high)`

$O(n^2)$ – Quadratic Time

Complexity class $O(n^2)$ – quadratic time



— $O(n^2)$




Material taken from: <https://www.happycoders.eu/algorithms/big-o-notation-time-complexity/>

$O(n \log n)$ – Quasilinear Time

- “big O of $n \log n$ ”
- Runtime grows linearly and logarithmically with the input size
- Example: Good Sort
 - `haystack.sort(low->high)`

$O(n \log n)$ – Quasilinear Time

Complexity class $O(n \log n)$ – quasilinear time




— $O(n \log n)$



Material taken from: <https://www.happycoders.eu/algorithms/big-o-notation-time-complexity/>

O() Complexities

Comparing the complexity classes $O(1)$, $O(\log n)$, $O(n)$, $O(n \log n)$, $O(n^2)$



Conclusion #3: Think about your data structure!

- How will you be accessing your data?
 - Randomly? Sequentially?
- How will you be updating your data?
- Pick a data structure to minimize overheads for your access patterns

Find: The needle in the haystack.

Find: The needle in the haystack.

```
def find_ignore_case( needle, haystack):
    results = []
    for hi in range(len(haystack)):
        match = True
        for ni in range(len(needle)-1):
            h = haystack[hi + ni].lower()
            n = needle[ni].lower()
            if h != n:
                match=False
        if match:
            results.append(hi)
    return results
```

• No libraries!

```
28 sz=20
29 haystack = random_str(sz)
30 needle = haystack[int(sz/2):int(sz/2)+2]
31 results = find_ignore_case(needle, haystack)
32
33 print (needle)
34 print (haystack)
35 print (results)
36
```

```
sk
eiPPzDAnWiskaumnqYpl
[10]
```

Find: The needle in the haystack.

```
1 def find_ignore_case( needle, haystack):
2     results = []
3     for hi in range(len(haystack)-len(needle)):
4         match = True
5         for ni in range(len(needle)):
6             h = haystack[hi + ni].lower()
7             n = needle[ni].lower()
8             if h != n:
9                 match=False
10        if match:
11            results.append(hi)
12    return results
13
14 random.seed(1)
15 sz=1000000
16 haystack = random_str(sz)
17 needle = haystack[int(sz/2):int(sz/2)+2]
18
19 start_time = time.time()
20 results = find_ignore_case(needle, haystack)
21 end_time = time.time()
22 print("True List: %f seconds" % (end_time - start_time))
```

True List: 1.109001 seconds

Find: The needle in the haystack.


```
: 1 import cProfile
   2 cProfile.run('find_ignore_case(needle, haystack)')
```

5001486 function calls in 1.715 seconds

Ordered by: standard name

ncalls	totttime	percall	cumtime	percall	filename:lineno(function)
1	1.276	1.276	1.715	1.715	<ipython-input-187-411d0d74434b>:1(find_ignore_case)
1	0.000	0.000	1.715	1.715	<string>:1(<module>)
1	0.000	0.000	1.715	1.715	{built-in method builtins.exec}
1000000	0.072	0.000	0.072	0.000	{built-in method builtins.len}
1490	0.000	0.000	0.000	0.000	{method 'append' of 'list' objects}
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}
3999992	0.367	0.000	0.367	0.000	{method 'lower' of 'str' objects}


Find: The needle in the haystack.



Find: The needle in the haystack.

```
1 def find_ignore_case( needle, haystack):
2     results = []
3     for hi in range(len(haystack)-len(needle)):
4         match = True
5         for ni in range(len(needle)):
6             h = haystack[hi + ni].lower()
7             n = needle[ni].lower()
8             if h != n:
9                 match=False
10        if match:
11            results.append(hi)
12    return results
13
14 random.seed(1)
15 sz=1000000
16 haystack = random_str(sz)
17 needle = haystack[int(sz/2):int(sz/2)+2]
18
19 start_time = time.time()
20 results = find_ignore_case(needle, haystack)
21 end_time = time.time()
22 print("True List: %f seconds" % (end_time - start_time))
```

True List: 1.109001 seconds



• No libraries!

Find: The needle in the haystack.

```
def find_ignore_case( needle, haystack):
    results = []
    for hi in range(len(haystack)-len(needle)):
        match = True
        for ni in range(len(needle)):
            h = haystack[hi + ni].lower()
            n = needle[ni].lower()
            if h != n:
                match=False
        if match:
            results.append(hi)
    return results
```


```
def find_ignore_case2( needle, haystack):
    results = []
    needle = needle.lower() # new
    haystack = haystack.lower() # new
    for hi in range(len(haystack)-len(needle)):
        match = True
        for ni in range(len(needle)):
            h = haystack[hi + ni]#.lower()
            n = needle[ni]#.lower()
            if h != n:
                match=False
                break # new
        if match:
            results.append(hi)
    return results
```

Find: 0.917540 seconds

Find2: 0.440155 seconds

Anything else?

```
def find_ignore_case2( needle, haystack):
    results = []
    needle = needle.lower() # new
    haystack = haystack.lower() # new
    for hi in range(len(haystack)-len(needle)):
        match = True
        for ni in range(len(needle)):
            h = haystack[hi + ni]#.lower()
            n = needle[ni]#.lower()
            if h != n:
                match=False
                break # new
        if match:
            results.append(hi)
    return results
```



```

def find_ignore_case3( needle, haystack):
    results = []
    needle = needle.lower() # new
    haystack = haystack.lower() # new
    ➔ r = range(len(needle)) # new

    for hi in range(len(haystack)-len(needle)):
        match = True

        if haystack[hi] == needle[0]:
            for ni in r: # update
                h = haystack[hi + ni].lower()
                n = needle[ni].lower()
                if h != n:
                    match=False
                    break # new
            if match:
                results.append(hi)
    return results

```

Find: 0.370030 seconds

Find2: 0.057817 seconds

Find3: 0.053763 seconds

[New Mac Times]

```

def find_ignore_case4( needle, haystack):
    results = []
    needle = needle.lower() # new
    haystack = haystack.lower() # new
    r = range(len(needle)-1) # new

    for hi in range(len(haystack)-len(needle)):
        #match = False

        if haystack[hi] == needle[0]:
            for ni in r: # update
                h = haystack[hi + ni].lower()
                n = needle[ni].lower()
                if h == n: # new
                    #match=False
                    results.append(hi) # new
                    break # new
            #if match:
            #results.append(hi)
    return results

```

Find: 0.259516 seconds
 Find2: 0.057128 seconds
 Find3: 0.053053 seconds
 Find4: 0.048197 seconds

Q: How is this so much faster?

```
def find_ignore_case5( needle, haystack):  
    return [haystack.find(needle)]
```

Using built-in libraries is usually the fastest...

```
def find_ignore_case5( needle, haystack):  
    return [haystack.find(needle)]
```

```
Find: 0.259516 seconds  
Find2: 0.057128 seconds  
Find3: 0.053053 seconds  
Find4: 0.048197 seconds  
Find5: 0.000172 seconds
```

Coding for Performance

Engr 315: Hardware / Software Codesign
Andrew Lukefahr
Indiana University

