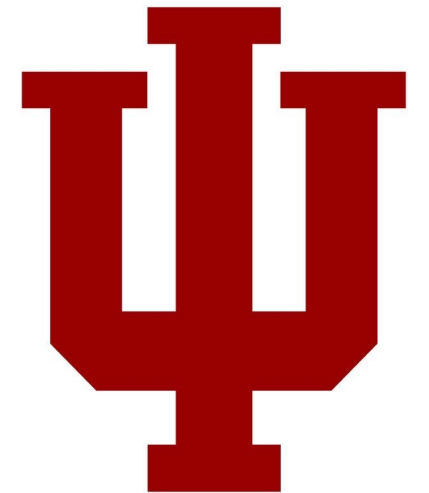


# Coding for Performance

Engr 315: Hardware / Software Codesign  
Andrew Lukefahr  
*Indiana University*



Course Website

engr315.github.io

Write that down!

# Autograder:

- Use this:
- <https://autograder.luddy.indiana.edu>
- Not this:
- <https://ag.luddy.indiana.edu>

It's a long story...

Slack - <https://e315-fall2022.slack.com>

- Thanks Nicole!
- [https://join.slack.com/t/slack-9wx7802/shared\\_invite/zt-1ep5fpz2b-bNXM9830VdBlKQGcMb3Quw](https://join.slack.com/t/slack-9wx7802/shared_invite/zt-1ep5fpz2b-bNXM9830VdBlKQGcMb3Quw)

# Office Hours

- Andrew: M/W 11-12 in 2032 Luddy or 4111 Luddy

Nicole Miller - [ncm1@iu.edu](mailto:ncm1@iu.edu)

Caleb Cook - [cookce@iu.edu](mailto:cookce@iu.edu)

## **Office Hours:**

- Tuesday: 10am-1pm
- Wednesday: 5pm-8pm
- Thursday: 5pm-8pm
- Friday: 2pm-5pm

**Location:** 3111 Luddy Hall (one floor below 4111)



# I try to post all the code I use in class

The screenshot shows the GitHub interface for the repository 'Engr315 / lecture\_slides'. The repository is public and has an 'Unwatch' button. The navigation bar includes links for Code, Issues, Pull requests, Actions, Projects, Wiki, and Security. Below the navigation bar, there are buttons for 'master' (selected), '1 branch', and '0 tags'. There are also buttons for 'Go to file', 'Add file', and a green 'Code' button. The commit history shows a commit by Andrew Lukefahr adding '00\_intro.pdf' 35 seconds ago, with 6 commits in total. Below the commit history, there is a table showing the files added in the commit.



File	Commit Message	Time
code_00	Adding 00_intro.pdf	35 seconds ago
00_intro.pdf	Adding 00_intro.pdf	35 seconds ago

Remind me if (when) I forget.

# Project 1: Circles

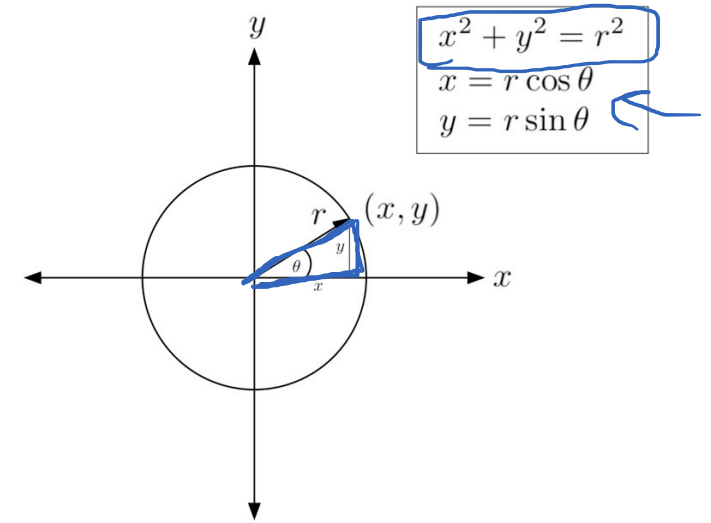
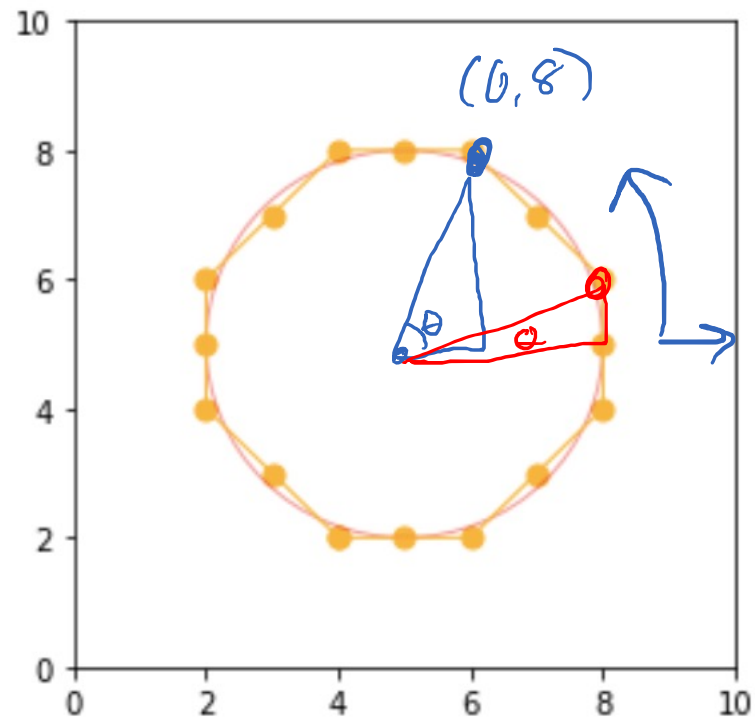


# Project 1: Circles





# Project 1: Circles



# Project 1: Circles

- This is optimized already:

```
def computeTheta(self, x,y, x_centre, y_centre):  
    return math.atan2(x-x_centre, y-y_centre)
```

- You aren't going to accelerate it. Don't try.
- Figure out how to call it less.

# Project 1: “Bonus”

- Bonus: 0.011 seconds
- Better Bonus: 0.007 seconds (best time from student)

Good luck!

# Code Profiling

- In software engineering, **profiling** ("program profiling", "software profiling") is a form of dynamic program analysis that **measures**, for example, the space (memory) or time complexity of a program, the usage of particular instructions, or **the frequency and duration of function calls**. Most commonly, profiling information serves to aid program optimization. [Wiki]

# Code Profiling can measure

- Program Runtimes
- **Function Call Numbers/Runtimes**
- Memory Usage
- Instruction Usage
- Others

# Profiling guide us on where to look to reduce runtime

```
1 def squares(n):  
2     if n <= 1:  
3         return [1]  
4     else:  
5         seq = squares(n-1)  
6         seq.append(n*n)  
7     return seq
```

40002 function calls (20003 primitive calls) in 0.021 seconds

Ordered by: standard name

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
20000/1	0.019	0.000	0.021	0.021	<ipython-input-8-50d13c5dd8df>:1(squares)
1	0.000	0.000	0.021	0.021	<string>:1(<module>)
1	0.000	0.000	0.021	0.021	{built-in method builtins.exec}
19999	0.002	0.000	0.002	0.000	{method 'append' of 'list' objects}
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}

# Conclusion #1: Function calls are not free!

- Setup/Return overheads with function calls
  - Small
- Recursion: small overheads \* many calls
  - Can add notable overheads
- Only use recursion if you must.

# Cutting recursion buys us $\sim 2x$

```
1 def squares(n):
2     if n <= 1:
3         return [1]
4     else:
5         seq = squares(n-1)
6         seq.append(n*n)
7         return seq
```

```
1 def squares2(n):
2     if n <= 1:
3         return [1]
4     else:
5         seq = []
6         for i in range(1,n):
7             seq.append(i*i)
8         return seq
```

```
1 import time
2 import sys
3 sys.setrecursionlimit(21000)
4
5 start_time = time.time()
6 squares(20000)
7 end_time = time.time()
8
9 # at the end of the program:
10 print("%f seconds" % (end_time - start_time))
```

0.009825 seconds

```
1 import time
2
3 start_time = time.time()
4 squares2(20000)
5 end_time = time.time()
6
7 # at the end of the program:
8 print("%f seconds" % (end_time - start_time))
```

0.004209 seconds



# Can we make it go even faster?

```
1 def squares2(n):
2     if n <= 1:
3         return [1]
4     else:
5         seq = []
6         for i in range(1,n):
7             seq.append(i*i)
8         return seq
```

```
1 import time
2
3 start_time = time.time()
4 squares2(20000)
5 end_time = time.time()
6
7 # at the end of the program:
8 print("%f seconds" % (end_time - start_time))
```

0.004209 seconds

```
1 import numpy as np
2 def squares3(n):
3
4     seq = np.zeros(n, dtype=np.int)
5     for i in range(1, n+1):
6         seq[i-1] = i * i
7     return seq
```

```
1 import time
2
3 start_time = time.time()
4 squares3(20000)
5 end_time = time.time()
6
7 # at the end of the program:
8 print("%f seconds" % (end_time - start_time))
```

0.003960 seconds

... And I'm bested!

```
1  #Thanks Drason!
2  def squares4(n):
3      return [i * i for i in range(1, n+1)]
4
5  start_time = time.time()
6  squares4(20000)
7  end_time = time.time()
8
9  # at the end of the program:
10 print("%.f seconds" % (end_time - start_time))
```

0.003010 seconds

## Conclusion #2: memory preallocation is *usually* faster

- Numpy reallocates a large contiguous block
- List.append() allocates new memory as needed
- Python's "List" is weird.

# Array vs. Linked List: Which is faster?

- Randomly accessing a specific element?
- Appending new values?

# Array vs. Linked List: Random Access

```
1 lst = collections.deque(nums)
2 arr = np.array(nums)
3 print (lst)
4 print (arr)
```

```
deque([5, 1, 9, 0, 3, 2, 6, 4, 8, 7])
[5 1 9 0 3 2 6 4 8 7]
```

```
1 def traverse( thing, times):
2     idx = 0
3     for i in range(times):
4         nidx = thing[idx]
5         print (i, ': ', idx, '->', nidx)
6         idx = nidx
```

```
1 trips = 10
2 traverse(lst, trips)
```

```
0 : 0 -> 5
1 : 5 -> 2
2 : 2 -> 9
3 : 9 -> 7
4 : 7 -> 4
5 : 4 -> 3
6 : 3 -> 0
7 : 0 -> 5
8 : 5 -> 2
9 : 2 -> 9
```

# Array vs. Linked List: Random Access

```
1 start_time = time.time()
2 traverse(lst, trips)
3 end_time = time.time()
4
5 # at the end of the program:
6 print("True List: %f seconds" % (end_time - start_time))
7
8 start_time = time.time()
9 traverse(arr, trips)
10 end_time = time.time()
11
12 # at the end of the program:
13 print("Array: %f seconds" % (end_time - start_time))
```

```
0 : 0 -> 5
1 : 5 -> 2
2 : 2 -> 9
3 : 9 -> 7
4 : 7 -> 4
5 : 4 -> 3
6 : 3 -> 0
7 : 0 -> 5
8 : 5 -> 2
9 : 2 -> 9
True List: 0.001251 seconds
0 : 0 -> 5
1 : 5 -> 2
2 : 2 -> 9
3 : 9 -> 7
4 : 7 -> 4
5 : 4 -> 3
6 : 3 -> 0
7 : 0 -> 5
8 : 5 -> 2
9 : 2 -> 9
Array: 0.006385 seconds
```

# Array vs. Linked List: Random Access

```
1 def traverse( thing, times):
2     idx = 0
3     for i in range(times):
4         idx = thing[idx]
5
6 random.seed(1)
7 sz = 1000000
8 nums = [x for x in range(sz)]
9 random.shuffle(nums)
10 random.shuffle(nums)
11 lst = collections.deque(nums)
12 arr = np.array(nums)
13 trips = 1000
14
15 start_time = time.time()
16 traverse(lst, trips)
17 end_time = time.time()
18 print("True List: %f seconds" % (end_time - start_time))
19
20 start_time = time.time()
21 traverse(arr, trips)
22 end_time = time.time()
23 print("Array: %f seconds" % (end_time - start_time))
24
25 start_time = time.time()
26 traverse(nums, trips)
27 end_time = time.time()
28 print("Python List: %f seconds" % (end_time - start_time))
```

True List: 0.037878 seconds

Array: 0.000312 seconds

Python List: 0.000410 seconds

# Python's “List” isn't actually a “List”

- It's a list of arrays!



# Array vs. Linked List: Sequential Insert

```
1 def insert(thing, idx, values):
2     print (thing)
3     for value in values:
4         thing.insert(idx, value)
5     print (thing)
6
7 random.seed(1)
8 sz = 10
9 nums = [x for x in range(sz)]
10 random.shuffle(nums)
11 random.shuffle(nums)
12 lst = collections.deque(nums)
13 arr = np.array(nums)
14
15 idxs = int(sz/2)
16 insert(nums, idxs, [-1,-2,-3,-4])
```

[5, 1, 9, 0, 3, 2, 6, 4, 8, 7]

[5, 1, 9, 0, 3, -4, -3, -2, -1, 2, 6, 4, 8, 7]

# Array vs. Linked List: Sequential Insert

Insert at: 0

True List: 0.000085 seconds

Array: 0.335853 seconds

Python List: 0.115629 seconds

Insert at: 750000

True List: 0.054327 seconds

Array: 0.336377 seconds

Python List: 0.022257 seconds

# Big O Complexity

- Computational time complexity describes the change in the runtime of an algorithm, depending on the change in the input data's size.
- "How much does an algorithm's performance change when the amount of input data changes?"

# $O(1)$ – Constant Time

- “big O of 1”
- Runtime is constant, regardless of input size
- Example: `x = array[n]`

# $O(1)$ – Constant Time

Complexity class  $O(1)$  – constant time



–  $O(1)$




# $O(n)$ – Linear Time

- “big O of n”
- Runtime grows linearly with input size
- Example: Linked Lists!

# $O(n)$ – Linear Time

Complexity class  $O(n)$  – linear time

more time



—  $O(n)$

bigger

size



HappyCoders.eu


# $O(n^2)$ – Quadratic Time

- “big O of n squared”
- Runtime grows linearly with square of the input size
- Example: Bubble Sort
  - `haystack.sort(low->high)`



# $O(n^2)$ – Quadratic Time

Complexity class  $O(n^2)$  – quadratic time



—  $O(n^2)$




Material taken from: <https://www.happycoders.eu/algorithms/big-o-notation-time-complexity/>

# $O(n \log n)$ – Quasilinear Time

- “big O of  $n \log n$ ”
- Runtime grows linearly and logarithmically with the input size
- Example: Good Sort
  - `haystack.sort(low->high)`

# $O(n \log n)$ – Quasilinear Time

Complexity class  $O(n \log n)$  – quasilinear time




—  $O(n \log n)$



Material taken from: <https://www.happycoders.eu/algorithms/big-o-notation-time-complexity/>

# O() Complexities

Comparing the complexity classes  $O(1)$ ,  $O(\log n)$ ,  $O(n)$ ,  $O(n \log n)$ ,  $O(n^2)$



# Conclusion #3: Think about your data structure!

- How will you be accessing your data?
  - Randomly? Sequentially?
- How will you be updating your data?
- Pick a data structure to minimize overheads for your access patterns

Find: The needle in the haystack.

# Find: The needle in the haystack.

```
1 def find_ignore_case( needle, haystack):
2     results = []
3     for hi in range(len(haystack)):
4         match = True
5         for ni in range(len(needle)):
6             h = haystack[hi + ni].lower()
7             n = needle[ni].lower()
8             if h != n:
9                 match=False
10                break
11        if match:
12            results.append(hi)
13    return results
14
```

```
28 sz=20
29 haystack = random_str(sz)
30 needle = haystack[int(sz/2):int(sz/2)+2]
31 results = find_ignore_case(needle, haystack)
32
33 print (needle)
34 print (haystack)
35 print (results)
36
```

```
sk
eiPPzDAnWiskaumnqYpl
[10]
```

# Find: The needle in the haystack.

```
1 def find_ignore_case( needle, haystack):
2     results = []
3     for hi in range(len(haystack)-len(needle)):
4         match = True
5         for ni in range(len(needle)):
6             h = haystack[hi + ni].lower()
7             n = needle[ni].lower()
8             if h != n:
9                 match=False
10        if match:
11            results.append(hi)
12    return results
13
14 random.seed(1)
15 sz=1000000
16 haystack = random_str(sz)
17 needle = haystack[int(sz/2):int(sz/2)+2]
18
19 start_time = time.time()
20 results = find_ignore_case(needle, haystack)
21 end_time = time.time()
22 print("True List: %f seconds" % (end_time - start_time))
```

True List: 1.109001 seconds



# Find: The needle in the haystack.

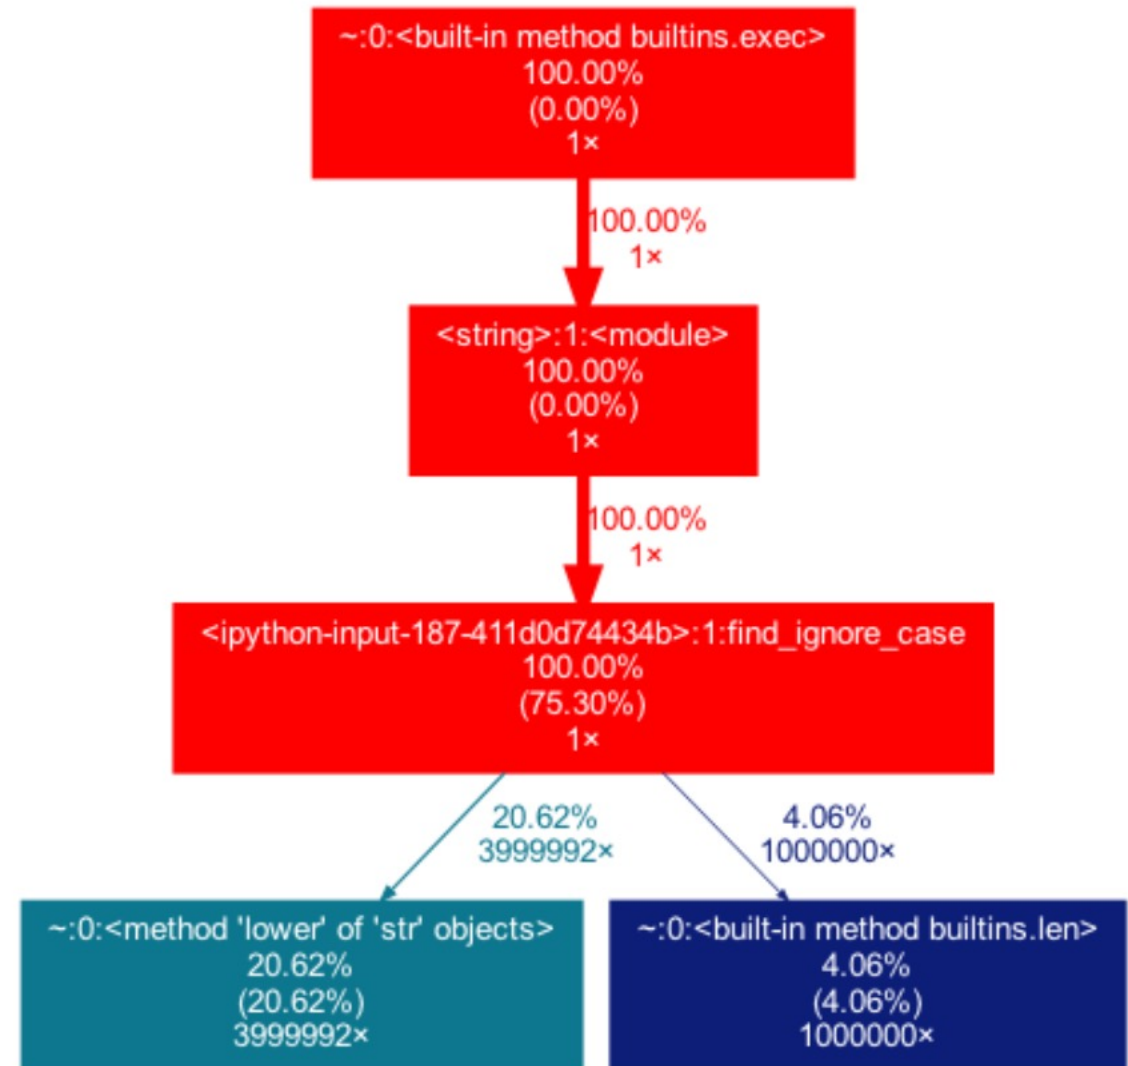
```
: 1 import cProfile
   2 cProfile.run('find_ignore_case(needle, haystack)')
```

5001486 function calls in 1.715 seconds

Ordered by: standard name

ncalls	totttime	percall	cumtime	percall	filename:lineno(function)
1	1.276	1.276	1.715	1.715	<ipython-input-187-411d0d74434b>:1(find_ignore_case)
1	0.000	0.000	1.715	1.715	<string>:1(<module>)
1	0.000	0.000	1.715	1.715	{built-in method builtins.exec}
1000000	0.072	0.000	0.072	0.000	{built-in method builtins.len}
1490	0.000	0.000	0.000	0.000	{method 'append' of 'list' objects}
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}
3999992	0.367	0.000	0.367	0.000	{method 'lower' of 'str' objects}

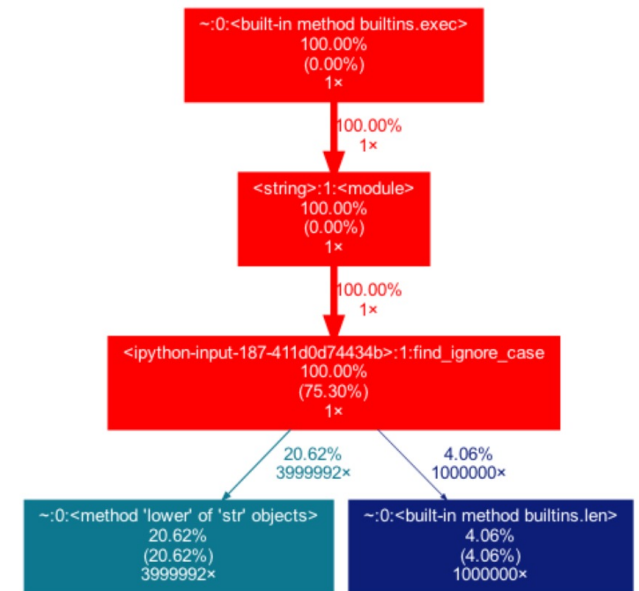
# Find: The needle in the haystack.



# Find: The needle in the haystack.

```
1 def find_ignore_case( needle, haystack):
2     results = []
3     for hi in range(len(haystack)-len(needle)):
4         match = True
5         for ni in range(len(needle)):
6             h = haystack[hi + ni].lower()
7             n = needle[ni].lower()
8             if h != n:
9                 match=False
10        if match:
11            results.append(hi)
12    return results
13
14 random.seed(1)
15 sz=1000000
16 haystack = random_str(sz)
17 needle = haystack[int(sz/2):int(sz/2)+2]
18
19 start_time = time.time()
20 results = find_ignore_case(needle, haystack)
21 end_time = time.time()
22 print("True List: %f seconds" % (end_time - start_time))
```

True List: 1.109001 seconds



• No libraries!

# Find: The needle in the haystack.

```
def find_ignore_case( needle, haystack):
    results = []
    for hi in range(len(haystack)-len(needle)):
        match = True
        for ni in range(len(needle)):
            h = haystack[hi + ni].lower()
            n = needle[ni].lower()
            if h != n:
                match=False
        if match:
            results.append(hi)
    return results
```

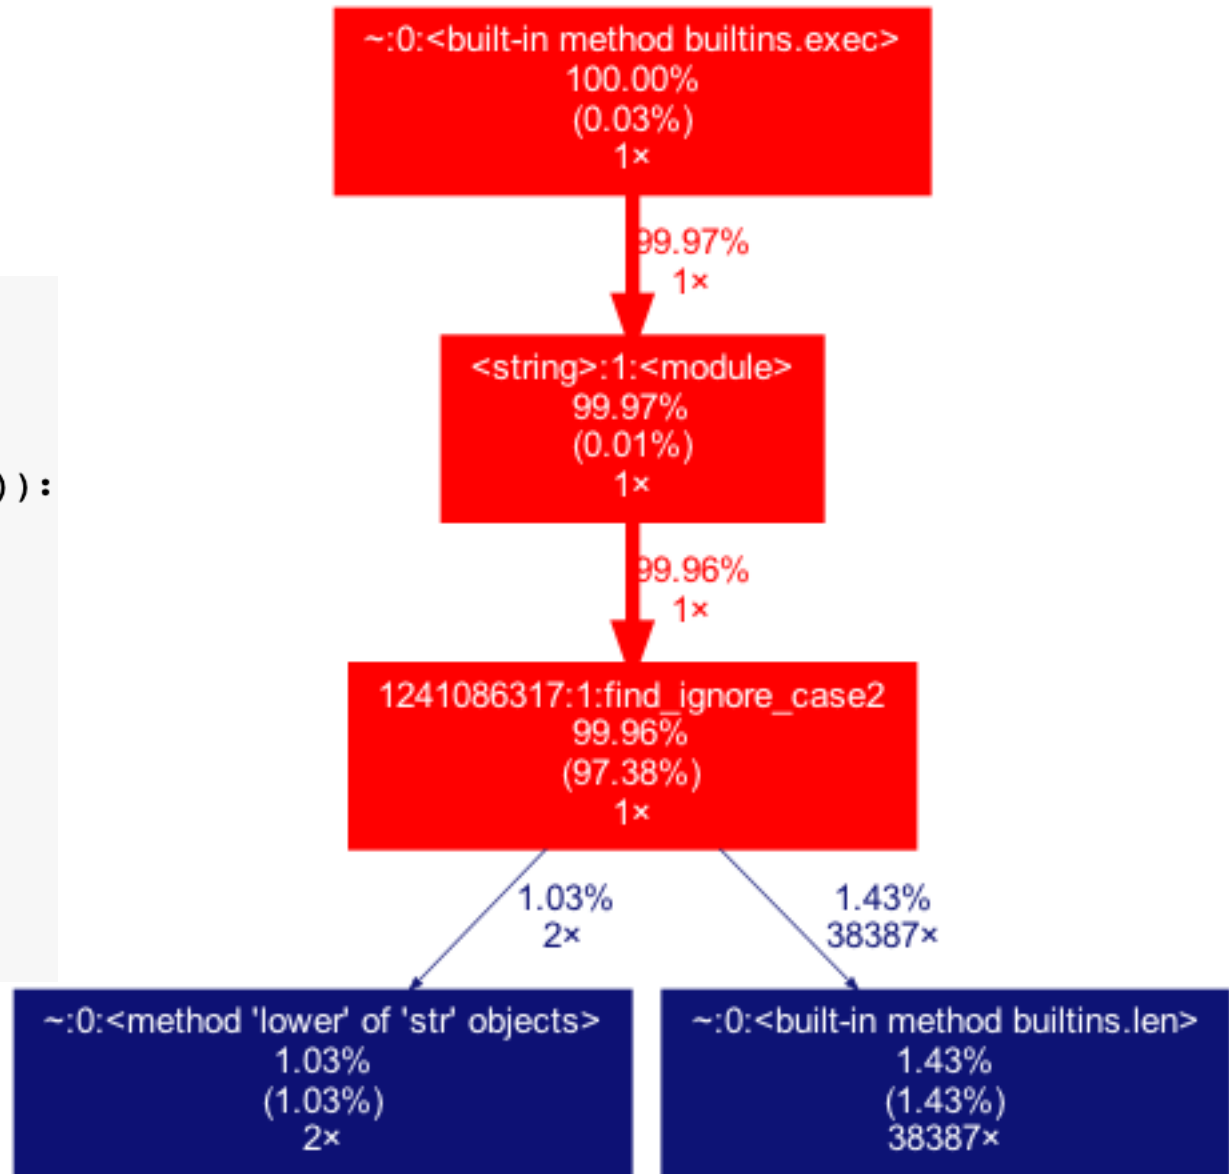
```
def find_ignore_case2( needle, haystack):
    results = []
    needle = needle.lower() # new
    haystack = haystack.lower() # new
    for hi in range(len(haystack)-len(needle)):
        match = True
        for ni in range(len(needle)):
            h = haystack[hi + ni]#.lower()
            n = needle[ni]#.lower()
            if h != n:
                match=False
                break # new
        if match:
            results.append(hi)
    return results
```

Find: 0.917540 seconds

Find2: 0.440155 seconds

# Anything else?

```
def find_ignore_case2( needle, haystack):  
    results = []  
    needle = needle.lower() # new  
    haystack = haystack.lower() # new  
    for hi in range(len(haystack)-len(needle)):  
        match = True  
        for ni in range(len(needle)):  
            h = haystack[hi + ni]#.lower()  
            n = needle[ni]#.lower()  
            if h != n:  
                match=False  
                break # new  
        if match:  
            results.append(hi)  
    return results
```



```

def find_ignore_case3( needle, haystack):
    results = []
    needle = needle.lower() # new
    haystack = haystack.lower() # new
    r = range(len(needle)) # new

    for hi in range(len(haystack)-len(needle)):
        match = True

        if haystack[hi] == needle[0]:
            for ni in r: # update
                h = haystack[hi + ni].lower()
                n = needle[ni].lower()
                if h != n:
                    match=False
                    break # new
            if match:
                results.append(hi)
    return results

```

Find: 0.370030 seconds

Find2: 0.057817 seconds

Find3: 0.053763 seconds

[New Mac Times]

```

def find_ignore_case4( needle, haystack):
    results = []
    needle = needle.lower() # new
    haystack = haystack.lower() # new
    r = range(len(needle)-1) # new

    for hi in range(len(haystack)-len(needle)):
        #match = False

        if haystack[hi] == needle[0]:
            for ni in r: # update
                h = haystack[hi + ni].lower()
                n = needle[ni].lower()
                if h == n: # new
                    #match=False
                    results.append(hi) # new
                    break # new
            #if match:
            #results.append(hi)
    return results

```

```

Find: 0.259516 seconds
Find2: 0.057128 seconds
Find3: 0.053053 seconds
Find4: 0.048197 seconds

```

# Using built-in libraries is usually the fastest...

```
def find_ignore_case5( needle, haystack):  
    return [haystack.find(needle)]
```

```
Find: 0.259516 seconds  
Find2: 0.057128 seconds  
Find3: 0.053053 seconds  
Find4: 0.048197 seconds  
Find5: 0.000172 seconds
```



Q: How is this so much faster?

```
def find_ignore_case5( needle, haystack):  
    return [haystack.find(needle)]
```

Q: How is this so much faster?

```
def find_ignore_case5( needle, haystack):  
    return [haystack.find(needle)]
```

A: It built into Python. So it runs in assembly!

# Coding for Performance

Engr 315: Hardware / Software Codesign  
Andrew Lukefahr  
*Indiana University*

