


09: Memory Translation

Engr 315: Hardware / Software Codesign

Andrew Lukefahr

Indiana University



Some material taken from EECS370 at U. of Michigan

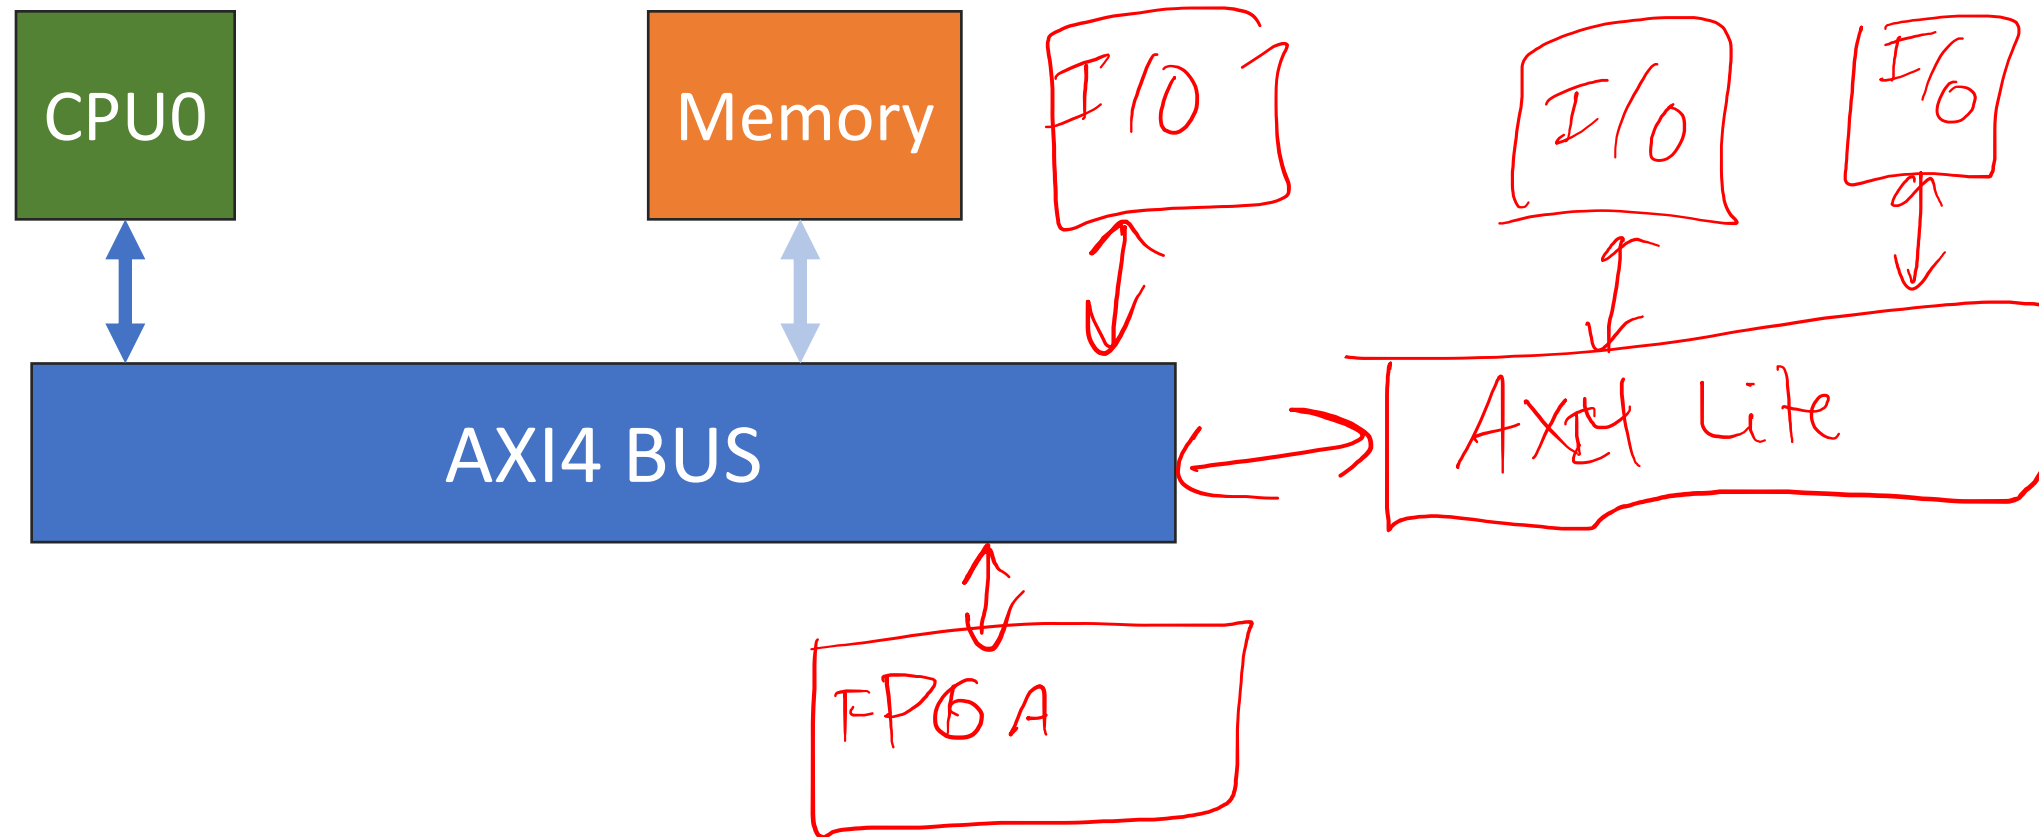
Announcements

- P3 is out
 - Pushed back to 11:59 Wednesday (2/23)
- P4 is out.
 - Expect some revisions

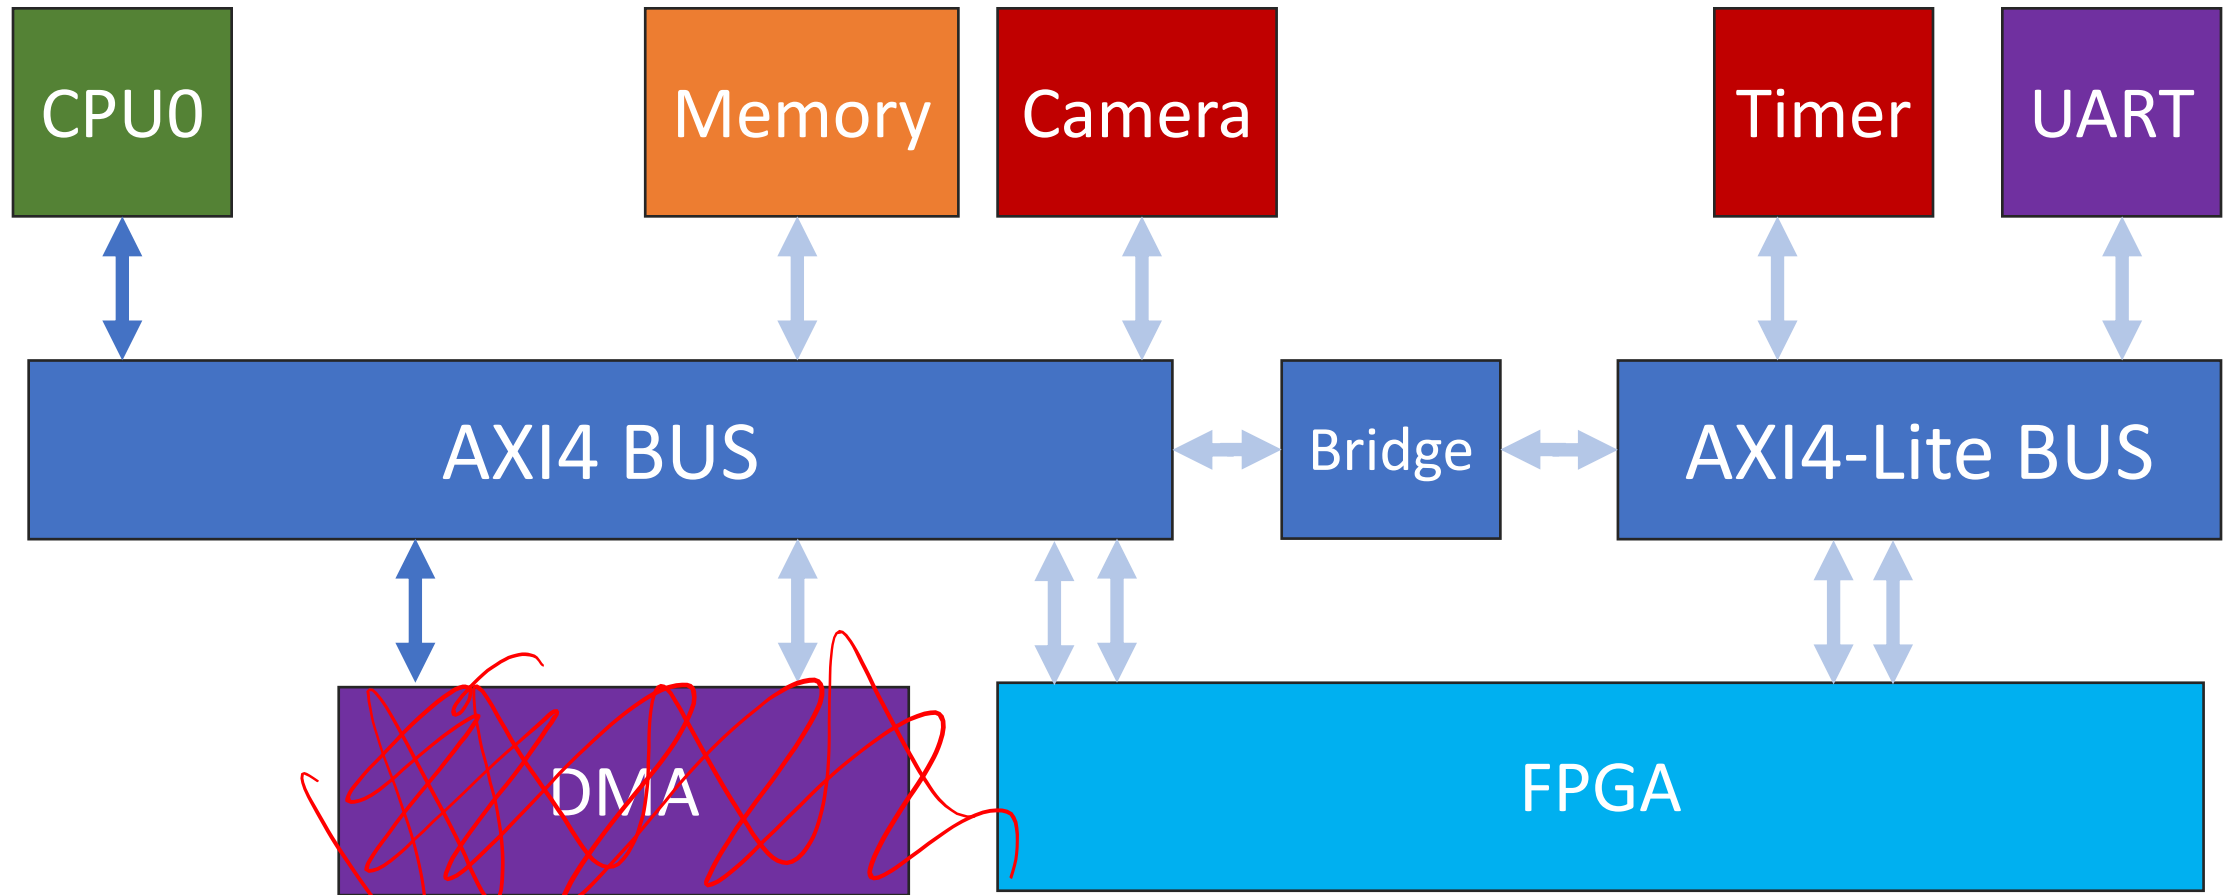
Machine Model, Version 0



Machine Model, V1

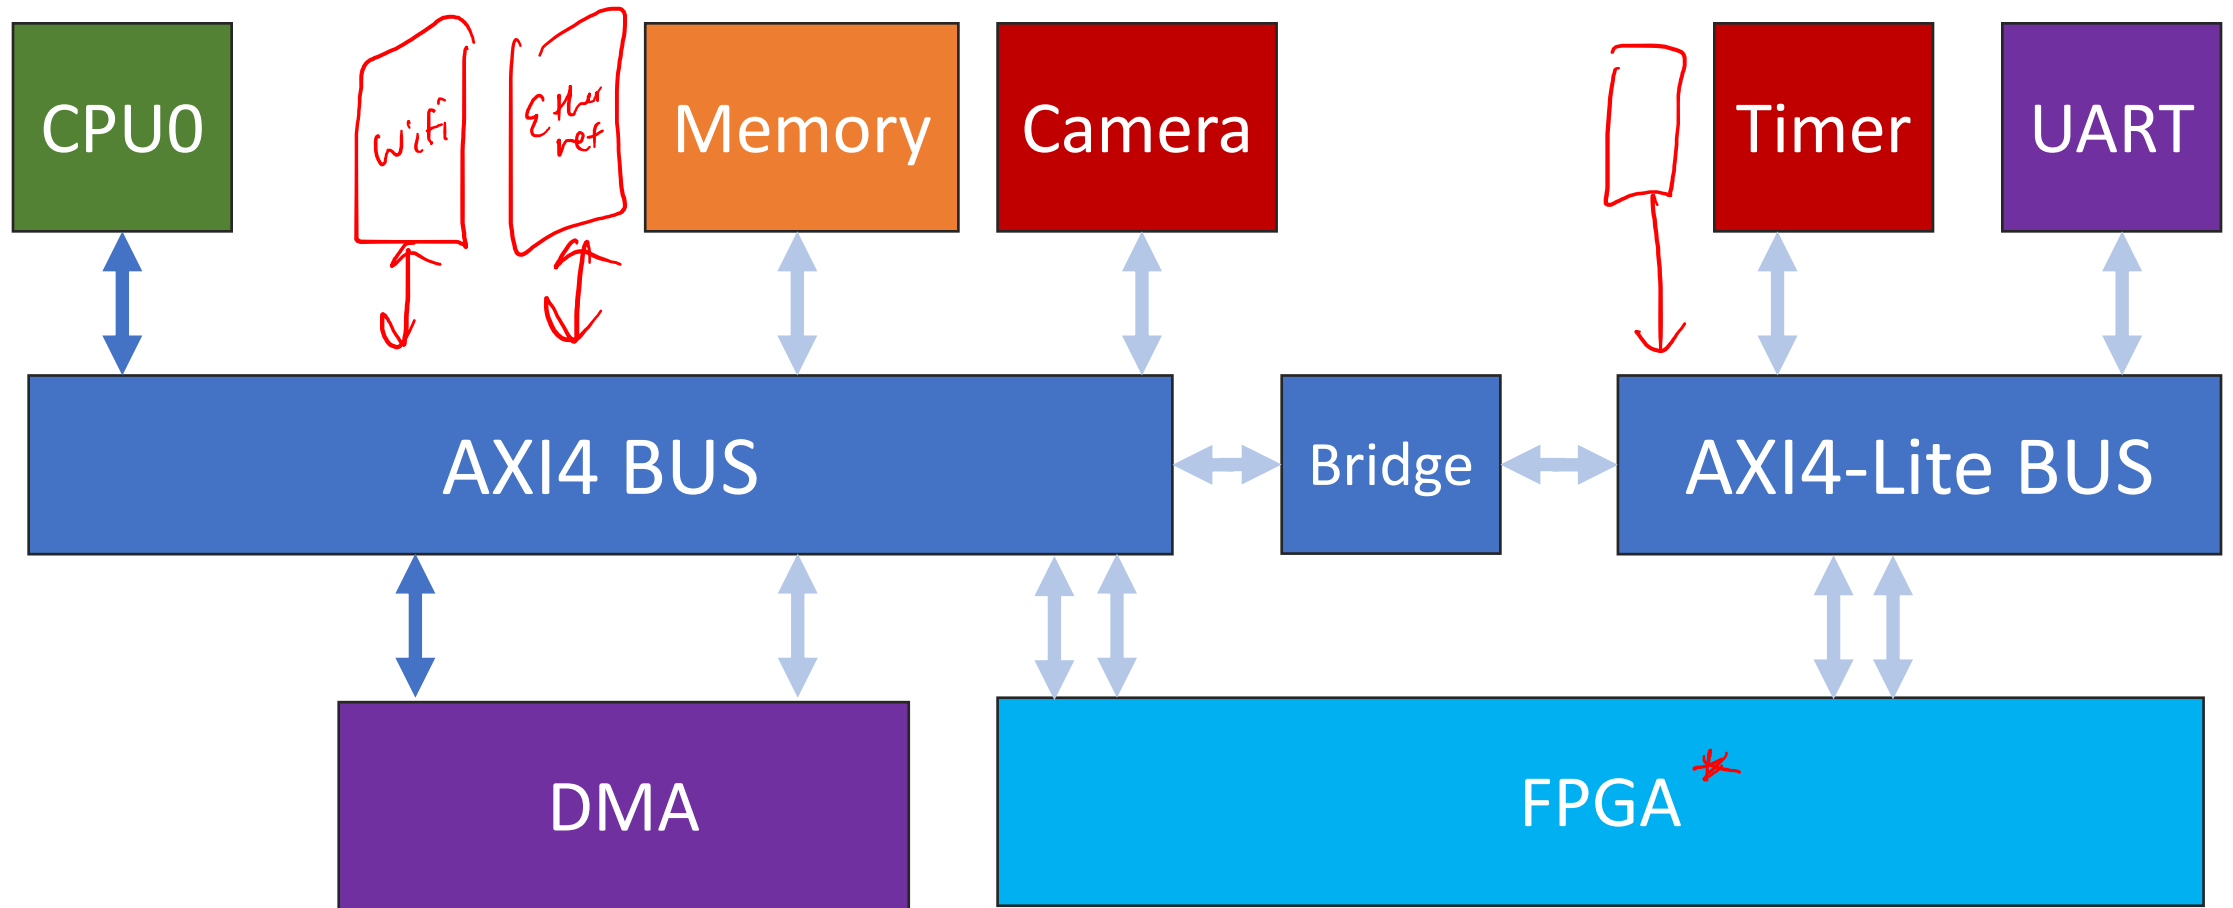


Machine Model, V2




Ignore for now...

Machine Model, V2



MMIO from C.

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#define EMA_MMIO 0x40000000
int main () {
    volatile uint32_t * ema_ptr = (uint32_t*)(EMA_MMIO);
    int32_t val = 0x1000;
    while (1) {
        //push new value into EMA
        *ema_ptr = val;
        //load value from EMA
        val = *ema_ptr;
        printf("Val: %d\n", val);
    }
    return 0;
}
```



MMIO from C.

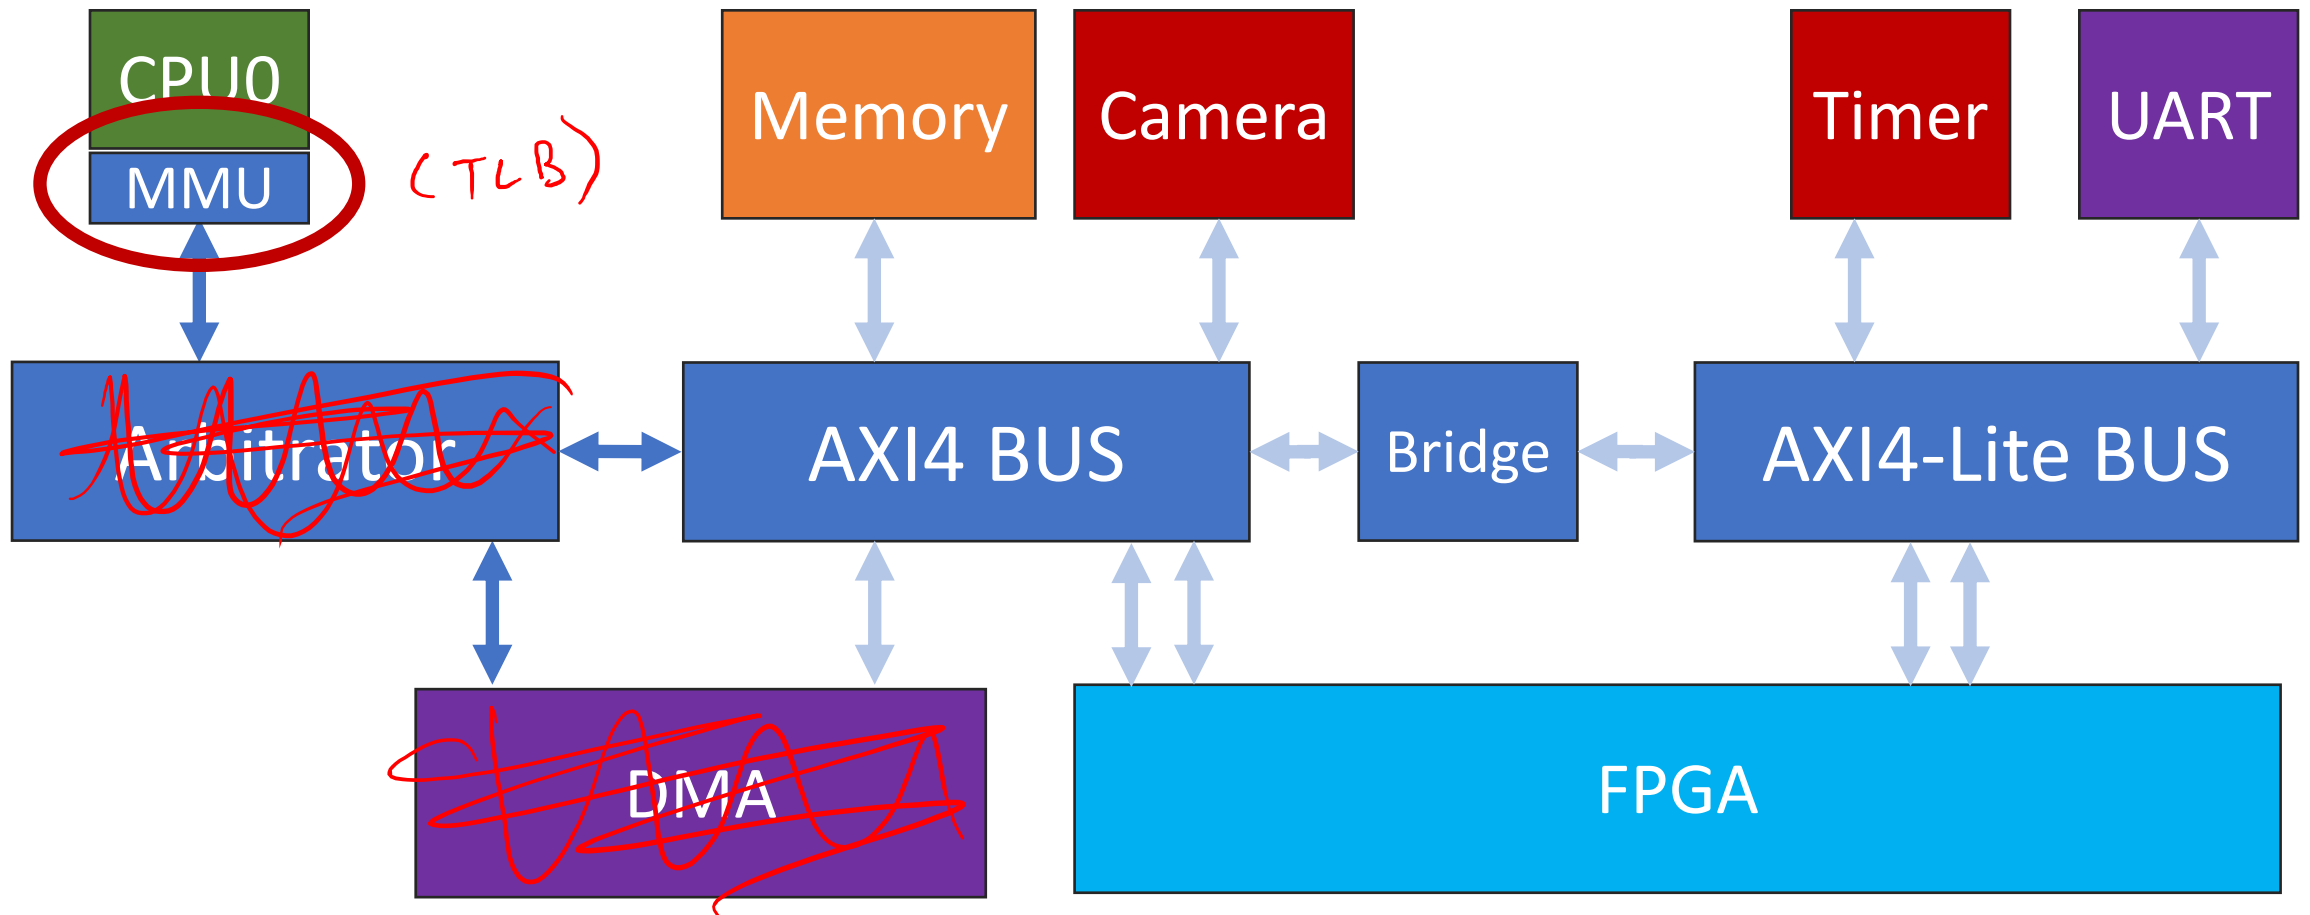
```
#define EMA_MMIO 0x400000000
volatile uint32_t * ema_ptr = (uint32_t*)(EMA_MMIO);
//push new value into EMA
    *ema_ptr = val;
//load value from EMA
    val = *ema_ptr;
```

```
volatile uint32_t * ema_ptr = (uint32_t*)(EMA_MMIO);
8224:    e3a05101    mov r5, #1073741824 ; 0x400000000
    *ema_ptr = val;
822c:    e5854000    str r4, [r5]
    val = *ema_ptr;
8230:    e5954000    ldr r4, [r5]
```


MMIO from C


- DON'T WORK!
- Why?
- Linux... and MMIOs

Machine Model, V3: MMUs




MMU: Memory Management Unit (TLB)

- Rejects load + stores that are “unauthorized”
- Translates addresses (Later)



MMUs track the following things



- **BASE ADDRESS**: the start of a memory region that is allowed through the MMU

Blue: 0xf
Or: 0xAB

- **OFFSET**: the size of a memory region that is 8 bytes allowed through the MMU

Blue: $0xf + 0x8 = 0x17$ offset = 0x8
Orange: offset = 8 \Rightarrow 0xAB \Rightarrow B3

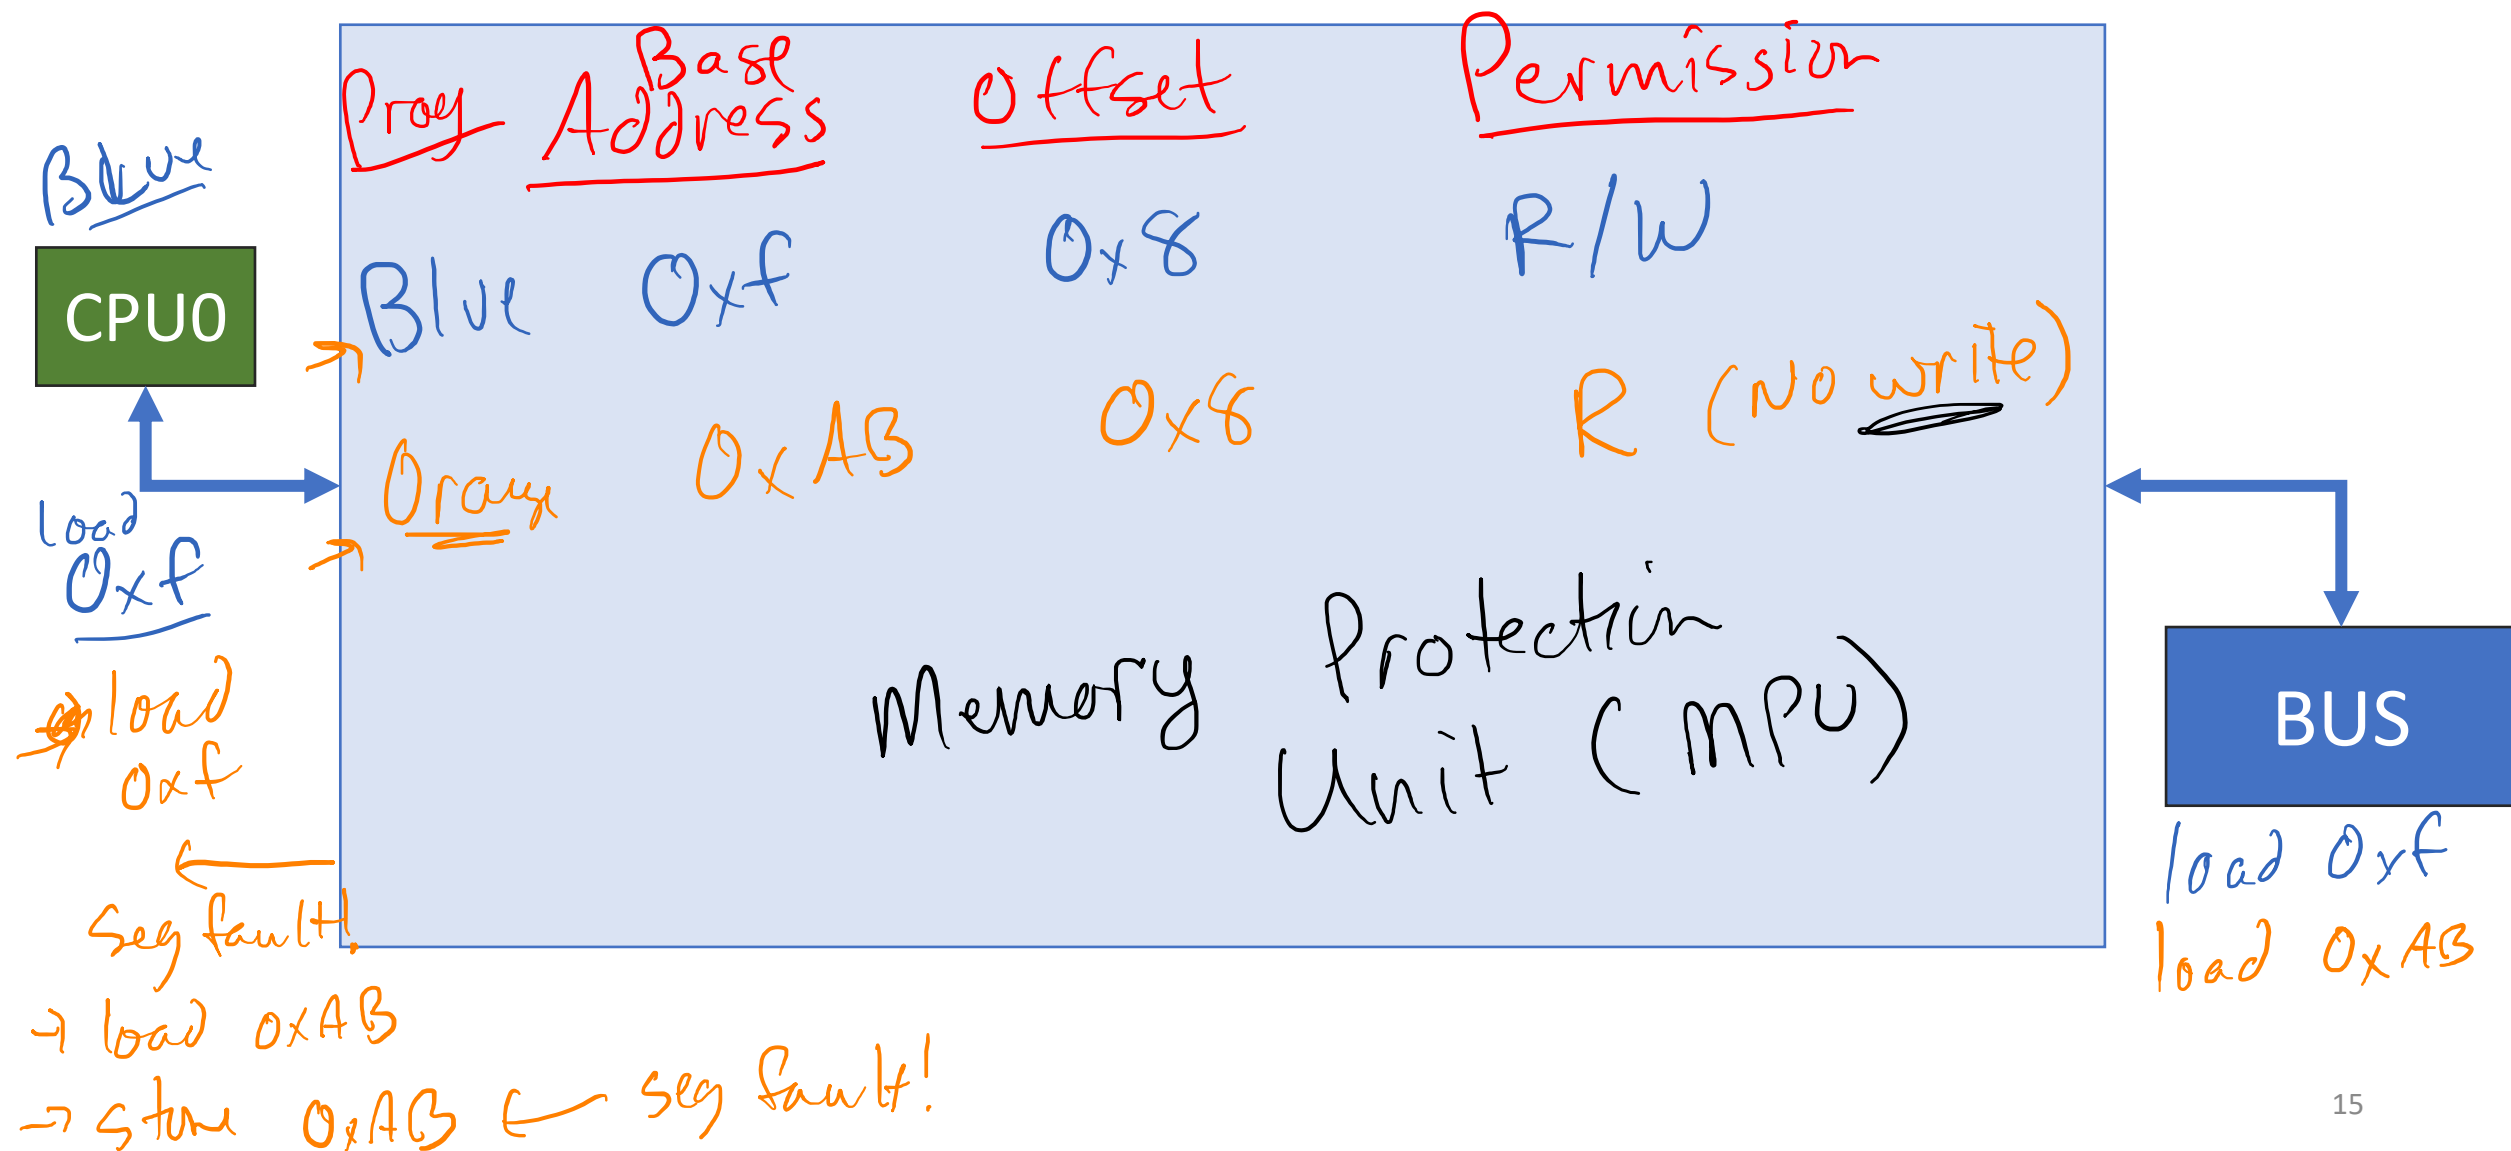
- **Permission**: the type of access that is allowed through the MMU

Read / Write


Blue: R/W for Blue None for Orange

Orange: None for Blue R/W for Orange

Basic MMU Table



Memory Protections Graphically



Why?

- Security
 - Keep you from modifying the code
 - Keep you from executing the data
- Separate multiple applications

Separating multiple applications

- A) What if two applications want to use the same memory address?
- B) How do I prevent your application from modifying my memory?

Two application test..

```
#include <stdio.h>
#include <stdlib.h>

volatile int avalue = 2;

int main ()
{
    while (avalue == 2) { ; }
    return 0;
}
```

```
#include <stdio.h>
#include <stdlib.h>

volatile int avalue = 3;

int main ()
{
    while (avalue == 3) { ; }
    return 0;
}
```

Two application test..

```
gcc -g -O0 test_1.c -o test_1.out  
objdump -DSs test_1.out > test_1.dis  
gcc -g -O0 test_2.c -o test_2.out  
objdump -DSs test_2.out > test_2.dis  
  
vi test_1.dis test_2.dis
```

Two application test..

00011008 <avalue>:

volatile int avalue = 2;

~~11008:~~ 00000002
ff11008

andeq r0, r0, r2

00011008 <avalue>:

volatile int avalue = 3;

~~11008:~~ 00000003
AA11008

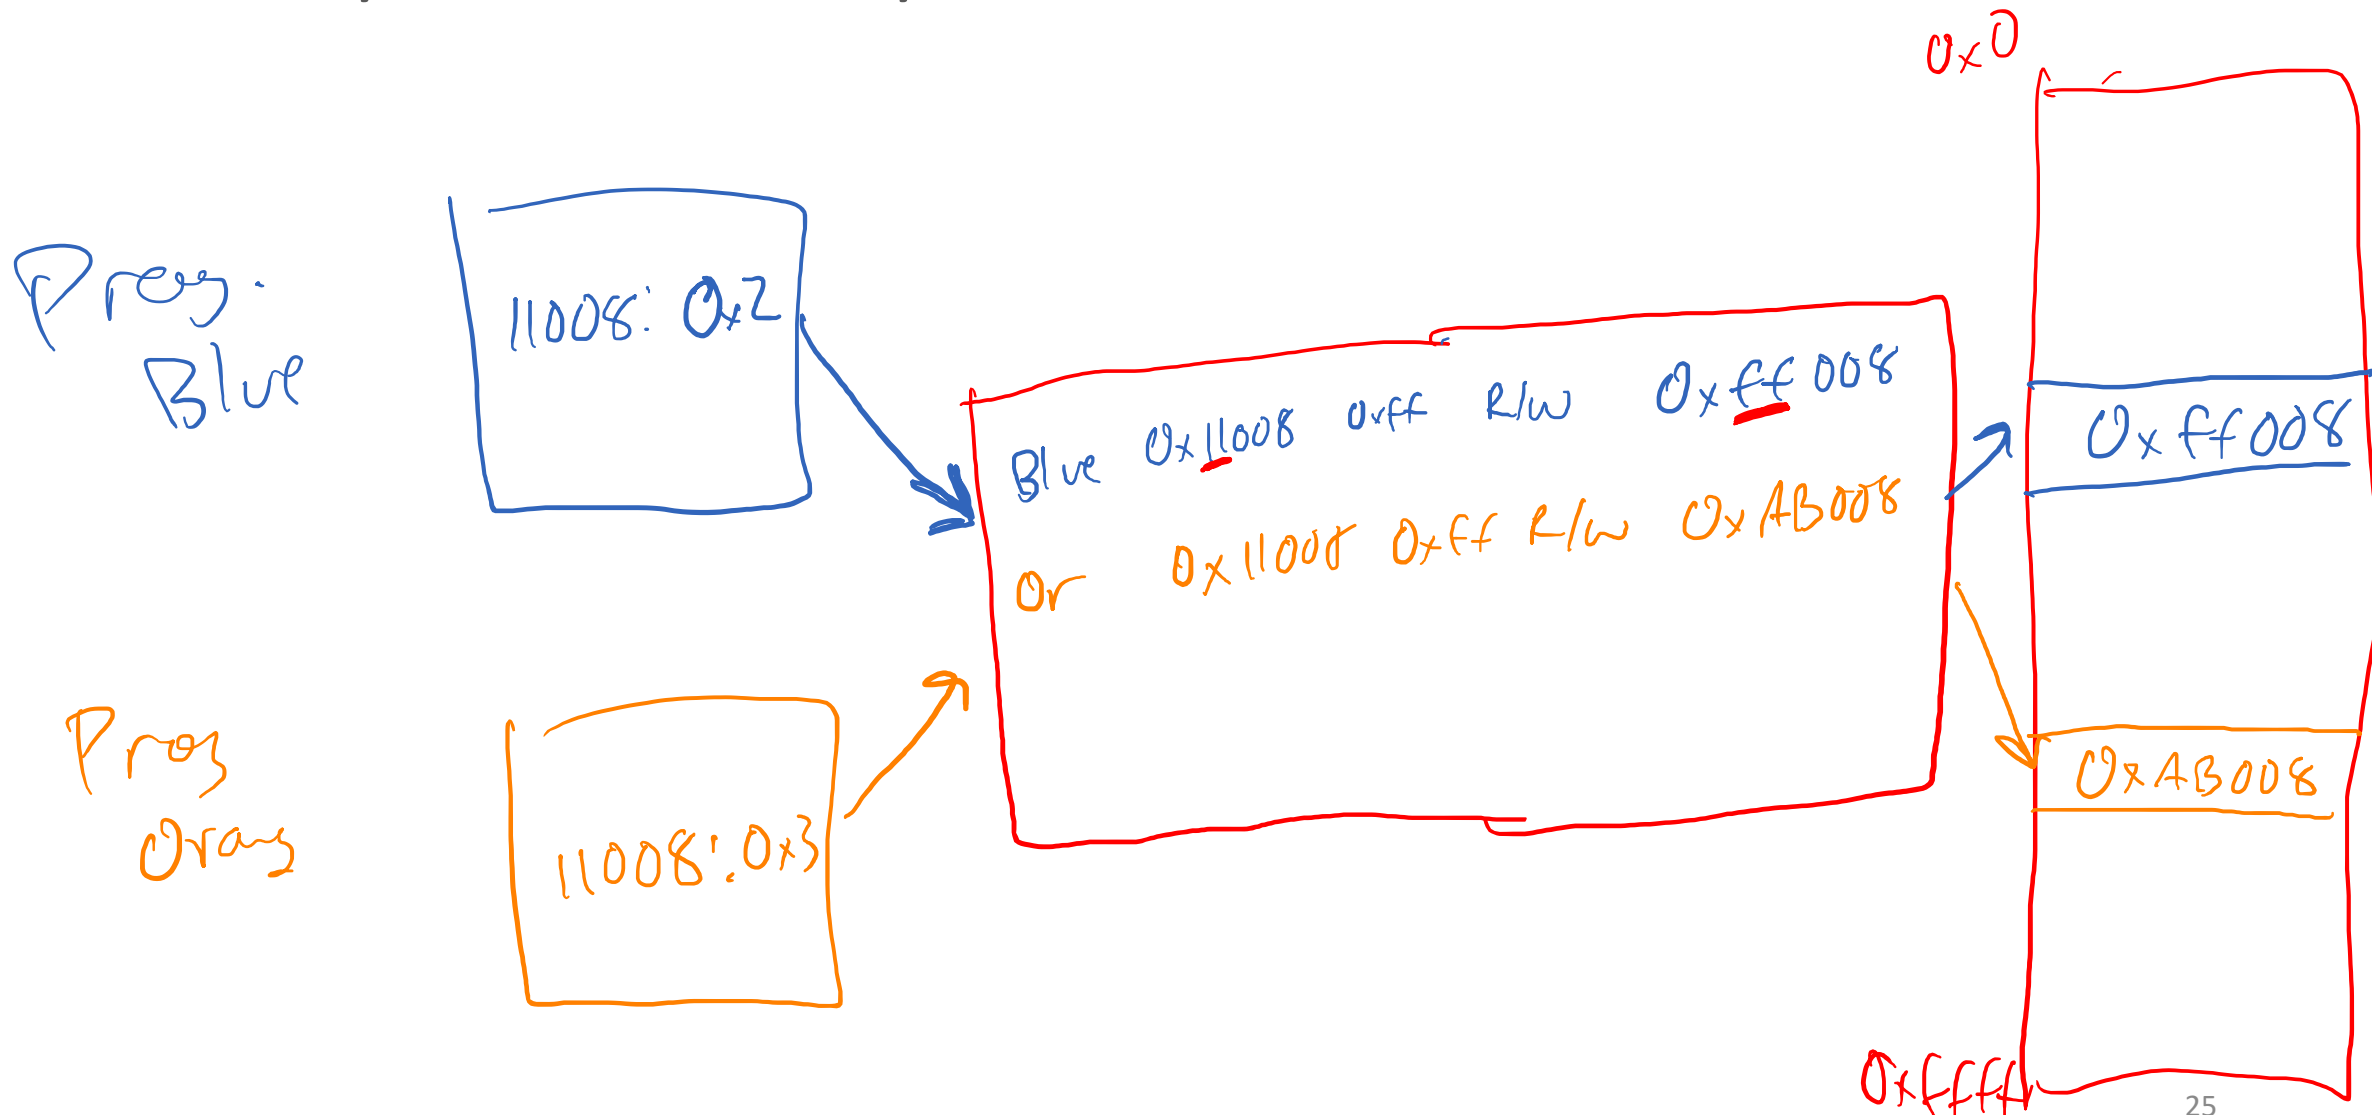
andeq r0, r0, r3

Two application test..


```
./test_1.out &  
./test_2.out &
```

```
top
```

Q: Why didn't they clobber each other?



Q: Why didn't they clobber each other?



Q: Why didn't they clobber each other?

- A: MMUs are doing something else... “**Virtual Memory**”

Virtual memory with an MMU

- MMU automatically translates each memory reference from a

(CPU) virtual address

(which the programmer sees as a huge array of bytes)


to a

(Memory) physical address


(which the hardware uses to identify where the storage actually resides)

Basics of Virtual Memory

- Any time you see the word virtual in computer science and architecture it means “using a level of indirection”
- Virtual memory hardware changes the virtual address the programmer sees into the physical one the memory chips see



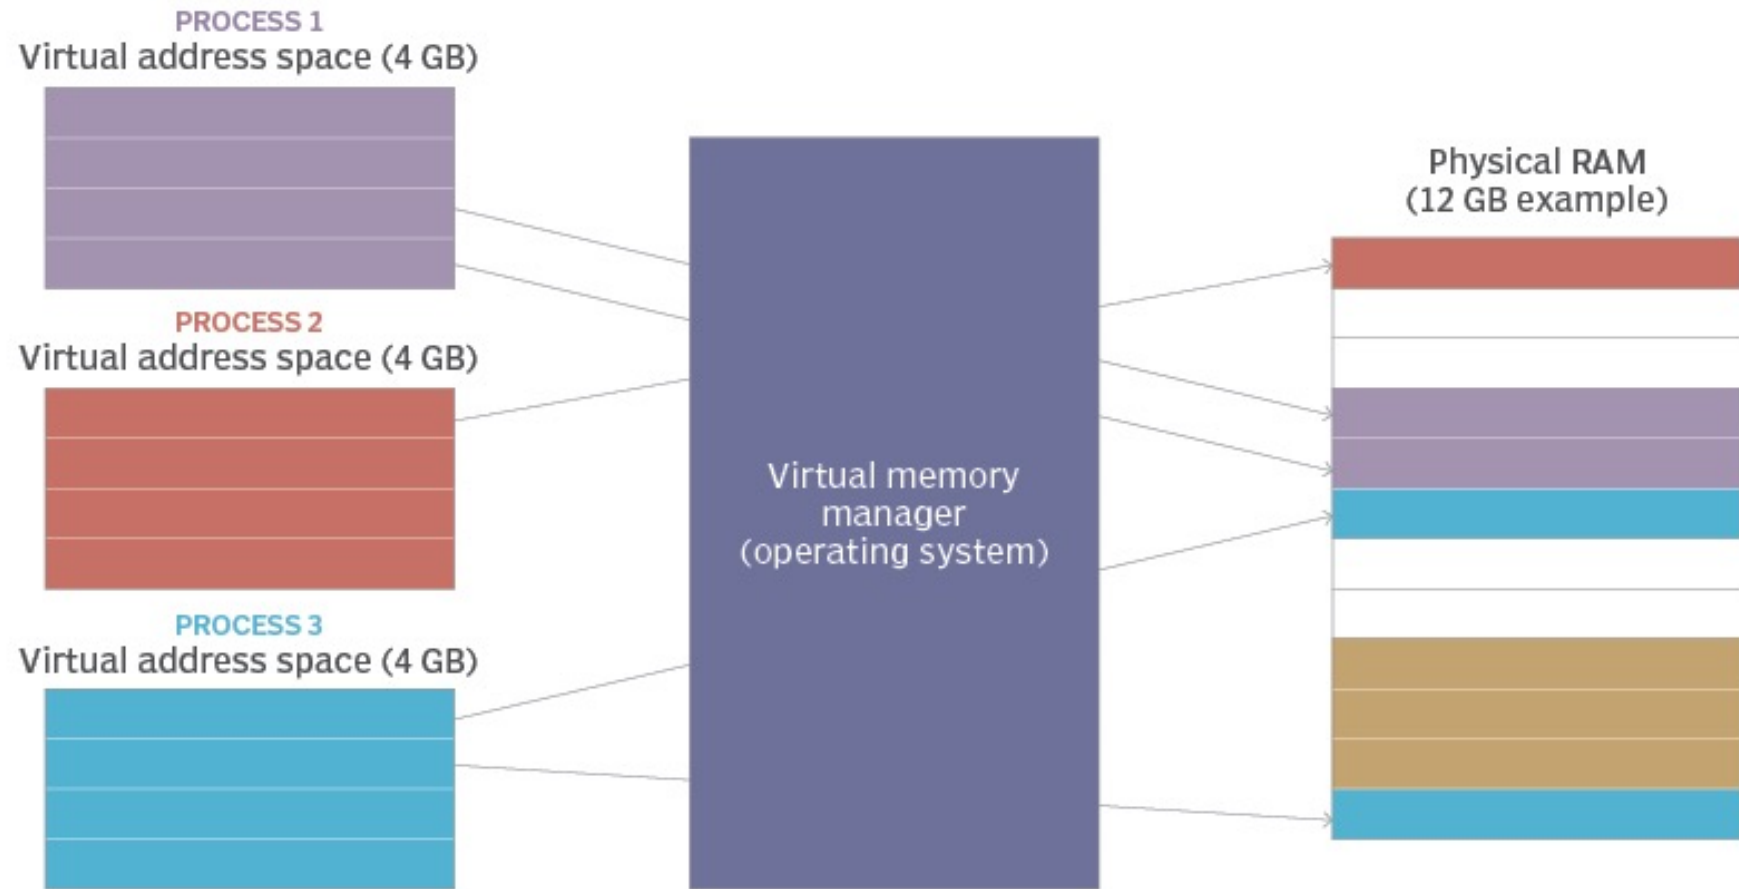
Virtual Memory View

- Virtual memory lets the programmer address a memory array **larger** than the DRAM available on a particular computer system OS (pages)
- Virtual memory enables multiple programs to share the physical memory without:
 -  Knowing other programs exist (**transparency**)
 - Worrying about one program modifying the data contents of another (**protection**)

Managing virtual memory


- Managed by hardware logic *and* operating system software
 - Hardware for speed
 - Software for flexibility and because disk storage is controlled by the operating system
- The hardware must be designed to support Virtual Memory

OS (Linux) mains full Virtual->Physical Mappings

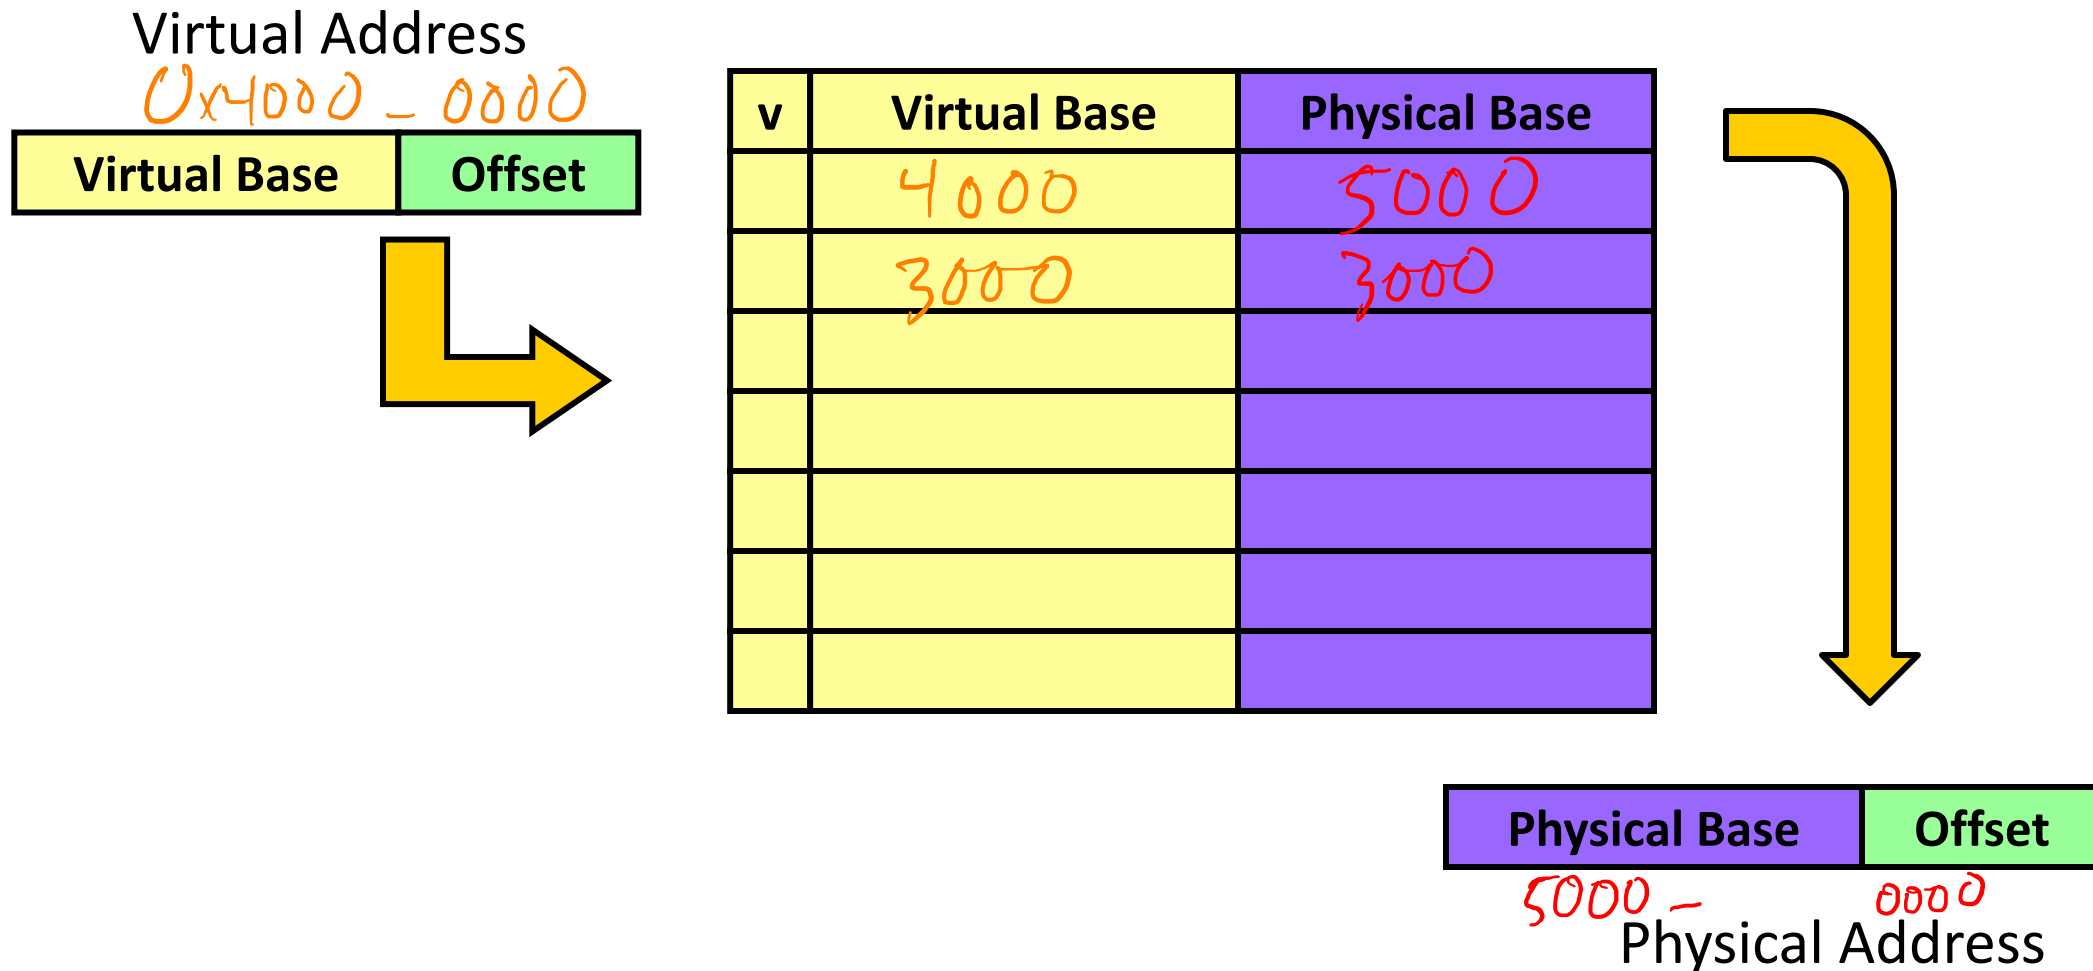


Hardware uses TLBs (Translation Look-aside Buffers)

- ❑ Buffer common Virtual->Physical translations in a **Translation Look-aside Buffer (TLB)**, a fast cache memory dedicated to storing a small subset of valid translations
- ❑ 16-512 entries common
- ❑ Generally has low miss rate (< 1%)



MMU Address Translation




MMU Address Translation




unity mapping

Virtual Address
(32-bit)

Virtual Base	Offset
--------------	--------



v	Virtual Base	Physical Base
	69ff	4000
	3f00	4f00
	2000	2000
	6000	8000
	7000	8000




0x4000 0004
2000 0008

Physical Base	Offset
---------------	--------

Physical Address

memory (mmio)

MMU Table



So how do we access the FPGA for P4?

- FPGA uses **physical** address
- CPU (w/Linux) uses **virtual** address
- Q: How do I talk to a physical address with Linux?
- A: Linux provides a special `/dev/mem` file to help us!

Next Time:

- /dev/mem
- /dev/uio

09: Memory Translation

Engr 315: Hardware / Software Codesign
Andrew Lukefahr
Indiana University

