


09: Memory Translation

Engr 315: Hardware / Software Codesign

Andrew Lukefahr

Indiana University



Some material taken from EECS370 at U. of Michigan

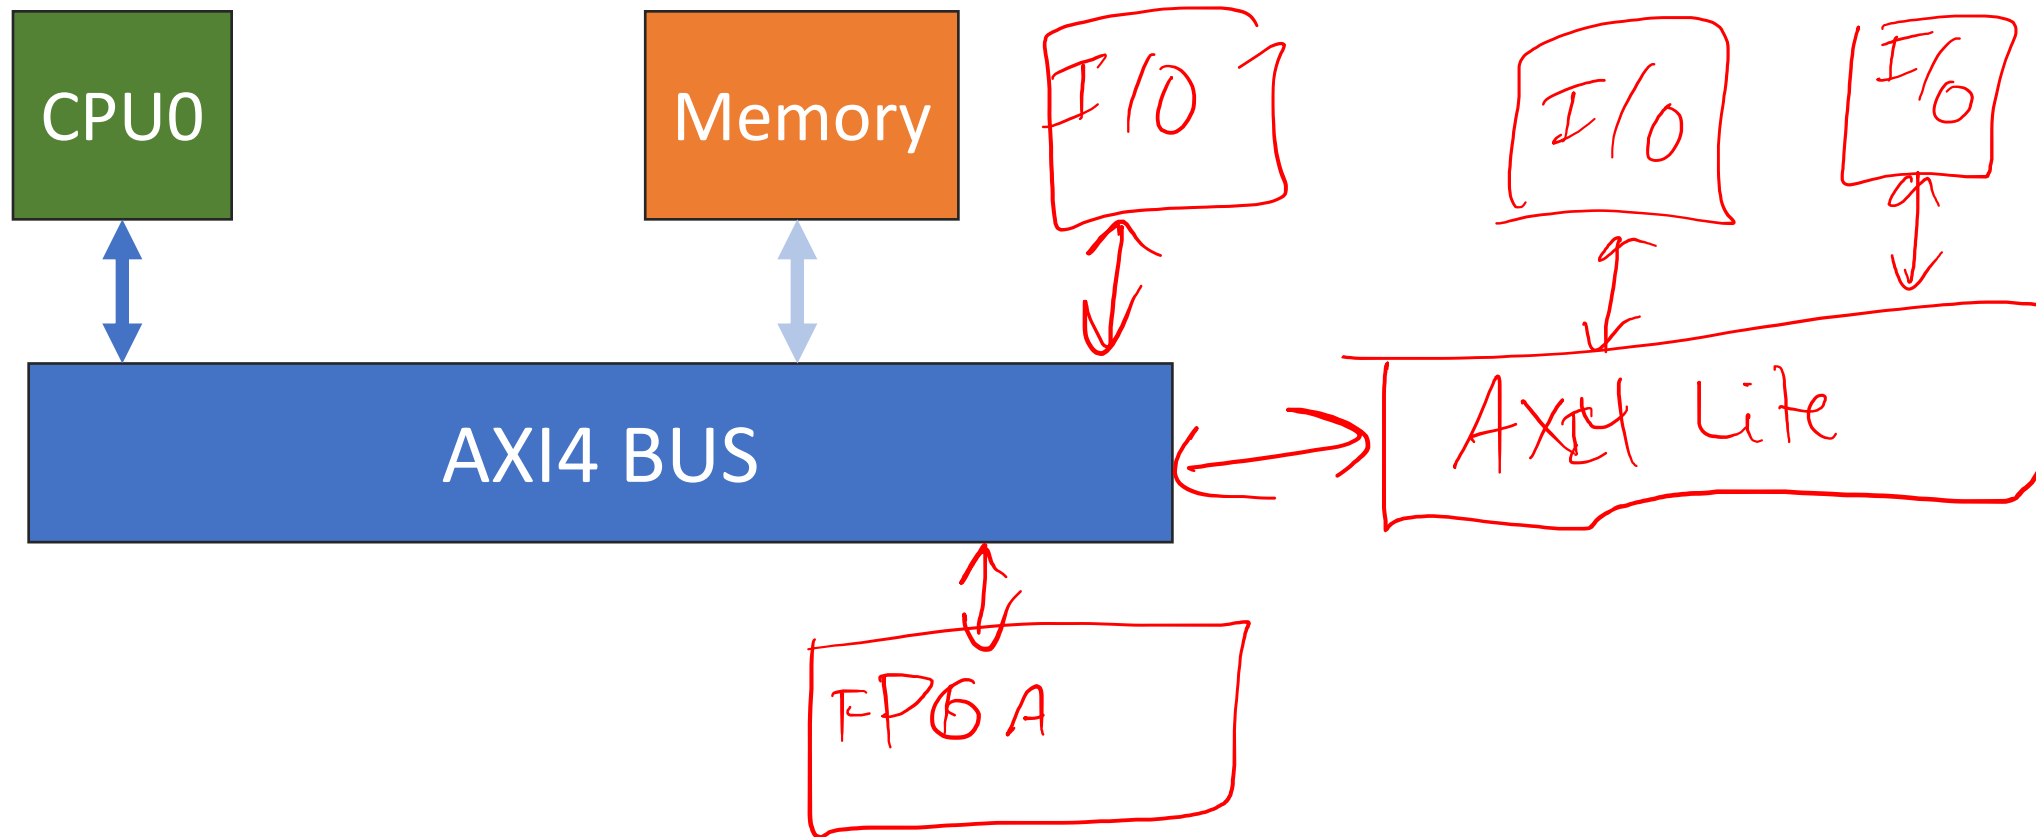
Announcements

- P4: Due Next Wednesday. *→ AG + 6 different from project.*
- P5: Out soon.

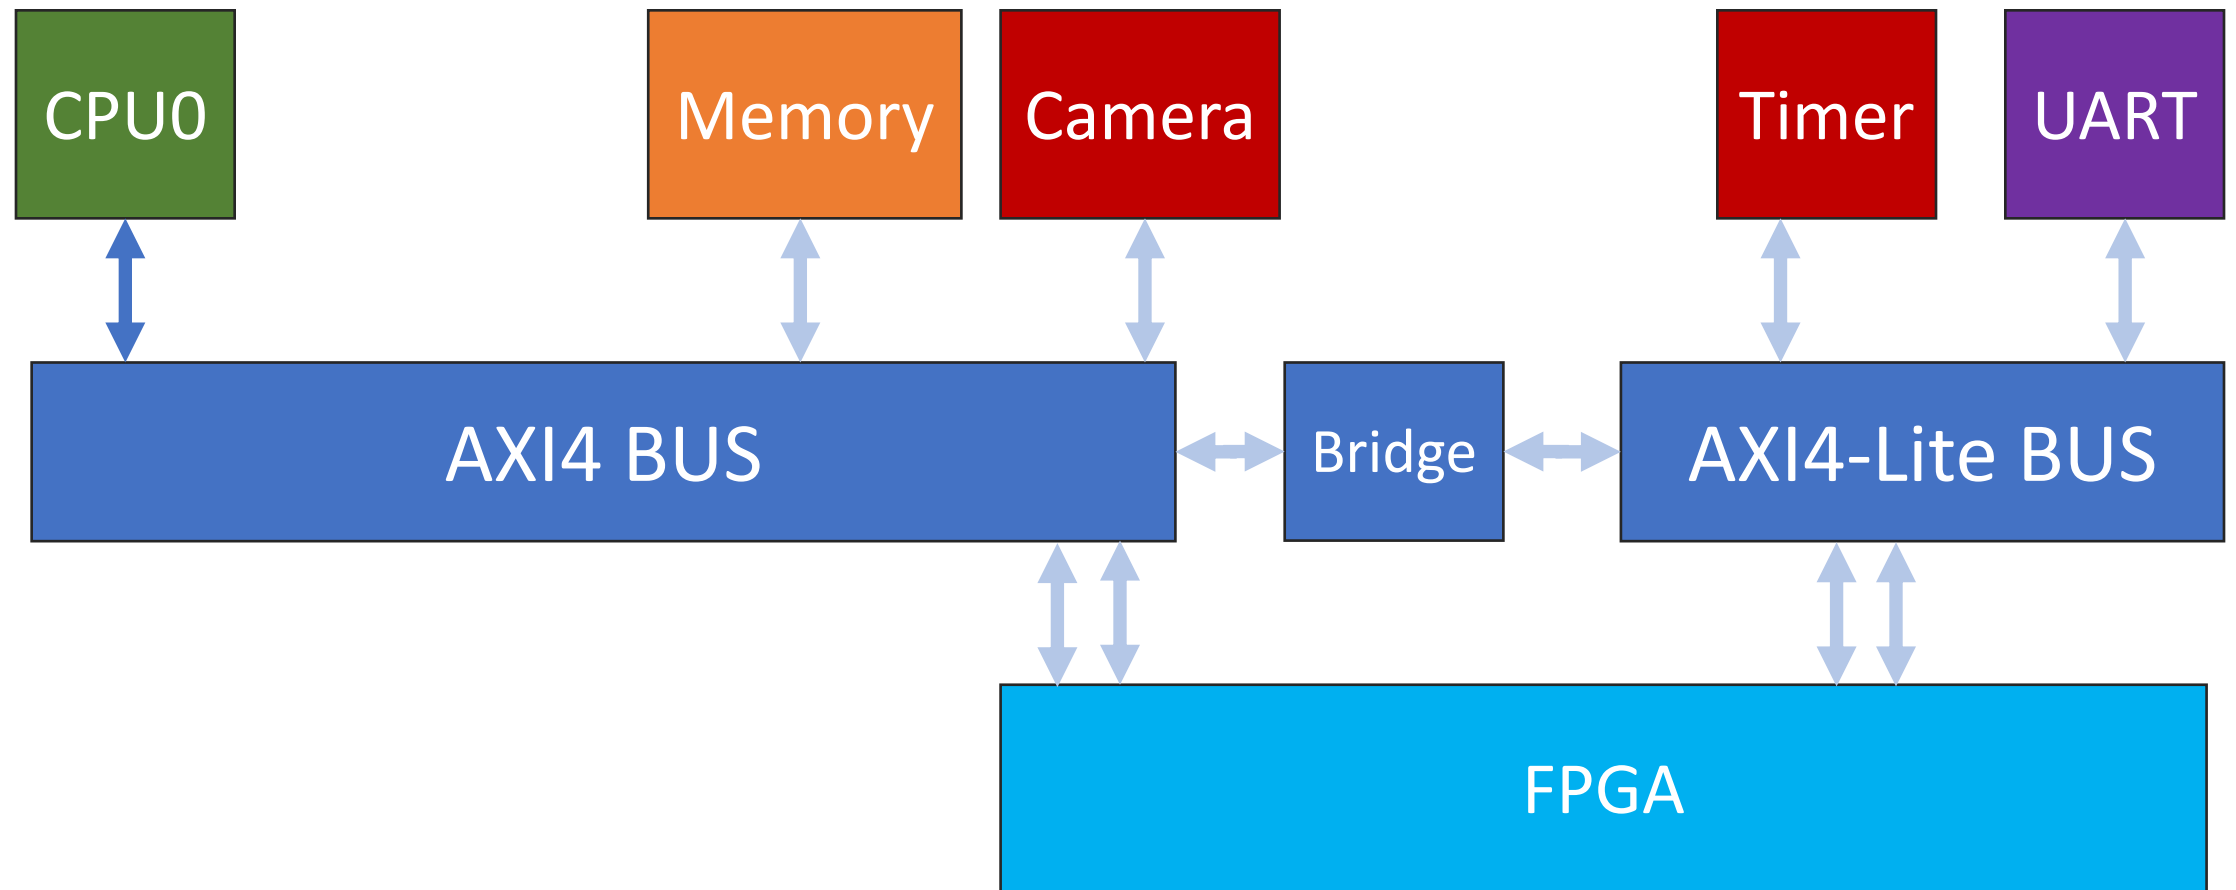
Machine Model, Version 0



Machine Model, V1



Machine Model, V2



MMIO from C.

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#define EMA_MMIO 0x40000000
int main () {
    volatile uint32_t * ema_ptr = (uint32_t*)(EMA_MMIO);
    int32_t val = 0x1000;
    while (1) {
        //push new value into EMA
        *ema_ptr = val;
        //load value from EMA
        val = *ema_ptr;
        printf("Val: %d\n", val);
    }
    return 0;
}
```

MMIO from C.


```
#define EMA_MMIO 0x400000000
volatile uint32_t * ema_ptr = (uint32_t*)(EMA_MMIO);
//push new value into EMA
    *ema_ptr = val;
//load value from EMA
    val = *ema_ptr;
```

```
volatile uint32_t * ema_ptr = (uint32_t*)(EMA_MMIO);
8224:    e3a05101    mov r5, #1073741824 ; 0x400000000
    *ema_ptr = val;
822c:    e5854000    str r4, [r5]
    val = *ema_ptr;
8230:    e5954000    ldr r4, [r5]
```

MMIO from C


- WONT WORK!
- Why?
- Linux... and MMIOs

Machine Model, V3: MMUs




MMU: Memory Management Unit (TLB)

- Rejects load + stores that are “unauthorized”
- Translates addresses (Later)




MMUs track the following things




- **BASE ADDRESS:** the start of a memory region that is allowed through the MMU
- **OFFSET:** the size of a memory region that is allowed through the MMU
- **Permission:** the type of access that is allowed through the MMU

Basic MMU Table



Memory Protections Graphically



Why?

- Security
 - Keep you from modifying the code
 - Keep you from executing the data
- Separate multiple applications

Separating multiple applications

- A) What if two applications want to use the same memory address?
- B) How do I prevent your application from modifying my memory?

Two application test..

```
#include <stdio.h>
#include <stdlib.h>

volatile int avalue = 2;

int main ()
{
    while (avalue == 2) { ; }
    return 0;
}
```

```
#include <stdio.h>
#include <stdlib.h>

volatile int avalue = 3;

int main ()
{
    while (avalue == 3) { ; }
    return 0;
}
```


Two application test..

```
gcc -g -O0 test_1.c -o test_1.out  
objdump -DSs test_1.out > test_1.dis  
gcc -g -O0 test_2.c -o test_2.out  
objdump -DSs test_2.out > test_2.dis  
  
vi test_1.dis test_2.dis
```

Two application test..

```
00011008 <avalue>:
```

```
volatile int avalue = 2;
```

```
11008:          00000002          andeq    r0, r0, r2
```

```
00011008 <avalue>:
```

```
volatile int avalue = 3;
```

```
11008:          00000003          andeq    r0, r0, r3
```

Two application test..

```
./test_1.out &  
./test_2.out &
```

```
top
```

Q: Why didn't they clobber each other?

Q: Why didn't they clobber each other?

- A: MMUs are doing something else... “**Virtual Memory**”

Virtual memory with an MMU

- MMU automatically translates each memory reference from a

(CPU) virtual address

(which the programmer sees as a huge array of bytes)


to a

(Memory) physical address

(which the hardware uses to identify where the storage actually resides)

Basics of Virtual Memory

- Any time you see the word virtual in computer science and architecture it means “using a level of indirection”
- Virtual memory hardware changes the virtual address the programmer sees into the physical one the memory chips see



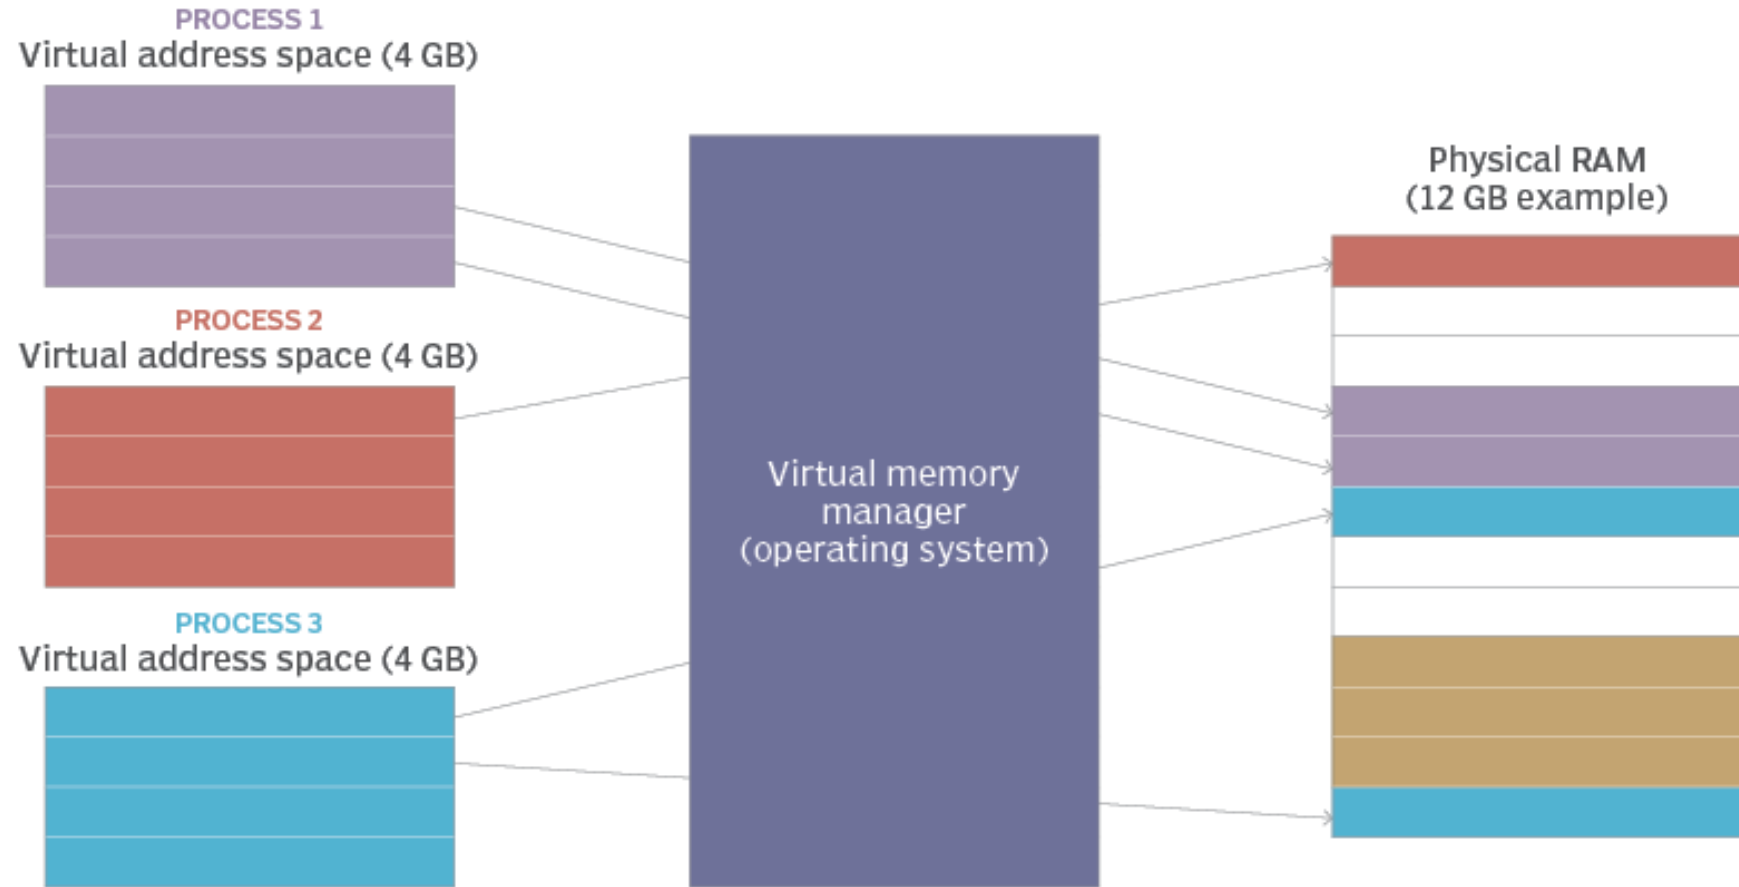
Virtual Memory View

- Virtual memory lets the programmer address a memory array **larger** than the DRAM available on a particular computer system
- Virtual memory enables multiple programs to share the physical memory without:
 - Knowing other programs exist (**transparency**)
 - Worrying about one program modifying the data contents of another (**protection**)

Managing virtual memory


- Managed by hardware logic *and* operating system software
 - Hardware for speed
 - Software for flexibility and because disk storage is controlled by the operating system
- The hardware must be designed to support Virtual Memory

OS (Linux) mains full Virtual->Physical Mappings




Hardware uses TLBs (Translation Look-aside Buffers)


- ❑ Buffer common Virtual->Physical translations in a **Translation Look-aside Buffer (TLB)**, a fast cache memory dedicated to storing a small subset of valid translations
- ❑ 16-512 entries common
- ❑ Generally has low miss rate (< 1%)



MMU Address Translation



MMU Table



So how do we access the FPGA for P4?

- FPGA uses **physical** address
- CPU (w/Linux) uses **virtual** address
- Q: How do I talk to a physical address with Linux?
- A: Linux provides a special `/dev/mem` file to help us!

P3: MMIO Popcount

Question:

- How do we access the FPGA from C w/Linux?
- FPGA uses **physical** address
- CPU (w/Linux) uses **virtual** address

Answer:

- Linux lets us cheat.

Python Example

```
from pynq import Overlay
from pynq import MMIO
class hw_ema():
    def __init__(self):
        self.overlay = Overlay('bitstream.bit')
        self.mmio = self.overlay.axi_popcount_0.S_AXI_LITE
    def ema(self, n):
        self.mmio.write(0x0, int(n))
        return self.mmio.read(0x0)
ema = hw_ema()
for i in range(1000, 6000, 1000):
    x = ema.ema(i)
    print ("In: ", i, " Out: ", x)
```

```
$ sudo python3 mmio_demo.py
```

```
In:  1000  Out:  250
```

```
In:  2000  Out:  687
```

```
In:  3000  Out: 1264
```

```
In:  4000  Out: 1948
```

```
In:  5000  Out: 2711
```

/dev/mem

- */dev/mem* is a character device file that **is an image of the main memory** of the computer. It may be used, for example, to examine (and even patch) the system.
- Byte **addresses in */dev/mem* are interpreted as physical memory addresses**. References to nonexistent locations cause errors to be returned.
- Requires root (sudo) access



/dev/mem

open file


read()
→
read()

file

super
super
powerful!
dangerous!

but what Jupyter or
Pygments does!

seek(0x40000000)
read() → 0x40000000
write()



Next Time:

- /dev/mem
- /dev/uio

09: Memory Translation

Engr 315: Hardware / Software Codesign

Andrew Lukefahr

Indiana University

