


11: Multi-Master Buses

Engr 315: Hardware / Software Codesign
Andrew Lukefahr
Indiana University



Announcements

- P3 is due Wednesday
- P4 is out.
 - Expect some revisions
 - Use P3 bitstream / hwh files
- P5 is out.

PG "commits"

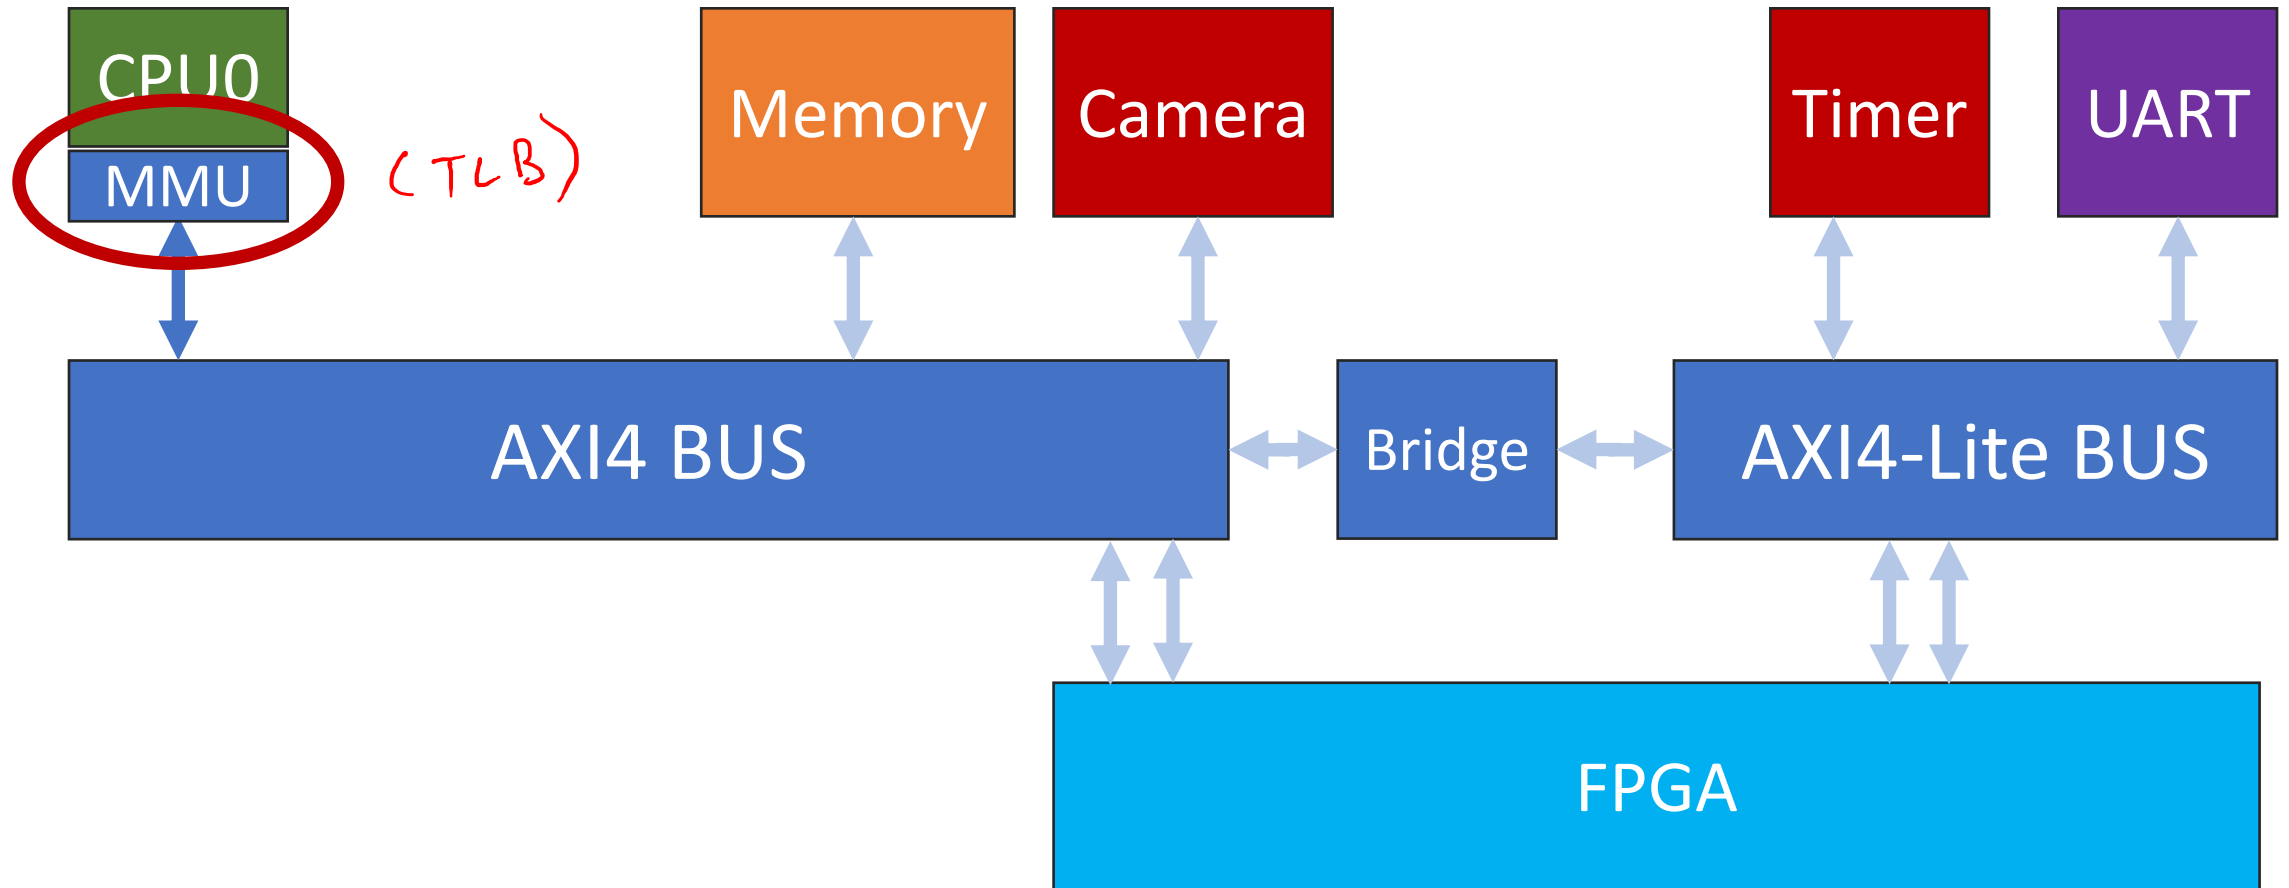
Exam Planning

3/15	Tuesday	–	SPRING BREAK	
3/17	Thursday	–	SPRING BREAK	P7 Dot (V)
3/22	Tuesday	17	Review	
3/24	Thursday	18	Exam Review	P7 Dot (V)
3/29	Tuesday	19	Linux Kernel I Exam	
3/31	Thursday	20	Linux Kernel II	PX Accel Dot (V)
4/05	Tuesday	24	Linux Kernel III	

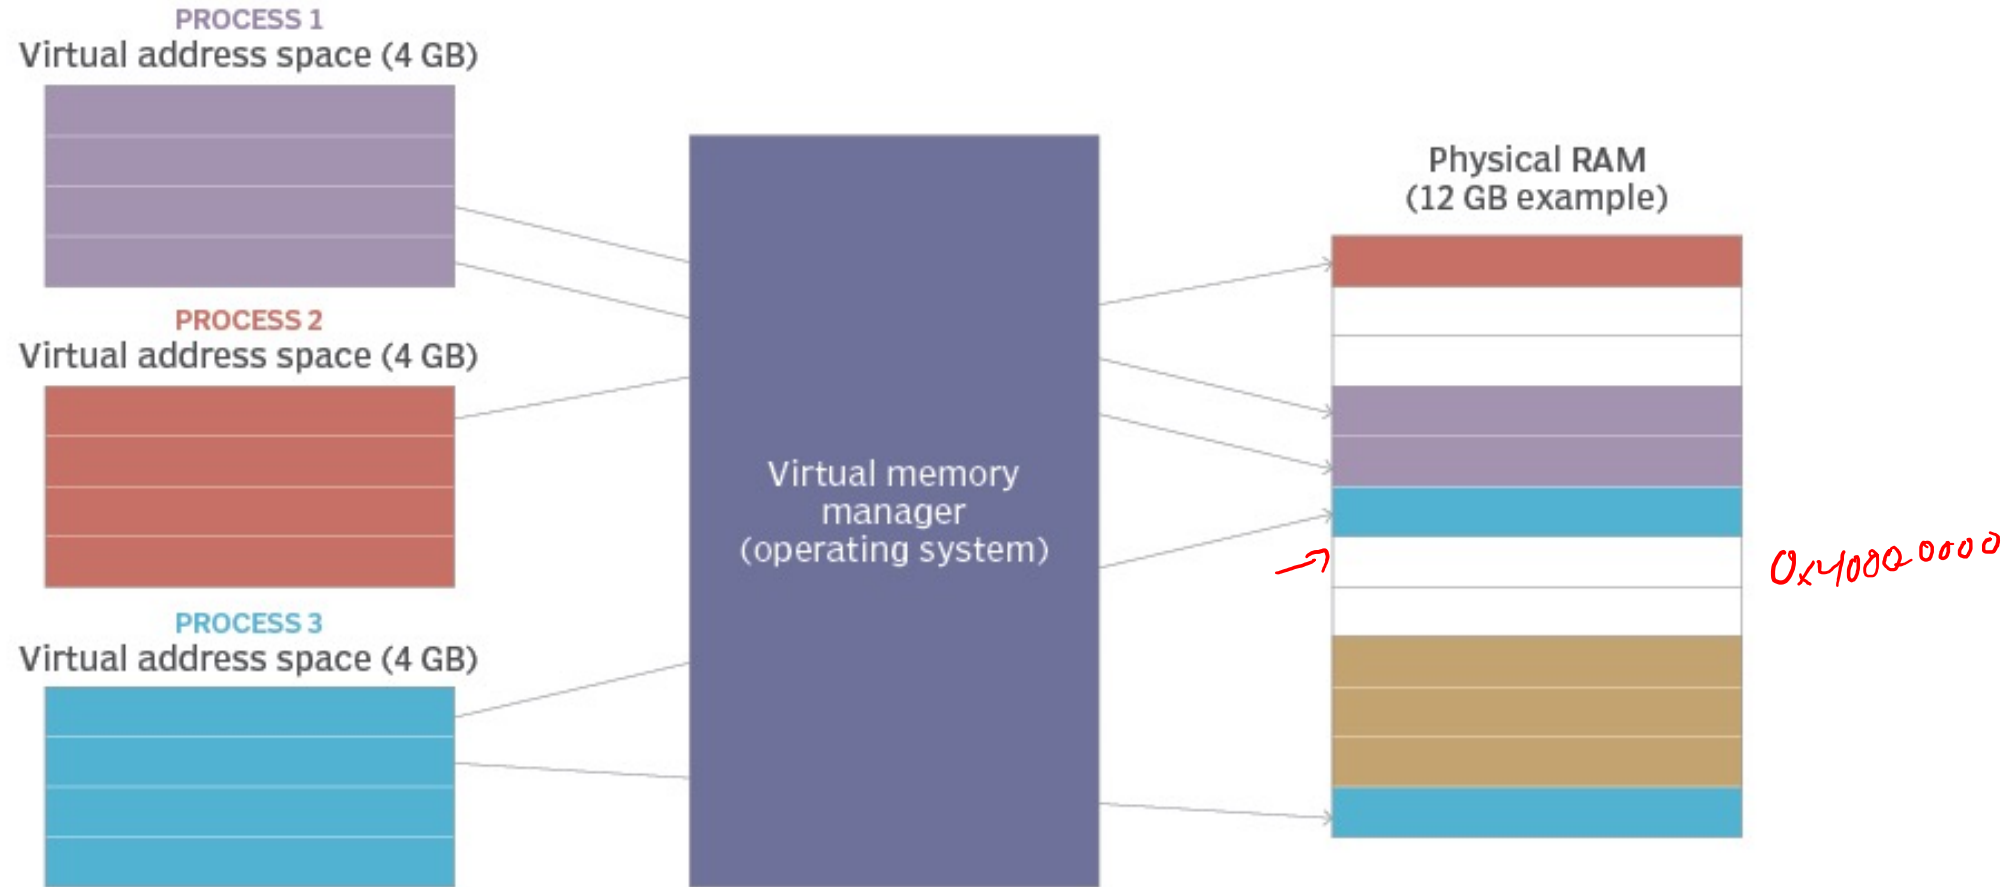
if objections:
please let me know by Tues

2/22

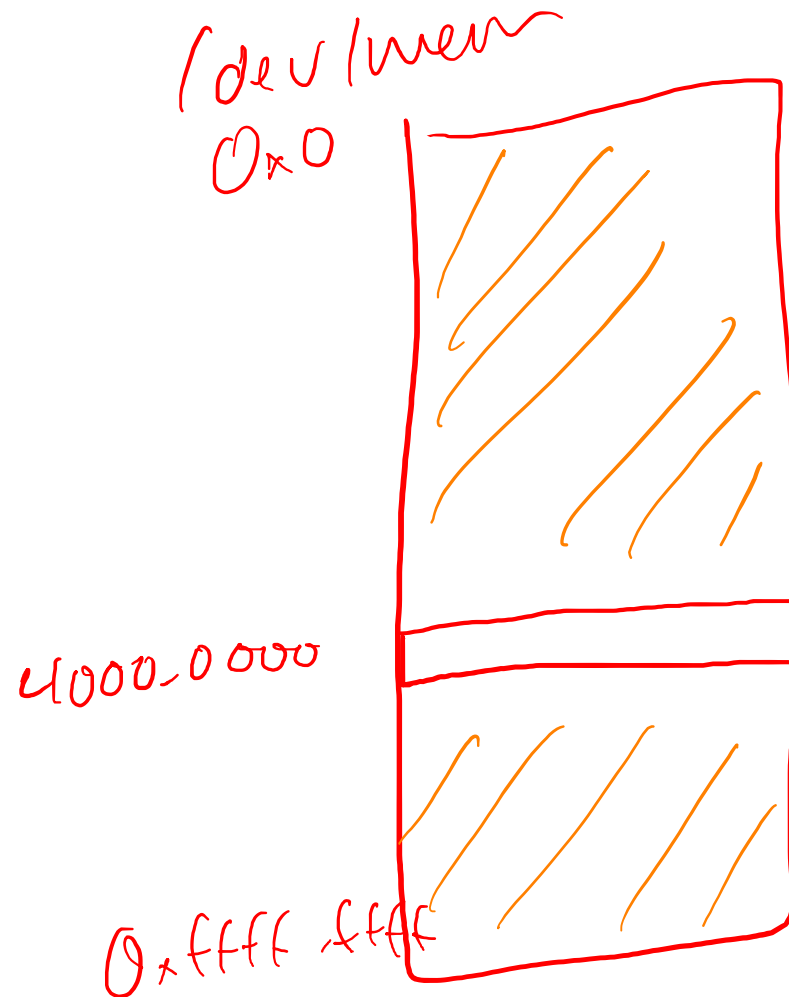
Machine Model, V3: MMUs



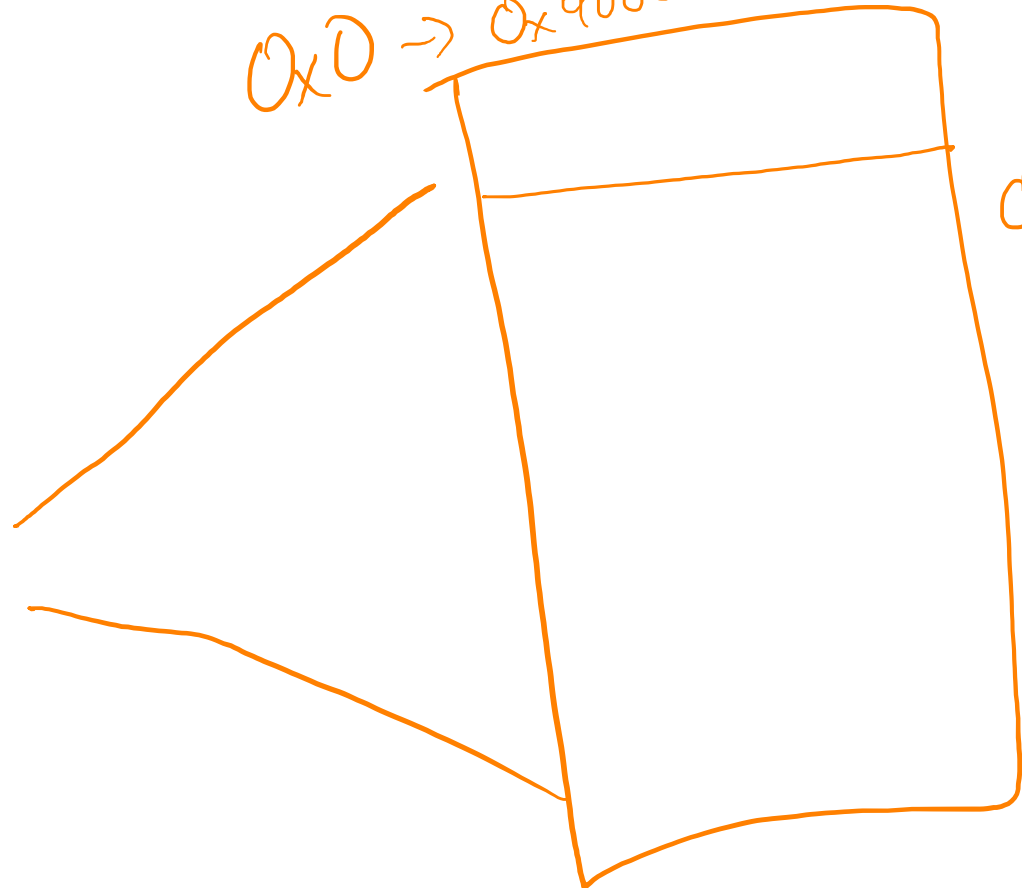
OS (Linux) mains full Virtual->Physical Mappings



/dev/urandom



/dev/urandom
0x0 → 0x4000-0000



0x4 →
~~0x4000-0000~~
0x4000-0004


mmap

physical
mem

Linux
/dev/mem

Linux
/dev/urandom

Linux
mmap



Userspace EMA w/C

```
int main (){

    int dev_mem_fd = -1;
    void * vaddr_base;

    //Mapping user-space I/O
    dev_mem_fd = open("/dev/uio0", O_RDWR|O_SYNC);
    if (dev_mem_fd < 0) { perror("open() /dev/uio0"); return 1; }

    // Map 1KB of physical memory starting at uio 0x0 (real 0x40000000)
    // to 1KB of virtual memory starting at vaddr_base
    vaddr_base = mmap(0, 1024, PROT_READ|PROT_WRITE,
        MAP_SHARED, dev_mem_fd, 0x0);

    if (vaddr_base == MAP_FAILED) { perror("mmap()"); return 1; }
```

```
volatile uint32_t * ema_reg = (uint32_t*) vaddr_base;
uint32_t tmp;

for (int i = 1000; i < 6000; i +=1000){
    printf ("Sending in: %d\n", i);
    *ema_reg = i; //mmio store
    tmp = *ema_reg; //mmio load
    printf("Receiving: %d\n", tmp);
}

if (munmap(vaddr_base, 1024) != 0) { perror("munmap()"); }
if (close(dev_mem_fd) != 0) { perror("close()"); }
dev_mem_fd = -1;

return 0;
}
```


Userspace EMA w/C

```
int main (){

    int dev_mem_fd = -1;
    void * vaddr_base;

    //Mapping user-space I/O
    dev_mem_fd = open("/dev/uio0", O_RDWR|O_SYNC);
    if (dev_mem_fd < 0) { perror("open() /dev/uio0"); return 1; }

    // Map 1KB of physical memory starting at uio 0x0 (real 0x40000000)
    // to 1KB of virtual memory starting at vaddr_base
    vaddr_base = mmap(0, 1024, PROT_READ|PROT_WRITE,
        MAP_SHARED, dev_mem_fd, 0x0);

    if (vaddr_base == MAP_FAILED) { perror("mmap()"); return 1; }
```

```
volatile uint32_t * ema_reg = (uint32_t*) vaddr_base;
uint32_t tmp;

for (int i = 1000; i < 6000; i +=1000){
    printf ("Sending in: %d\n", i);
    *ema_reg = i; //mmio store
    tmp = *ema_reg; //mmio load
    printf("Receiving: %d\n", tmp);
}

if (munmap(vaddr_base, 1024) != 0) { perror("munmap()"); }
if (close(dev_mem_fd) != 0) { perror("close()"); }
dev_mem_fd = -1;

return 0;
}
```

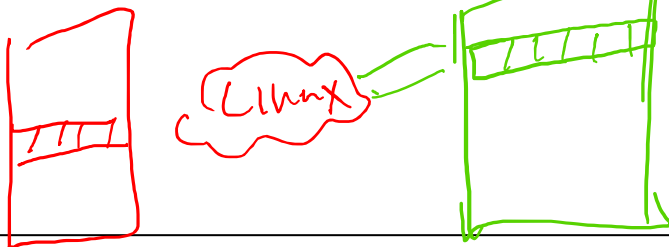
Userspace EMA w/C

```
int main (){

    int dev_mem_fd = -1;
    void * vaddr_base;

    //Mapping user-space I/O
    dev_mem_fd = open("/dev/uio0", O_RDWR|O_SYNC);
    if (dev_mem_fd < 0) { perror("open() /dev/uio0"); return 1; }

    // Map 1KB of physical memory starting at uio 0x0 (real 0x40000000)
    // to 1KB of virtual memory starting at vaddr_base
    vaddr_base = mmap(0, 1024, PROT_READ|PROT_WRITE,
    MAP_SHARED, dev_mem_fd, 0x0);
    40000000 w/ different name
    if (vaddr_base == MAP_FAILED) { perror("mmap()"); return 1; }
```



```
volatile uint32_t * ema_reg = (uint32_t*) vaddr_base;
uint32_t tmp;

for (int i = 1000; i < 6000; i +=1000){
    printf ("Sending in: %d\n", i);
    *ema_reg = i; //mmio store
    tmp = *ema_reg; //mmio load
    printf("Receiving: %d\n", tmp);
}

if (munmap(vaddr_base, 1024) != 0) { perror("munmap()"); }
if (close(dev_mem_fd) != 0) { perror("close()"); }
dev_mem_fd = -1;

return 0;
}
```

Userspace EMA w/C

```
int main (){

    int dev_mem_fd = -1;
    void * vaddr_base;

    //Mapping user-space I/O
    dev_mem_fd = open("/dev/uio0", O_RDWR|O_SYNC);
    if (dev_mem_fd < 0) { perror("open() /dev/uio0"); return 1; }

    // Map 1KB of physical memory starting at uio 0x0 (real 0x40000000)
    // to 1KB of virtual memory starting at vaddr_base
    vaddr_base = mmap(0, 1024, PROT_READ|PROT_WRITE,
        MAP_SHARED, dev_mem_fd, 0x0);

    if (vaddr_base == MAP_FAILED) { perror("mmap()"); return 1; }
```

```
volatile uint32_t * ema_reg = (uint32_t*) vaddr_base;
uint32_t tmp;

for (int i = 1000; i < 6000; i +=1000){
    printf ("Sending in: %d\n", i);
    *ema_reg = i; //mmio store
    tmp = *ema_reg; //mmio load
    printf("Receiving: %d\n", tmp);
}

if (munmap(vaddr_base, 1024) != 0) { perror("munmap()"); }
if (close(dev_mem_fd) != 0) { perror("close()"); }
dev_mem_fd = -1;

return 0;
}
```

Userspace EMA w/C

```
int main (){

    int dev_mem_fd = -1;
    void * vaddr_base;

    //Mapping user-space I/O
    dev_mem_fd = open("/dev/uio0", O_RDWR|O_SYNC);
    if (dev_mem_fd < 0) { perror("open() /dev/uio0"); return 1; }

    // Map 1KB of physical memory starting at uio 0x0 (real 0x40000000)
    // to 1KB of virtual memory starting at vaddr_base
    vaddr_base = mmap(0, 1024, PROT_READ|PROT_WRITE,
        MAP_SHARED, dev_mem_fd, 0x0);

    if (vaddr_base == MAP_FAILED) { perror("mmap()"); return 1; }
```

```
    volatile uint32_t * ema_reg = (uint32_t*) vaddr_base;
    uint32_t tmp;

    for (int i = 1000; i < 6000; i +=1000){
        printf ("Sending in: %d\n", i);
        *ema_reg = i; //mmio store to 0x4000_0000
        tmp = *ema_reg; //mmio load from 0x4000_0000
        printf("Receiving: %d\n", tmp);
    }

    if (munmap(vaddr_base, 1024) != 0) { perror("munmap()"); }
    if (close(dev_mem_fd) != 0) { perror("close()"); }
    dev_mem_fd = -1;

    return 0;
}
```

Complete EMA

```
#include <fcntl.h>
#include <stdlib.h>
#include <stdint.h>
#include <stdio.h>
#include <sys/mman.h>
#include <unistd.h>

int main (){

    int dev_mem_fd = -1;
    void * vaddr_base;

    //Mapping user-space I/O
    dev_mem_fd = open("/dev/uio0", O_RDWR|O_SYNC);
    if (dev_mem_fd < 0) { perror("open() /dev/uio0"); return 1; }

    // Map 1KB of physical memory starting at uio 0x0 (real 0x40000000)
    // to 1KB of virtual memory starting at vaddr_base
    vaddr_base = mmap(0, 1024, PROT_READ|PROT_WRITE,
                     MAP_SHARED, dev_mem_fd, 0x0);

    if (vaddr_base == MAP_FAILED) { perror("mmap()"); return 1; }
```

```
volatile uint32_t * ema_reg = (uint32_t*) vaddr_base;
uint32_t tmp;

for (int i = 1000; i < 6000; i +=1000){
    printf ("Sending in: %d\n", i);
    *ema_reg = i; //mmio store
    tmp = *ema_reg; //mmio load
    printf("Receiving: %d\n", tmp);
}

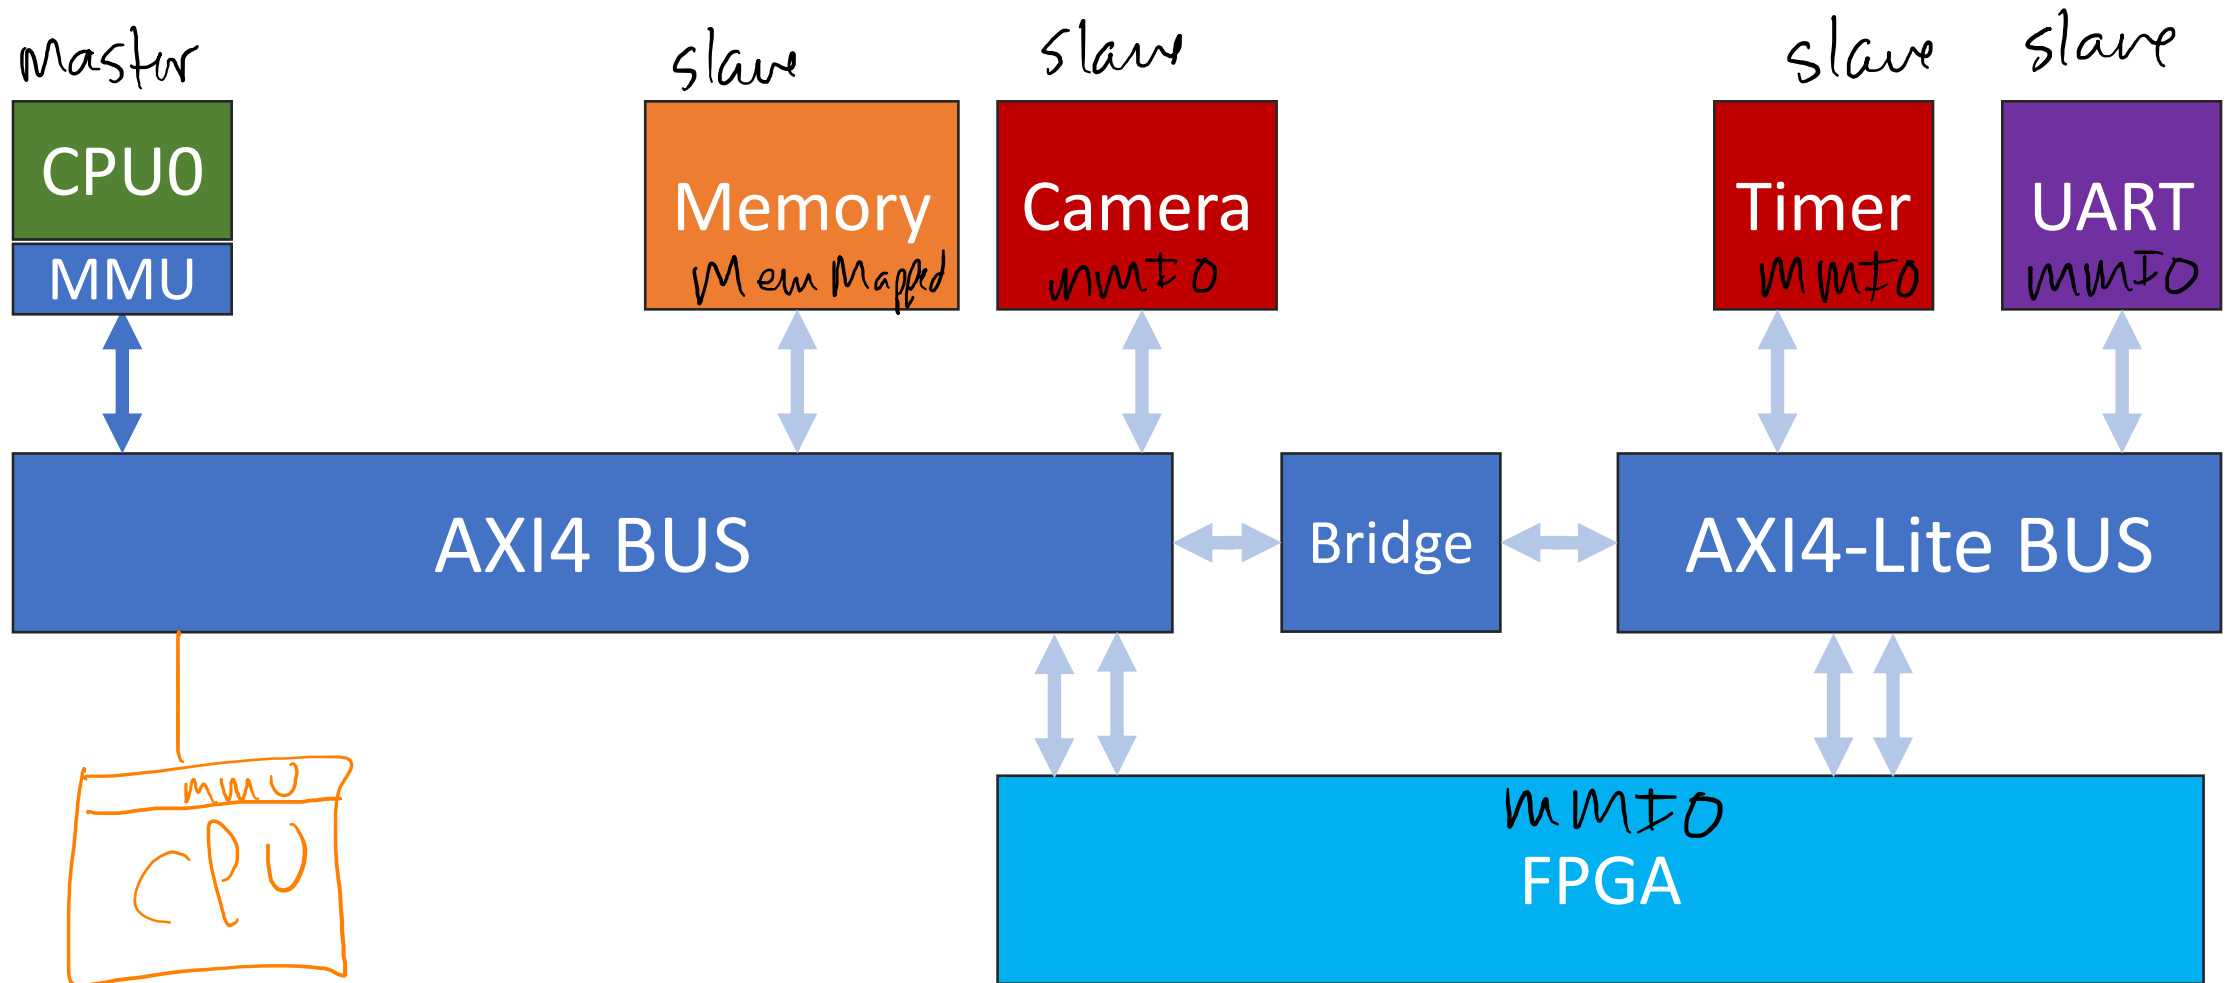
if (munmap(vaddr_base, 1024) != 0) { perror("munmap()"); }
if (close(dev_mem_fd) != 0) { perror("close()"); }
dev_mem_fd = -1;

return 0;
}
```

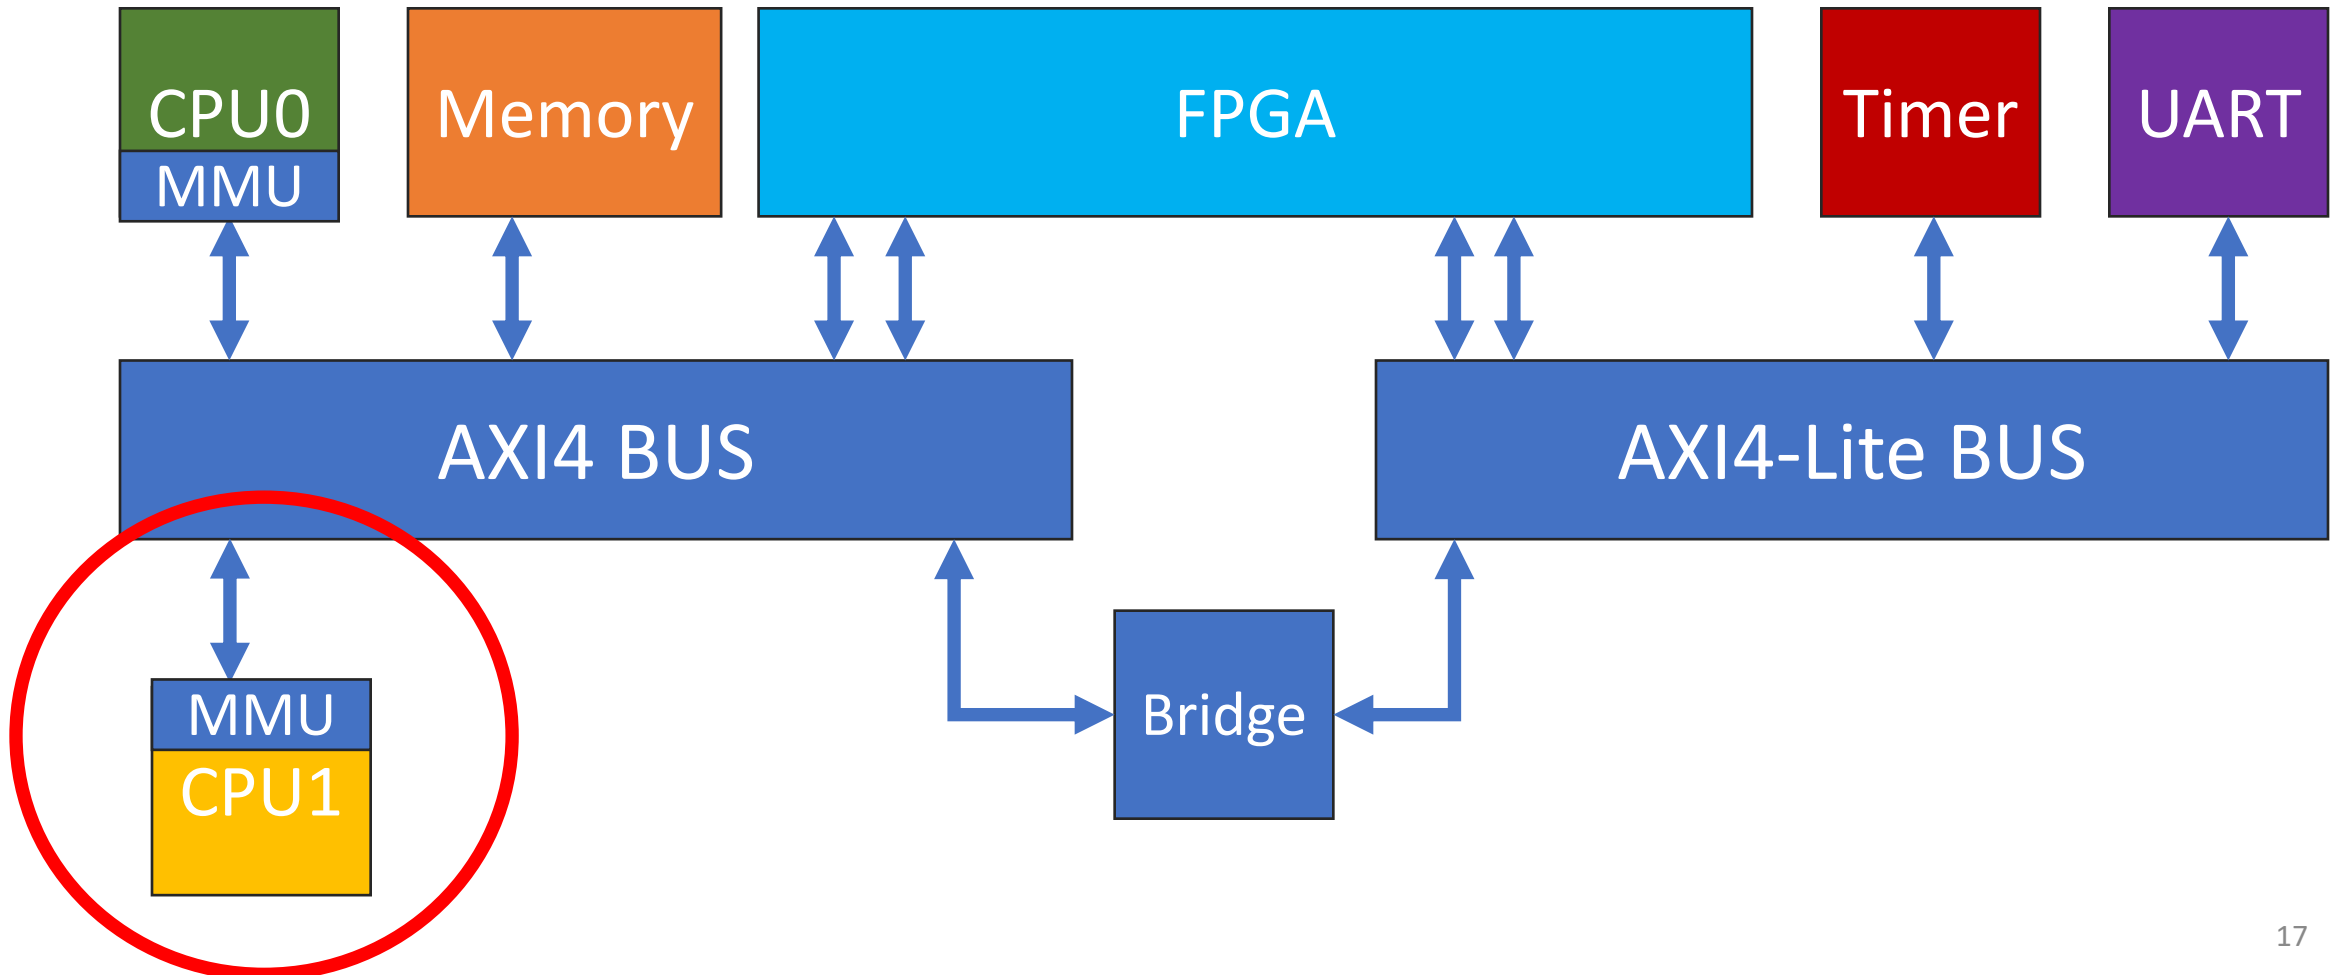
Demo Time

master: issue request
Slave: fulfil request

Which are AXI4 "Masters"?

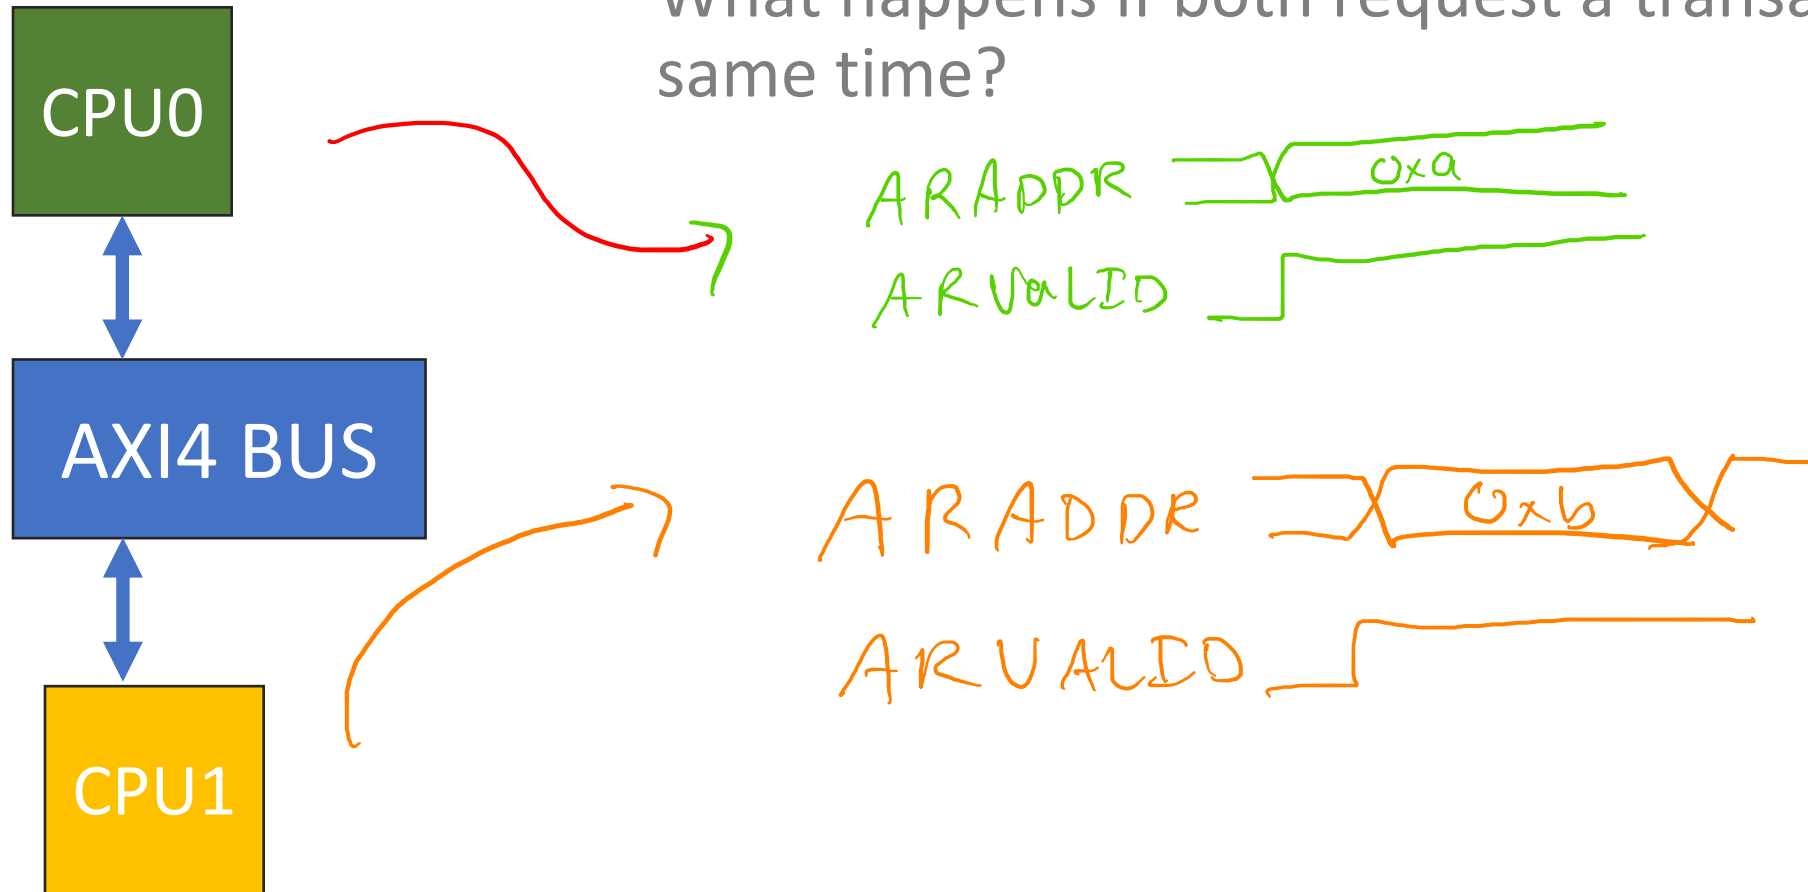


Multiple Masters

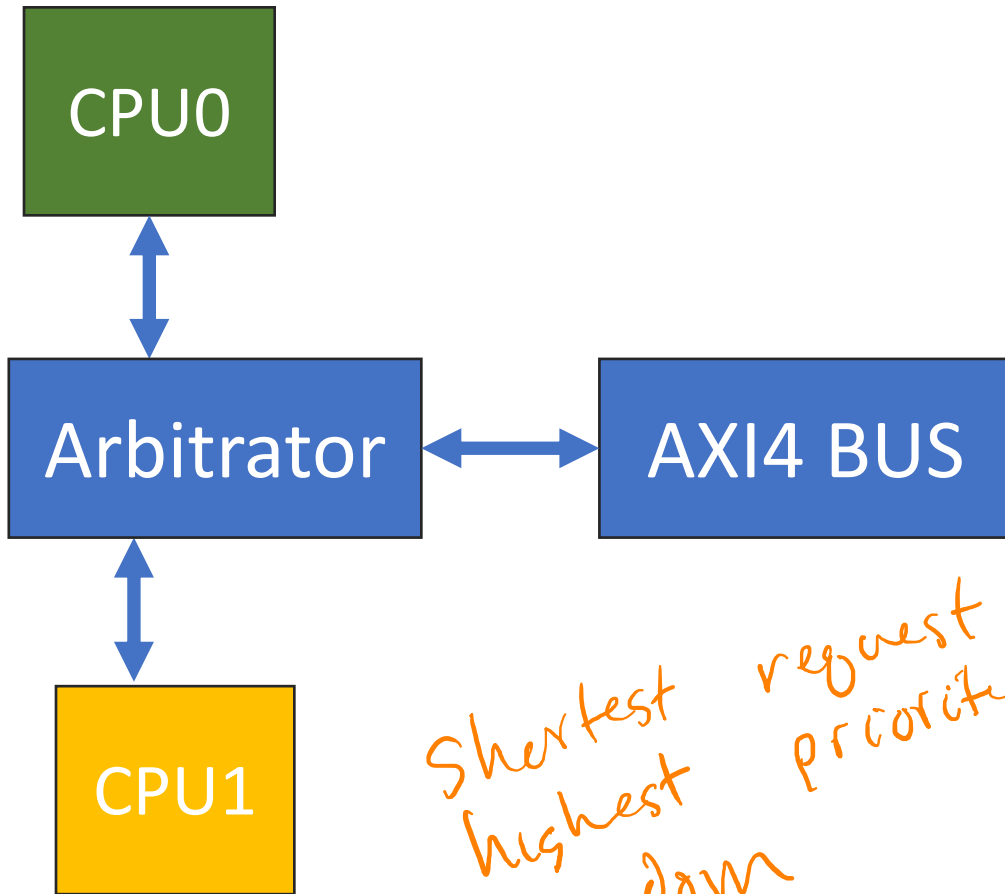


Multi-Master Buses

- What happens if both request a transaction at the same time?



An **Arbitrator** selects who gets to use the bus



- What happens if both request a transaction at the same time?

- **Arbitration:** Pick a winner!

- What Arbitration scheme to use?

*Shortest
highest
random*


*request
priority*

round-robin

Arbitration Options

- Highest-Priority First
- Round Robin

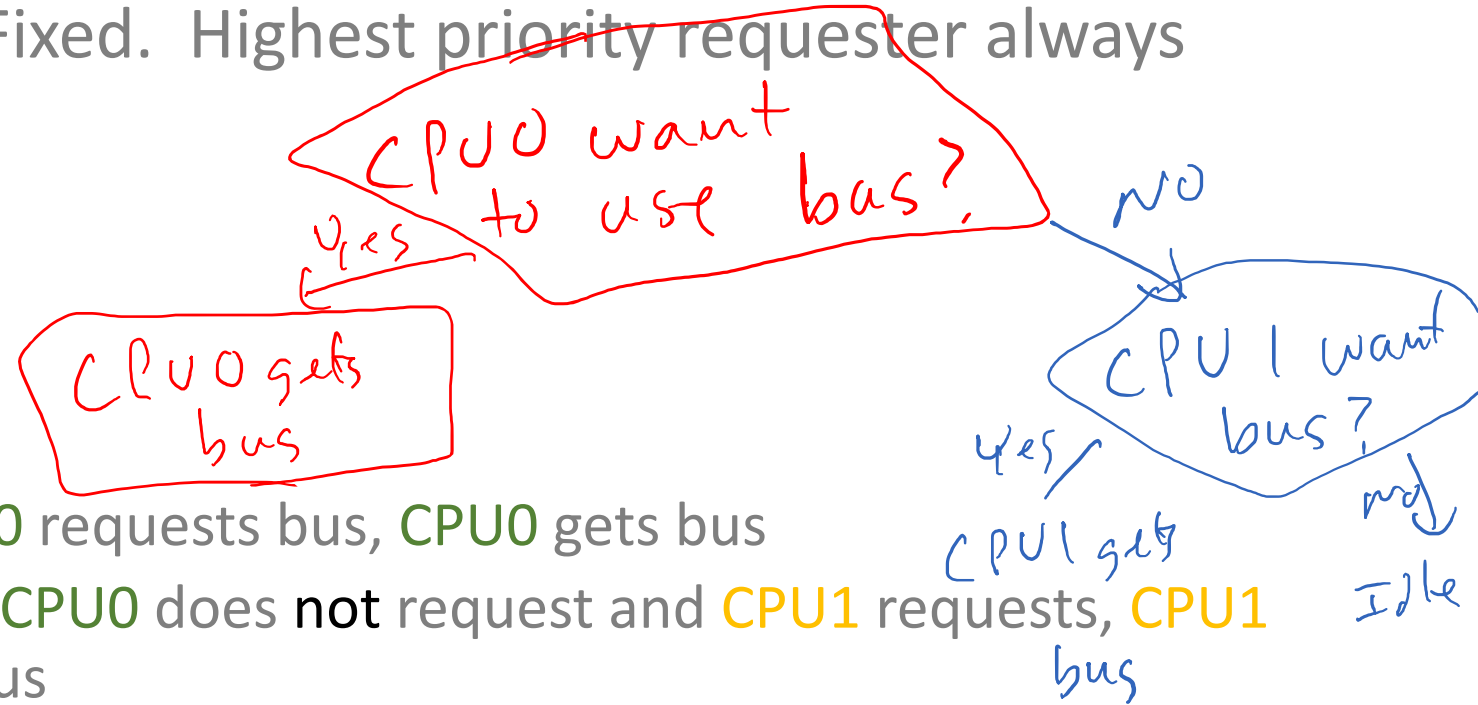
Highest-Priority First



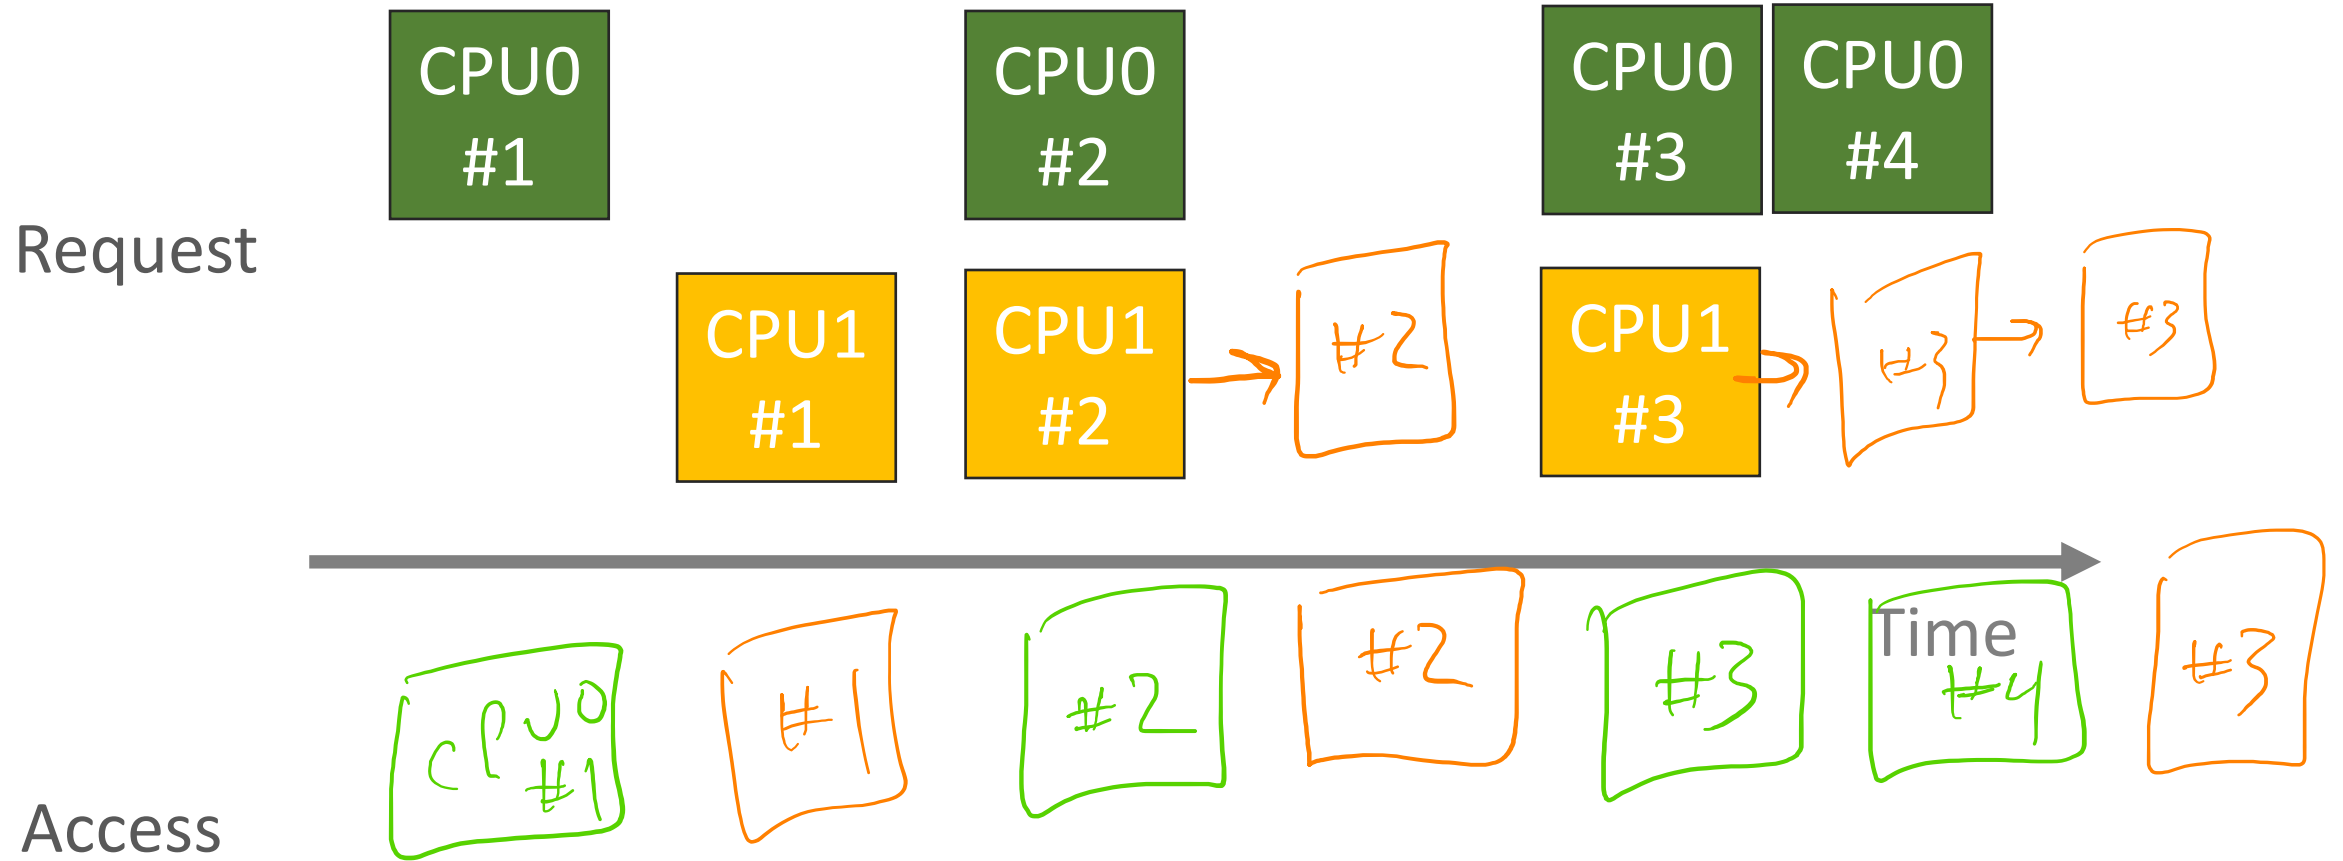

- Priority Fixed. Highest priority requester always wins.

- Rules:

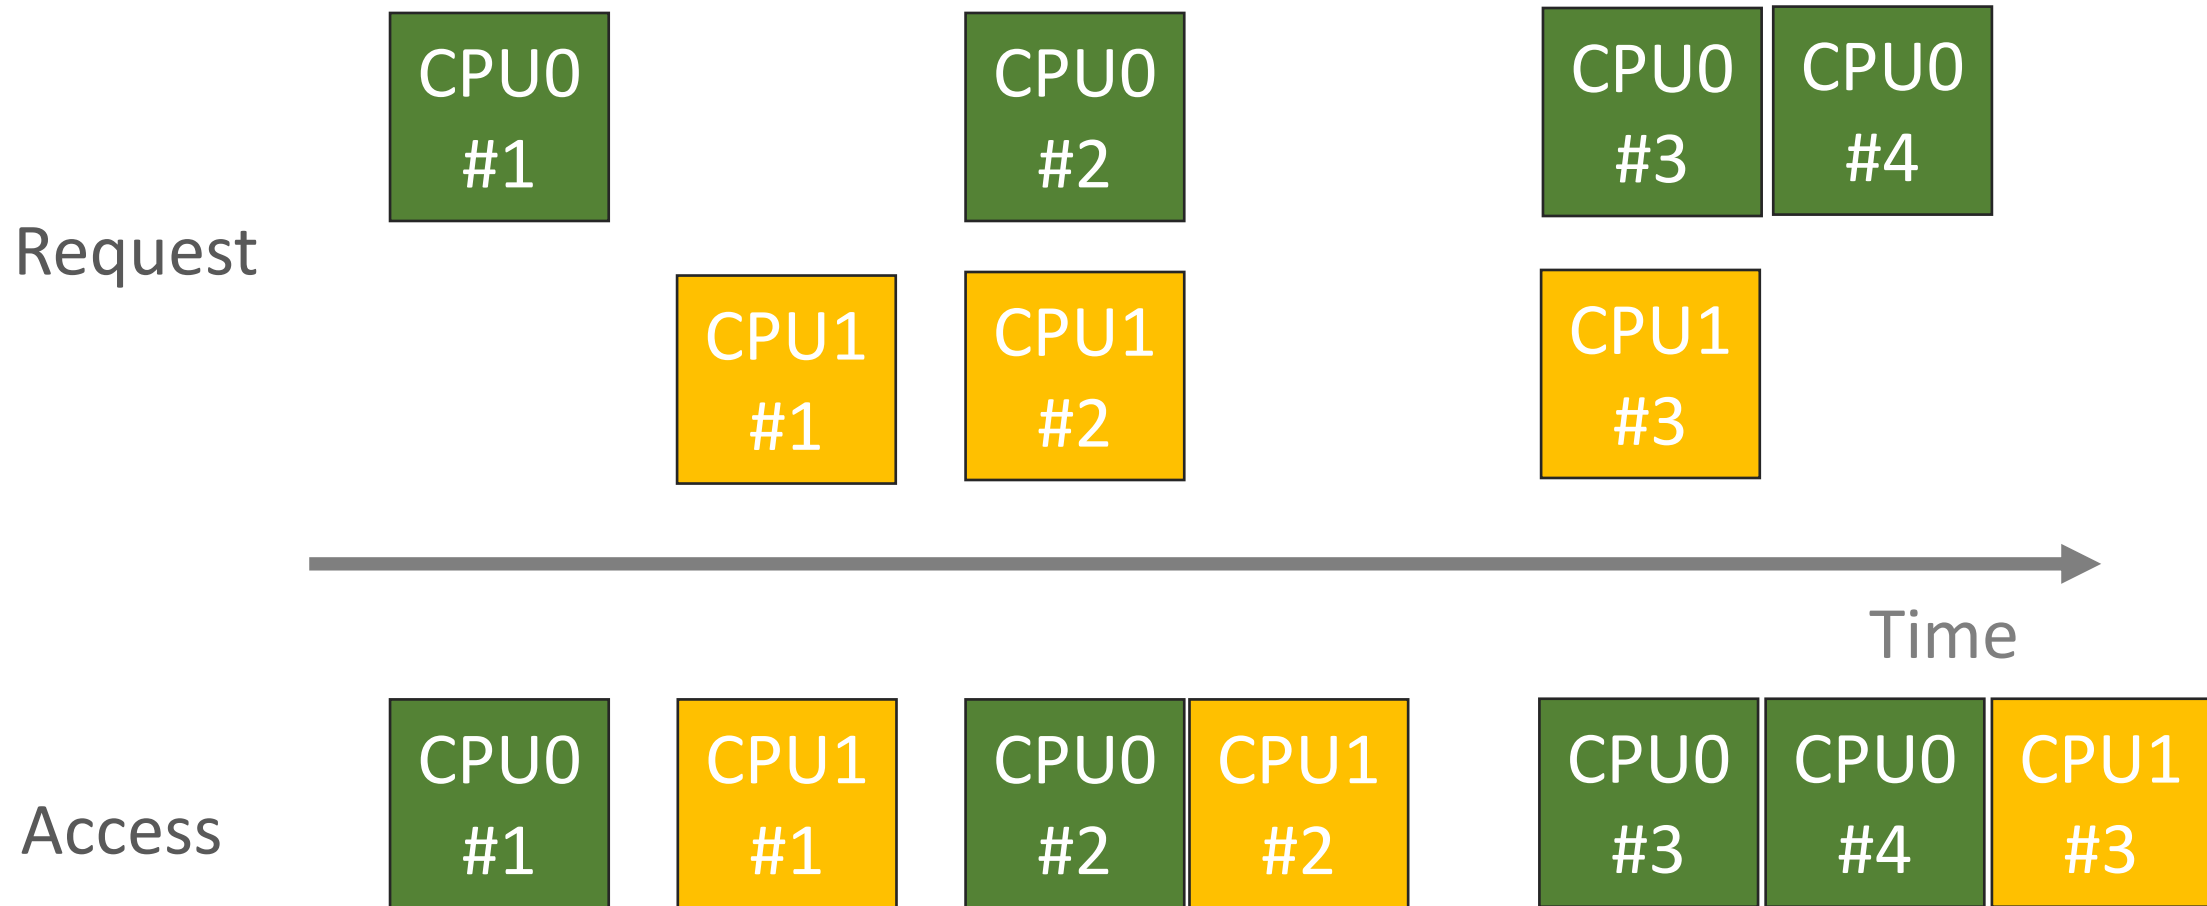

- If CPU0 requests bus, CPU0 gets bus
- Else if CPU0 does not request and CPU1 requests, CPU1 gets bus
- Else bus idle



Highest Priority First



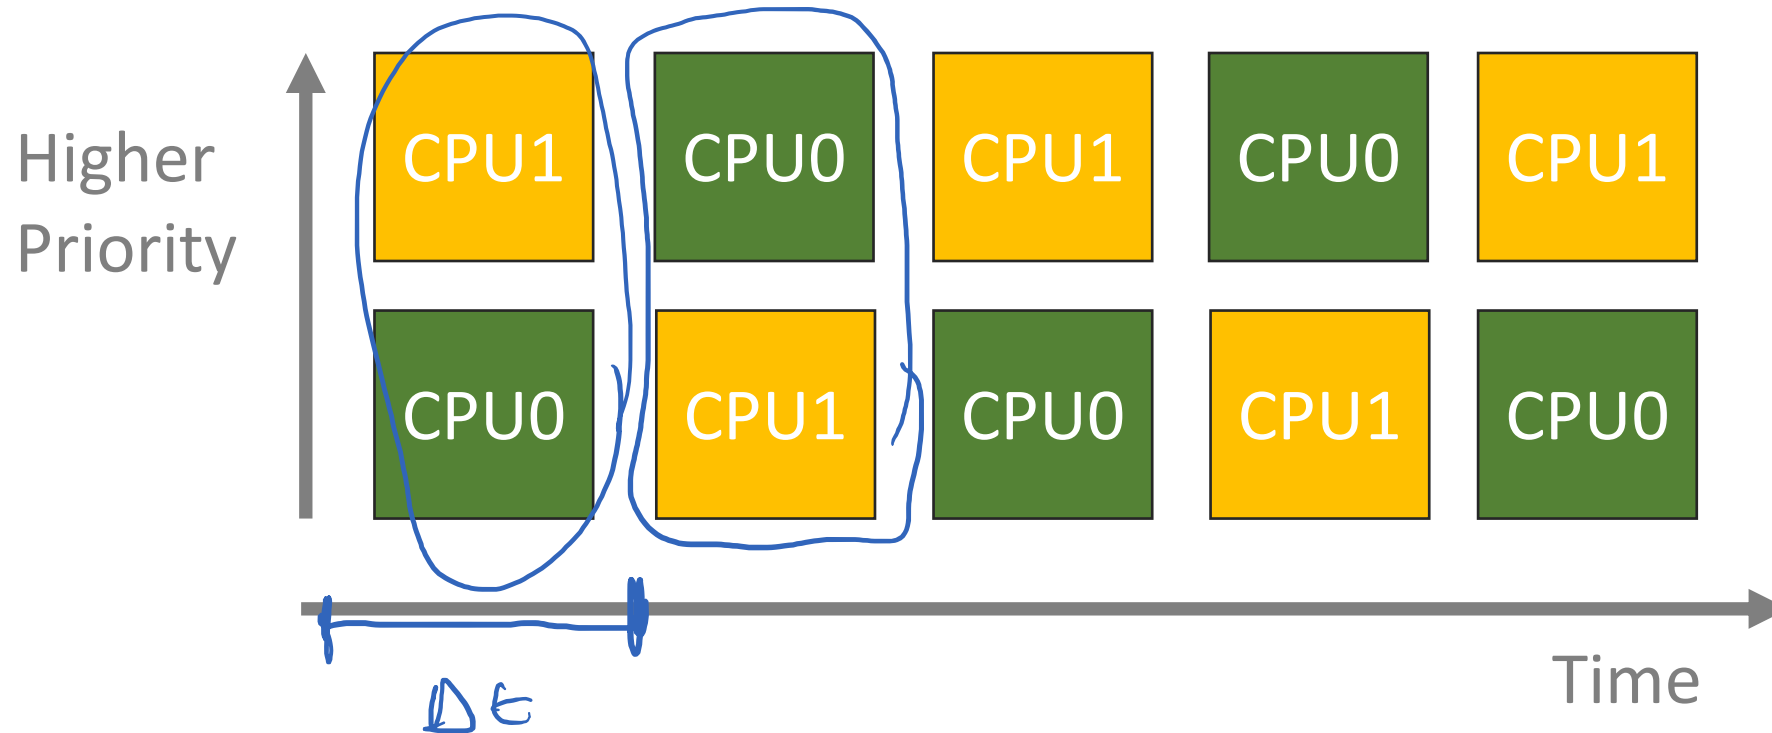
Highest Priority First



Arbitrator Circuit?


Round Robin

- Priority updates every cycle. Everyone get's equal access to highest priority



Round Robin

Higher
Priority



Request




Time


Access

Round Robin


Higher
Priority ↑



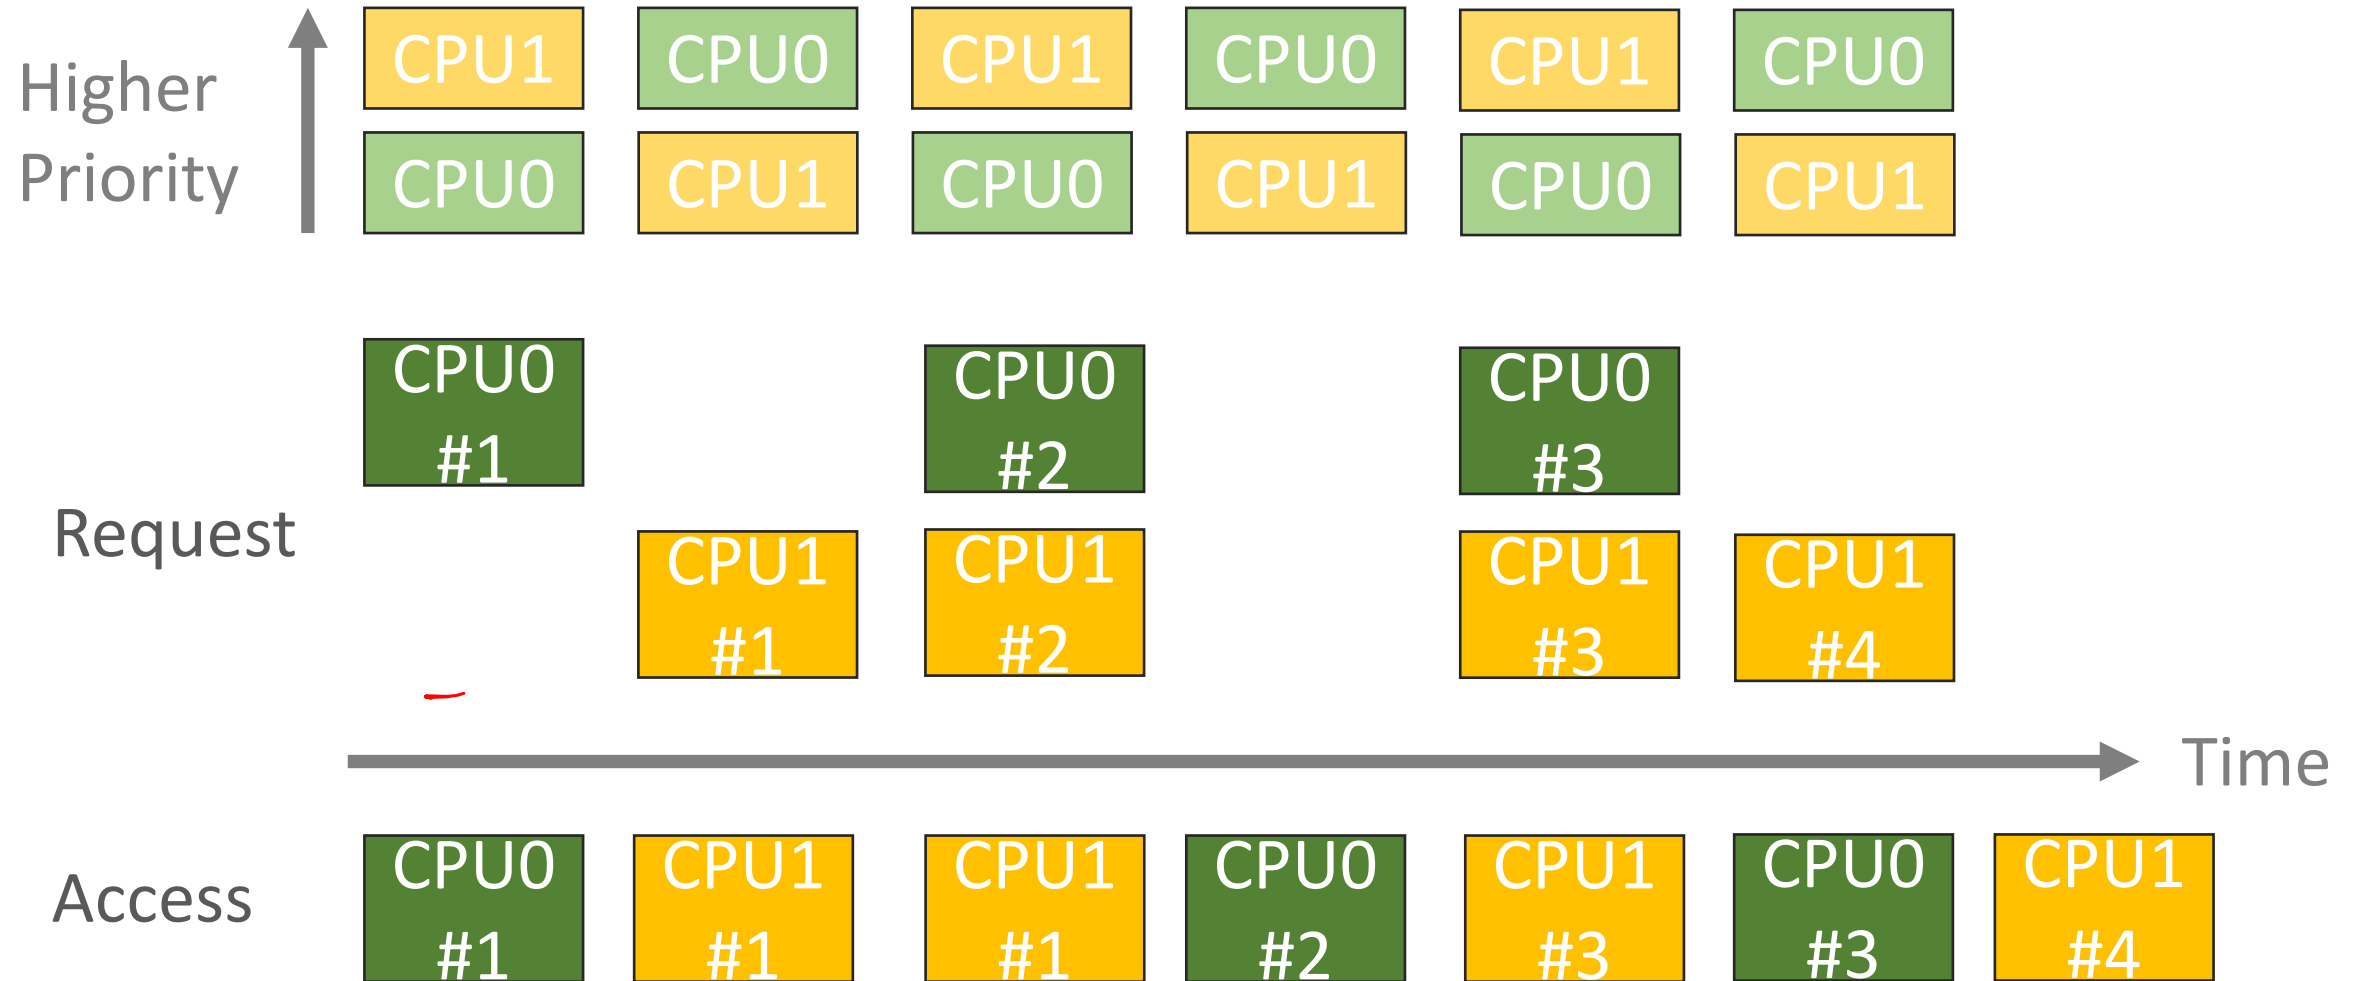
Request



Access

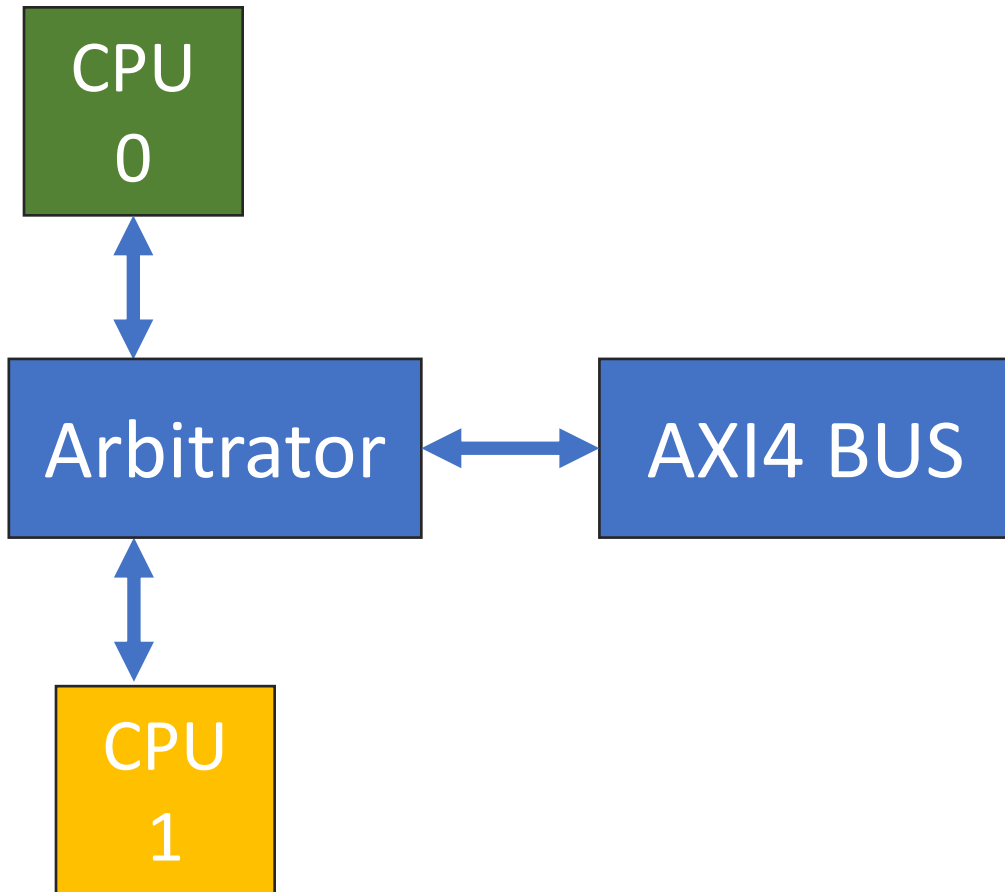


Round Robin



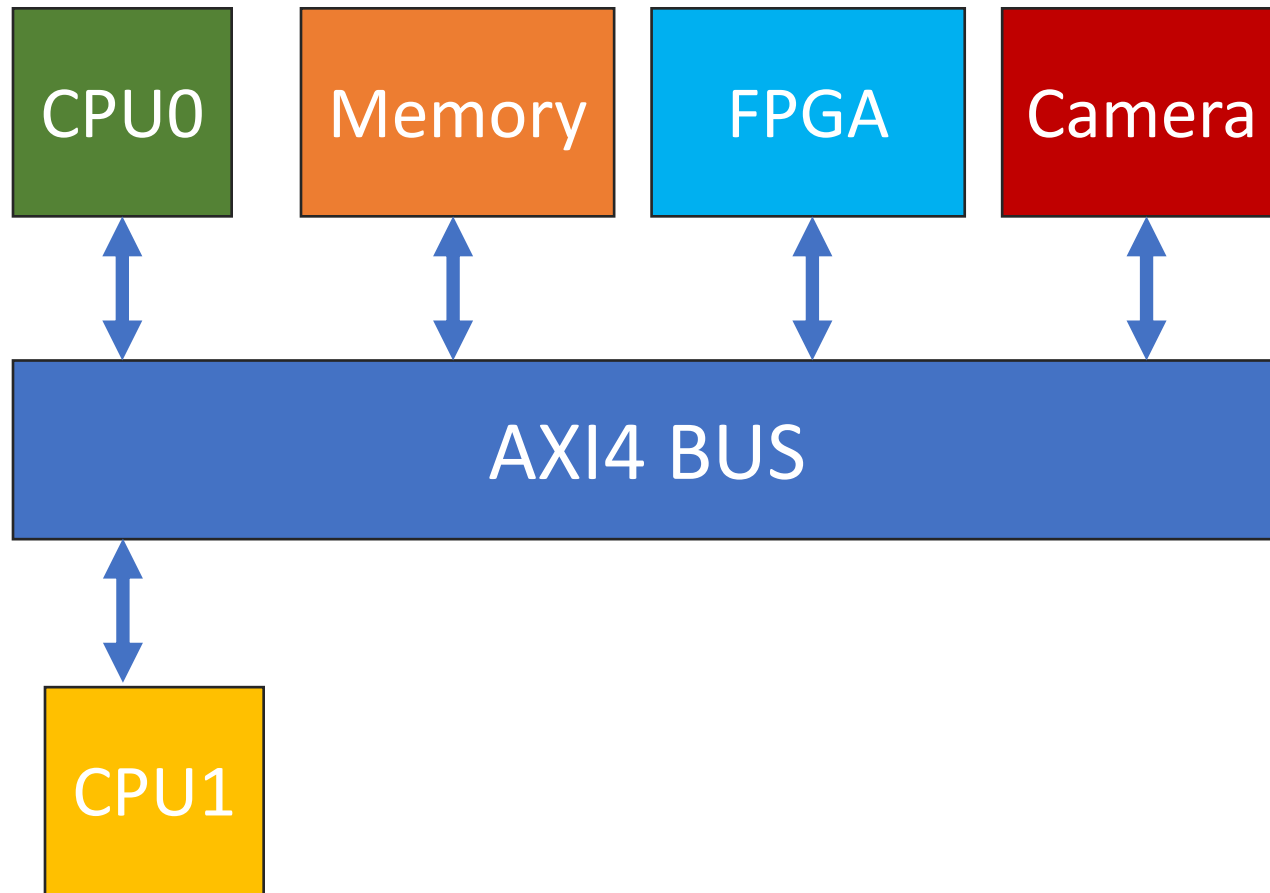
Other Arbitration ideas?

An **Arbitrator** selects who gets to use the bus



- What happens if both request a transaction at the same time?
- **Arbitration:**
 - Fixed-Priority
 - Round Robin
 - Many more...

Q: How do I move data between the Camera and Memory?



A: The CPU copies data from Camera to Memory

```
#define CAMERA_MMIO_ADDR 0x40000004
volatile uint32_t * camera =
    (uint32_t *) (CAMERA_MMIO_ADDR);
#define BUF_SIZE 1024;
uint32_t buf[BUF_SIZE];

int main () {
    //...
    while (true){
        copy_image(camera, buf, BUF_SIZE);
        detect_face(buf);
    }
}
```

```
void copy_image (uint32_t * from,
                 uint32_t * to,
                 uint32_t size)
{
    register uint32_t reg;

    for (int i = 0; i < size; ++i){
        reg = *from;
        to[i] = reg;
    }
}
```

A: The CPU copies data from Camera to Memory


```
#define CAMERA_MMIO_ADDR 0x40000004
volatile uint32_t * camera =
    (uint32_t *) (CAMERA_MMIO_ADDR);
#define BUF_SIZE 1024;
uint32_t buf[BUF_SIZE];

int main () {

    while (true){
        copy_image(camera, buf, BUF_SIZE);
        detect_face(buf);
    }
}
```

```
void copy_image (uint32_t * from,
                uint32_t * to,
                uint32_t size)
{
    register uint32_t reg;

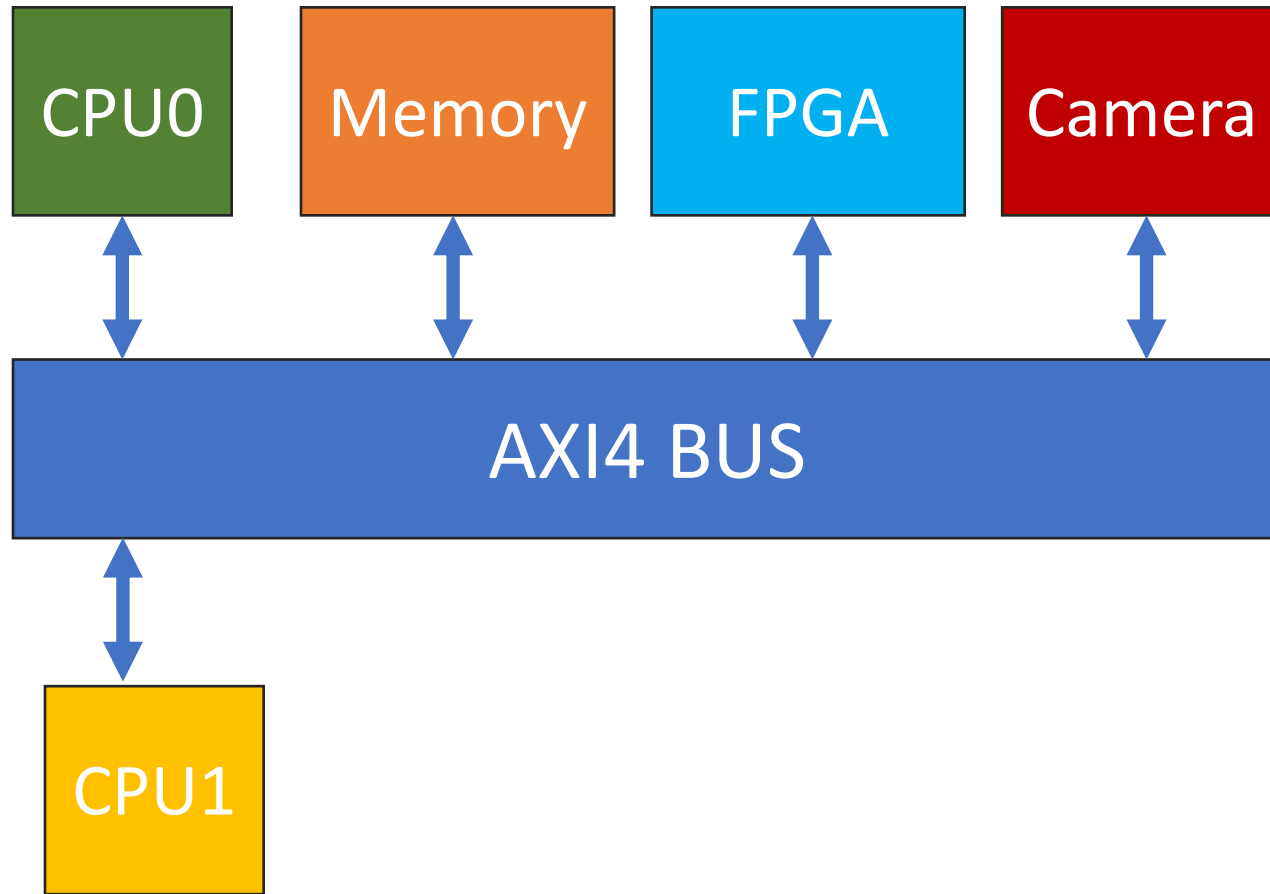
    for (int i = 0; i < size; ++i){
        ① reg = *from;
        ② to[i] = reg;
        ④
    }
}
```



① load
② store

③ increment
④ branch

What bus transactions result?



What else can the CPU do
while copying data?

CPU \Rightarrow 1 GHz \rightarrow 1 Billion cycle / second
 \rightarrow 1 Billion instructions / second

What else can the CPU do while copying data?

- CPU can do 1B instructions/second. (1GHz)
- 4 Instructions per loop
 - 1 load, 1 store, 1 increment, 1 branch
- 250M copies/second

4K Video: 1697 Mbps* = 212 MB / second

~85% CPU utilization for Copy!

What about Ethernet?

- CPU can do 1B instructions/second. (1GHz)
- 4 Instructions per loop
 - 1 load, 1 store, 1 increment, 1 branch
- 250M copies/second
- 1Gbps Ethernet:
- 1 Gbps Receive + 1Gbps Transmit = 2 Gbps
- 2Gbps = 250MB/second
- **Nothing. ~100% of CPU required?**

What if we do the copy on CPU1?

```
int main () {  
  
    while (true){  
  
        ask_cpu1_to_copy_image(camera, buf, BUF_SIZE);  
  
        detect_face(buf);  
    }  
}
```

What if we do the copy on CPU1?

cpu0

```
int main () {  
    buf1, buf2  
    while (true){  
        → ask_cpu1_to_copy_image(camera, buf1, BUF_SIZE); ←  
        wait()  
        detect_face(buf2);  
    }  
}
```

switch

buf2

buf1

CPU0

—

empt
Buf0

empty
Buf1

buf0

CPU1

full
Buf0

Buf1

full
Buf0

buf1

↑
overlap
(parallelism)

time

double
buffering

Copy on CPU1, Version 2.

```
int main () {
```


```
    ➔ ask_cpu1_to_copy_image(camera, buf1, BUF_SIZE);  
    ➔ wait_for_cpu1_done();
```

```
    while (true){  
        ask_cpu1_to_copy_image(camera, buf2, BUF_SIZE);  
        detect_face(buf1);  
        ➔ wait_for_cpu1_done();  
  
        ask_cpu1_to_copy_image(camera, buf1, BUF_SIZE);  
        detect_face(buf2);  
        wait_for_cpu1_done();  
    }
```

```
}
```


CPU0

CPU1



Why are we wasting an entire CPU for this?

```
void copy_image (uint32_t * from,
                 uint32_t * to,
                 uint32_t size)
{
    register uint32_t reg;

    for (int i = 0; i < size; ++i){

        reg = *from;

        to[i] = reg;

    }
}
```

DMA: Direct Memory Access

- A mini-CPU that does copy for you:

```
void copy (uint32_t * from,
           uint32_t * to,
           uint32_t size)
{
    register uint32_t reg;

    for (int i = 0; i < size; ++i){

        reg = *from;

        to[i] = reg;
    }
}
```

DMA

```
int main () {  
  
    dma_copy_image(camera, buf1, BUF_SIZE);  
    wait_for_dma_done();  
  
    while (true){  
        dma_copy_image(camera, buf2, BUF_SIZE);  
        detect_face(buf1);  
        wait_for_dma_done();  
  
        dma_copy_image(camera, buf1, BUF_SIZE);  
        detect_face(buf2);  
        wait_for_dma_done();  
    }  
}
```

Next Time:

- Direct Memory Access

Virt Linux → Python
Windows Vivado

References

- Zynq Book, Chapter 19 “AXI Interfacing”
- [Practical Introduction to Hardware/Software Codesign](#)
 - Chapter 10
- AMBA AXI Protocol v1.0
 - http://mazsola.iit.uni-miskolc.hu/~drdani/docs_arm/AMBAaxi.pdf
- <https://lauri.võsandi.com/hdl/zynq/axi-stream.html>

11: Multi-Master Buses

Engr 315: Hardware / Software Codesign
Andrew Lukefahr
Indiana University

