

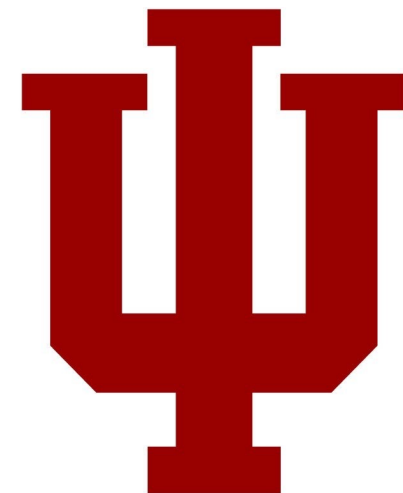
PS → } pushed to Monday
{ Demo Wed

13: DMA

Engr 315: Hardware / Software Codesign

Andrew Lukefahr

Indiana University



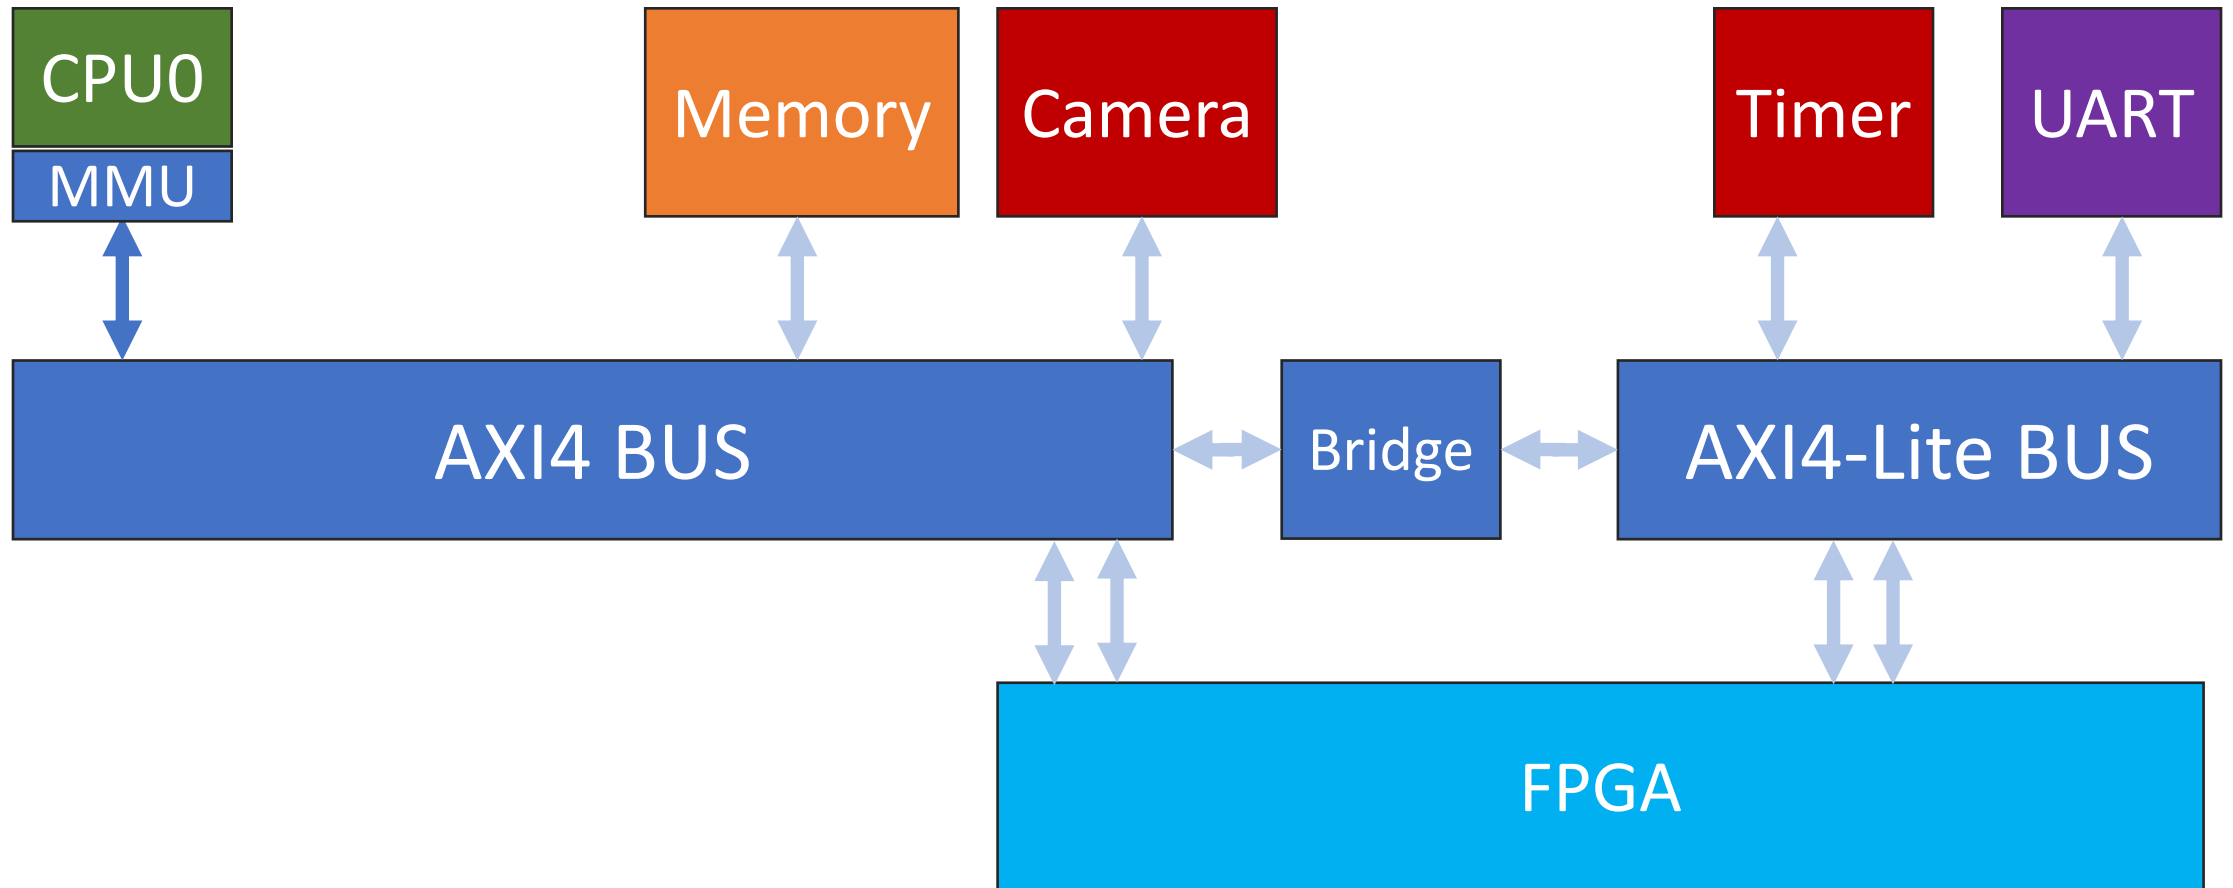
Announcements

- P5 is out! *↗ extended*

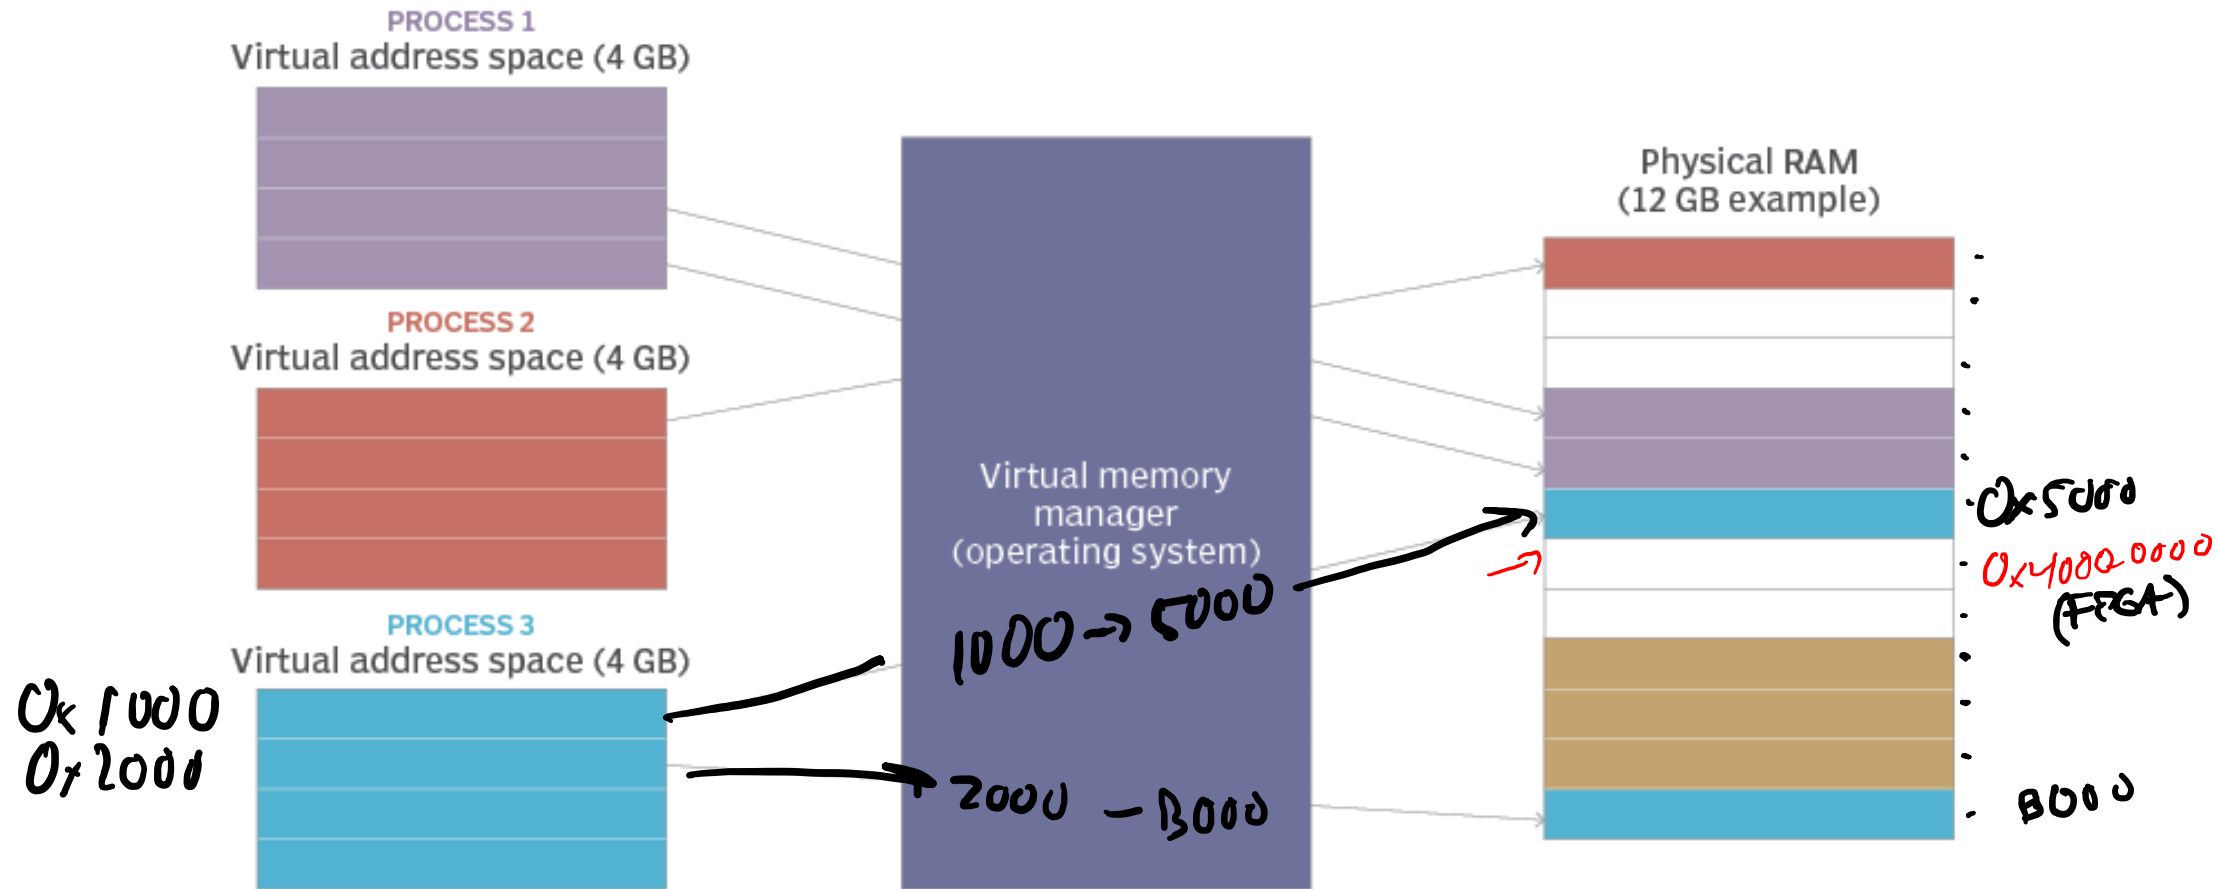
Exam Planning

11/01	Wednesday	18	Review	–
11/06	Monday	19	Exam	
11/08	Wednesday	20	Review	

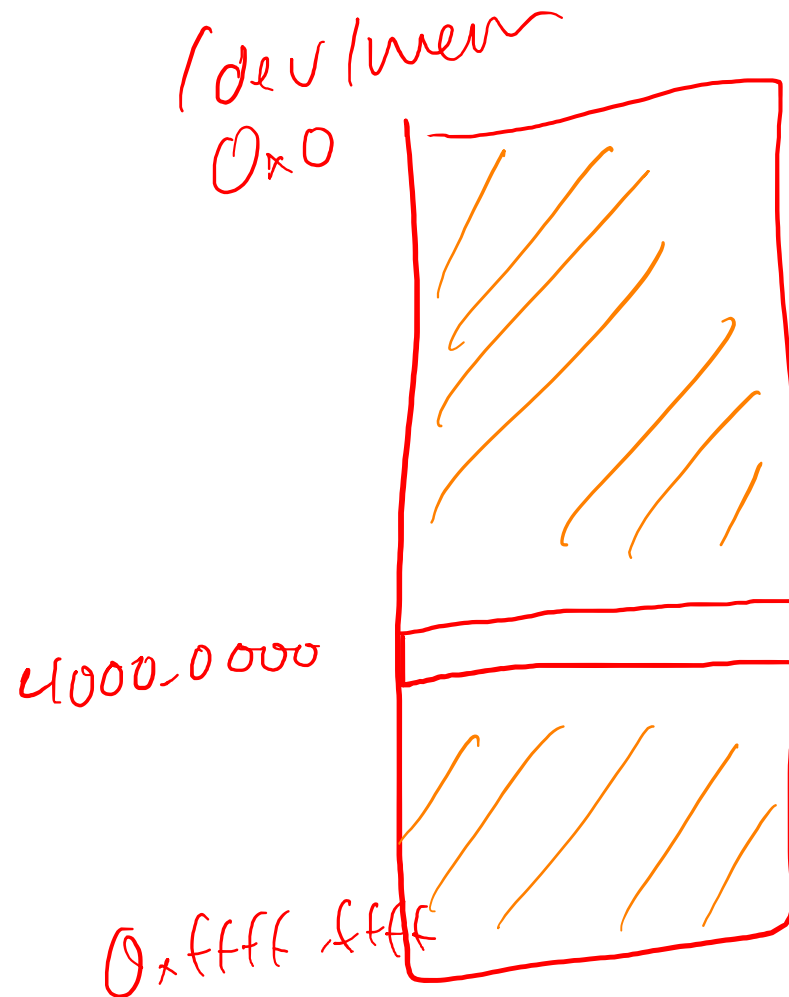
Machine Model, V3: MMUs



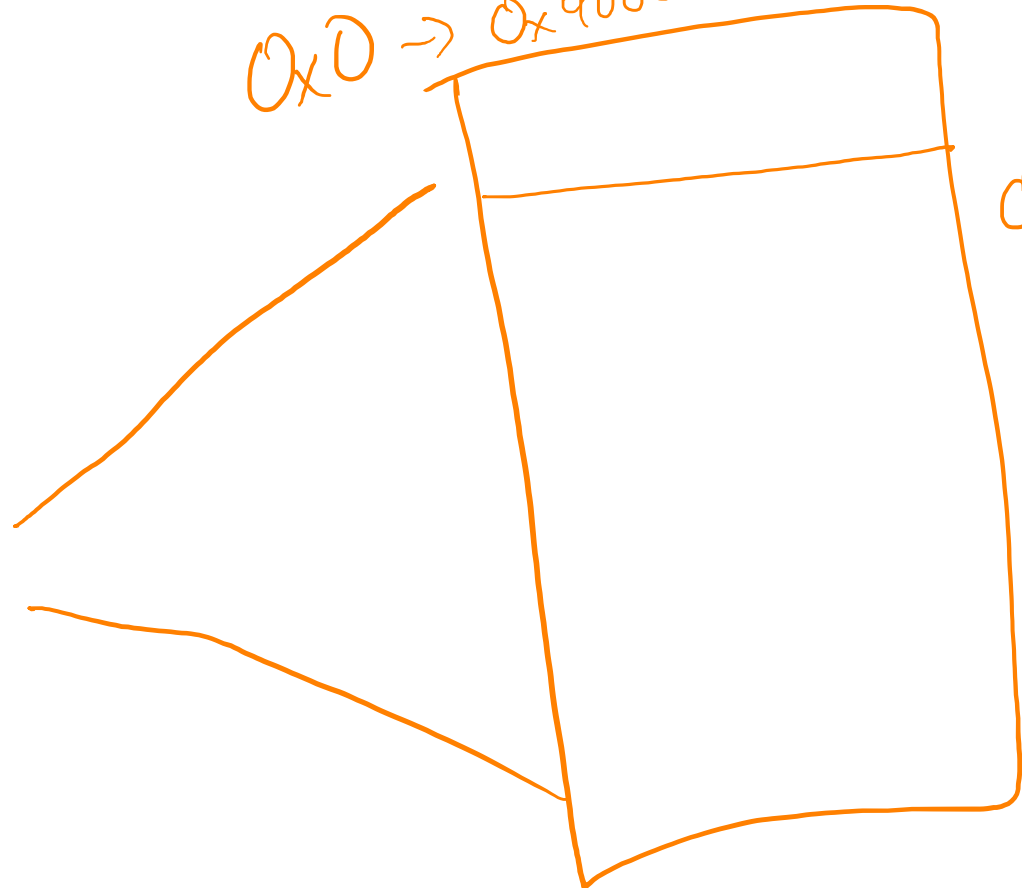
OS (Linux) mains full Virtual->Physical Mappings



/dev/urandom



/dev/urandom
0x0 → 0x4000-0000



0x4 →
~~0x4000-0000~~
0x4000-0004

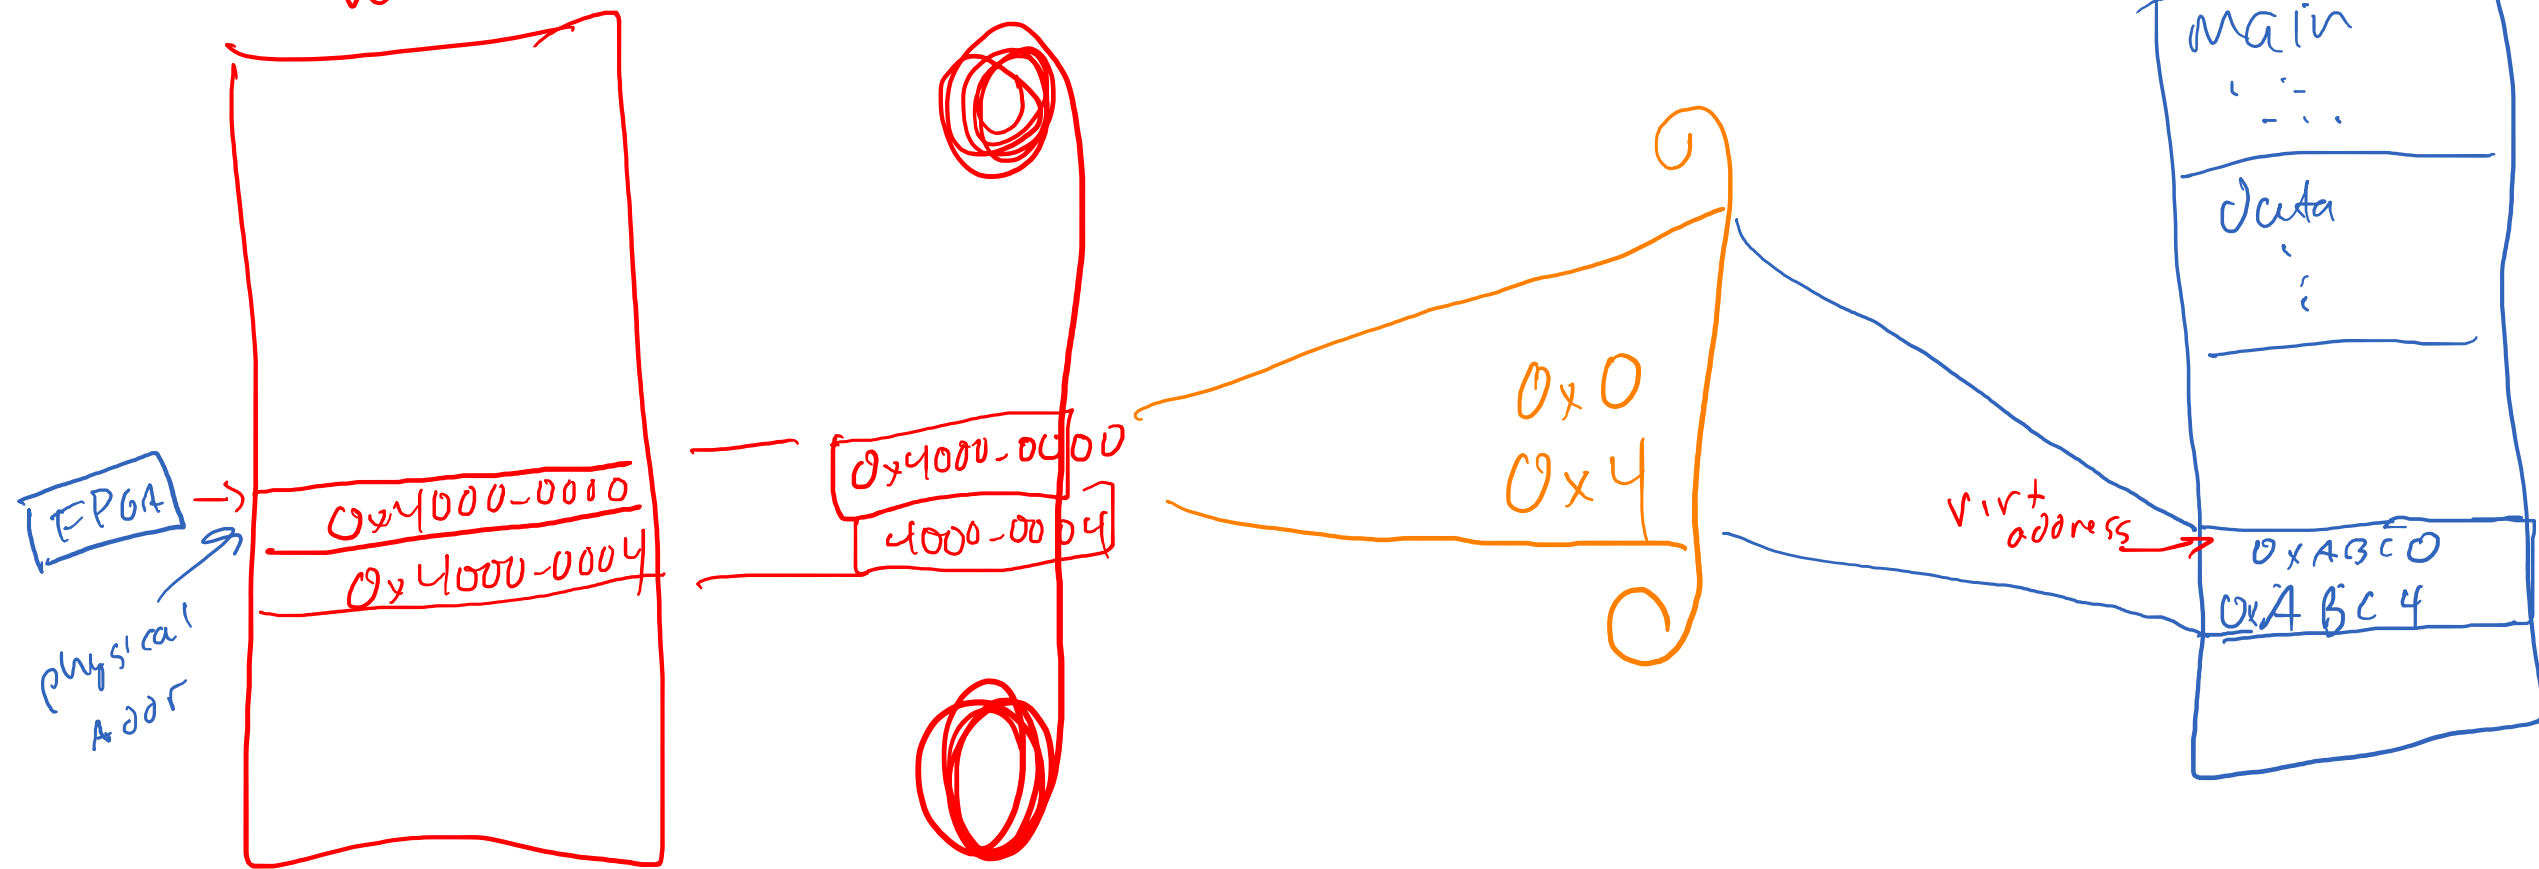
mmap

physical
mem

Linux
/dev/mem

Linux
/dev/urandom

Linux
mmap



Userspace EMA w/C

```
int main (){

    int dev_mem_fd = -1;
    void * vaddr_base;

    //Mapping user-space I/O
    dev_mem_fd = open("/dev/uio0", O_RDWR|O_SYNC);
    if (dev_mem_fd < 0) { perror("open() /dev/uio0"); return 1; }

    // Map 1KB of physical memory starting at uio 0x0 (real 0x40000000)
    // to 1KB of virtual memory starting at vaddr_base
    vaddr_base = mmap(0, 1024, PROT_READ|PROT_WRITE,
        MAP_SHARED, dev_mem_fd, 0x0);

    if (vaddr_base == MAP_FAILED) { perror("mmap()"); return 1; }
```

```
volatile uint32_t * ema_reg = (uint32_t*) vaddr_base;
uint32_t tmp;

for (int i = 1000; i < 6000; i +=1000){
    printf ("Sending in: %d\n", i);
    *ema_reg = i; //mmio store
    tmp = *ema_reg; //mmio load
    printf("Receiving: %d\n", tmp);
}

if (munmap(vaddr_base, 1024) != 0) { perror("munmap()"); }
if (close(dev_mem_fd) != 0) { perror("close()"); }
dev_mem_fd = -1;

return 0;
}
```


Userspace EMA w/C

```
int main (){

    int dev_mem_fd = -1;
    void * vaddr_base;

    //Mapping user-space I/O
    dev_mem_fd = open("/dev/uio0", O_RDWR|O_SYNC);
    if (dev_mem_fd < 0) { perror("open() /dev/uio0"); return 1; }

    // Map 1KB of physical memory starting at uio 0x0 (real 0x40000000)
    // to 1KB of virtual memory starting at vaddr_base
    vaddr_base = mmap(0, 1024, PROT_READ|PROT_WRITE,
        MAP_SHARED, dev_mem_fd, 0x0);

    if (vaddr_base == MAP_FAILED) { perror("mmap()"); return 1; }
```

```
volatile uint32_t * ema_reg = (uint32_t*) vaddr_base;
uint32_t tmp;

for (int i = 1000; i < 6000; i +=1000){
    printf ("Sending in: %d\n", i);
    *ema_reg = i; //mmio store
    tmp = *ema_reg; //mmio load
    printf("Receiving: %d\n", tmp);
}

if (munmap(vaddr_base, 1024) != 0) { perror("munmap()"); }
if (close(dev_mem_fd) != 0) { perror("close()"); }
dev_mem_fd = -1;

return 0;
}
```

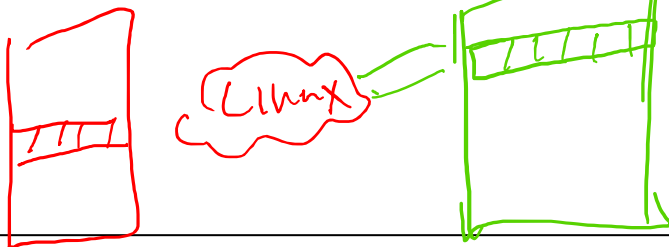
Userspace EMA w/C

```
int main (){

    int dev_mem_fd = -1;
    void * vaddr_base;

    //Mapping user-space I/O
    dev_mem_fd = open("/dev/uio0", O_RDWR|O_SYNC);
    if (dev_mem_fd < 0) { perror("open() /dev/uio0"); return 1; }

    // Map 1KB of physical memory starting at uio 0x0 (real 0x40000000)
    // to 1KB of virtual memory starting at vaddr_base
    vaddr_base = mmap(0, 1024, PROT_READ|PROT_WRITE,
    MAP_SHARED, dev_mem_fd, 0x0);
    40000000 w/ different name
    if (vaddr_base == MAP_FAILED) { perror("mmap()"); return 1; }
```



```
volatile uint32_t * ema_reg = (uint32_t*) vaddr_base;
uint32_t tmp;

for (int i = 1000; i < 6000; i +=1000){
    printf ("Sending in: %d\n", i);
    *ema_reg = i; //mmio store
    tmp = *ema_reg; //mmio load
    printf("Receiving: %d\n", tmp);
}

if (munmap(vaddr_base, 1024) != 0) { perror("munmap()"); }
if (close(dev_mem_fd) != 0) { perror("close()"); }
dev_mem_fd = -1;

return 0;
}
```

Userspace EMA w/C

virtual address base

```
int main (){

    int dev_mem_fd = -1;
    void * vaddr_base;

    //Mapping user-space I/O
    dev_mem_fd = open("/dev/uio0", O_RDWR|O_SYNC);
    if (dev_mem_fd < 0) { perror("open() /dev/uio0"); return 1; }

    // Map 1KB of physical memory starting at uio 0x0 (real 0x40000000)
    // to 1KB of virtual memory starting at vaddr_base
    vaddr_base = mmap(0, 1024, PROT_READ|PROT_WRITE,
        MAP_SHARED, dev_mem_fd, 0x0);

    if (vaddr_base == MAP_FAILED) { perror("mmap()"); return 1; }
```

```
volatile uint32_t * ema_reg = (uint32_t*) vaddr_base;
uint32_t tmp;

for (int i = 1000; i < 6000; i +=1000){
    printf ("Sending in: %d\n", i);
    *ema_reg = i; //mmio store
    tmp = *ema_reg; //mmio load
    printf("Receiving: %d\n", tmp);
}

if (munmap(vaddr_base, 1024) != 0) { perror("munmap()"); }
if (close(dev_mem_fd) != 0) { perror("close()"); }
dev_mem_fd = -1;

return 0;
}
```

Userspace EMA w/C

```
int main (){

    int dev_mem_fd = -1;
    void * vaddr_base;

    //Mapping user-space I/O
    dev_mem_fd = open("/dev/uio0", O_RDWR|O_SYNC);
    if (dev_mem_fd < 0) { perror("open() /dev/uio0"); return 1; }

    // Map 1KB of physical memory starting at uio 0x0 (real 0x40000000)
    // to 1KB of virtual memory starting at vaddr_base
    vaddr_base = mmap(0, 1024, PROT_READ|PROT_WRITE,
        MAP_SHARED, dev_mem_fd, 0x0);

    if (vaddr_base == MAP_FAILED) { perror("mmap()"); return 1; }
```

```
    volatile uint32_t * ema_reg = (uint32_t*) vaddr_base;
    uint32_t tmp;

    for (int i = 1000; i < 6000; i +=1000){
        printf ("Sending in: %d\n", i);
        *ema_reg = i; //mmio store to 0x4000_0000
        tmp = *ema_reg; //mmio load from 0x4000_0000
        printf("Receiving: %d\n", tmp);
    }

    if (munmap(vaddr_base, 1024) != 0) { perror("munmap()"); }
    if (close(dev_mem_fd) != 0) { perror("close()"); }
    dev_mem_fd = -1;

    return 0;
}
```

Complete EMA

```
#include <fcntl.h>
#include <stdlib.h>
#include <stdint.h>
#include <stdio.h>
#include <sys/mman.h>
#include <unistd.h>

int main (){

    int dev_mem_fd = -1;
    void * vaddr_base;

    //Mapping user-space I/O
    dev_mem_fd = open("/dev/uio0", O_RDWR|O_SYNC);
    if (dev_mem_fd < 0) { perror("open() /dev/uio0"); return 1; }

    // Map 1KB of physical memory starting at uio 0x0 (real 0x40000000)
    // to 1KB of virtual memory starting at vaddr_base
    vaddr_base = mmap(0, 1024, PROT_READ|PROT_WRITE,
                     MAP_SHARED, dev_mem_fd, 0x0);

    if (vaddr_base == MAP_FAILED) { perror("mmap()"); return 1; }
```

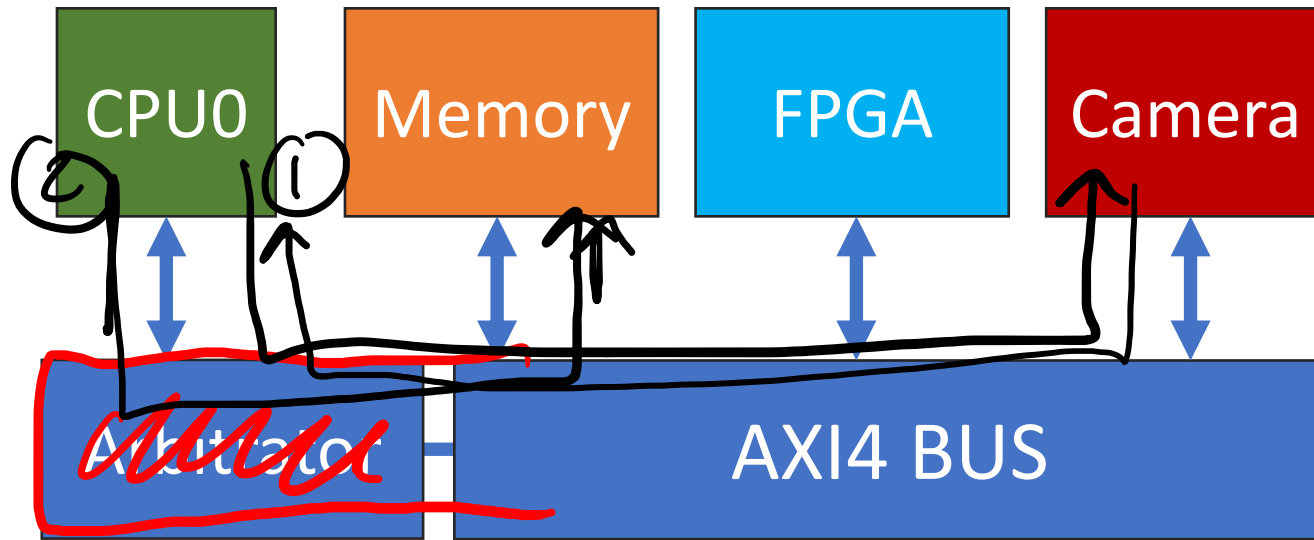
```
volatile uint32_t * ema_reg = (uint32_t*) vaddr_base;
uint32_t tmp;

for (int i = 1000; i < 6000; i +=1000){
    printf ("Sending in: %d\n", i);
    *ema_reg = i; //mmio store
    tmp = *ema_reg; //mmio load
    printf("Receiving: %d\n", tmp);
}

if (munmap(vaddr_base, 1024) != 0) { perror("munmap()"); }
if (close(dev_mem_fd) != 0) { perror("close()"); }
dev_mem_fd = -1;

return 0;
}
```

Q: How do I move data between the Camera and Memory?



① Load from Camera to CPU

② Store to memory

A: The CPU copies data from Camera to Memory

include "fancy-face-detect.h"

```
#define CAMERA_MMIO_ADDR 0x40000004
volatile uint32_t * camera =
    (uint32_t *) (CAMERA_MMIO_ADDR);
```

```
#define BUF_SIZE 1024;
uint32_t buf[BUF_SIZE];
```

```
int main () {
```

```
    for (j=0; j<BUF_SIZE; j++) {
        copy_image(camera, buf, BUF_SIZE);
        fancy-face-detect(buf);
    }
```

```
}
```

$i=2$
 $x = i++ \rightarrow x=2, i=3$
 $x = ++i \rightarrow x=3, i=3$

```
void copy_image (uint32_t * from,
                 uint32_t * to,
                 uint32_t size)
```

```
{
    register uint32_t tmp;
    for (int i = 0; i < size; i++) {
        tmp = *from;
        to[i] = tmp;
    }
}
```

A: The CPU copies data from Camera to Memory

#include "magic-face-detect.h"

```
#define CAMERA_MMIO_ADDR 0x40000004
volatile uint32_t * camera =
    (uint32_t *) (CAMERA_MMIO_ADDR);
#define BUF_SIZE 1024;
uint32_t buf[BUF_SIZE];

int main () {
    //...
    while (true){
        copy_image(camera, buf, BUF_SIZE);
        detect_face(buf);
    }
}
```

```
void copy_image (uint32_t * from,
                uint32_t * to,
                uint32_t size)
{
    uint32_t register uint32_t reg;

    for (int i = 0; i < size; ++i){
        reg = *from;

        to[i] = reg;
    }
}
```


What else can the CPU do while copying data?

CPU \rightarrow runs @ 1 GHz = f

$$P = \frac{1}{1 \text{ GHz}} = \frac{1}{10^9 \text{ Hz}} = 10^{-9} \text{ s}$$

$$= 10^9 \text{ operations / second} \quad \text{:- 1 ns/op}$$

What else can the CPU do while copying data?

- CPU can do 1B instructions/second. (1GHz)
- 4 Instructions per copy loop
 - 1 load, 1 store, 1 increment, 1 branch
- 250M copies/second

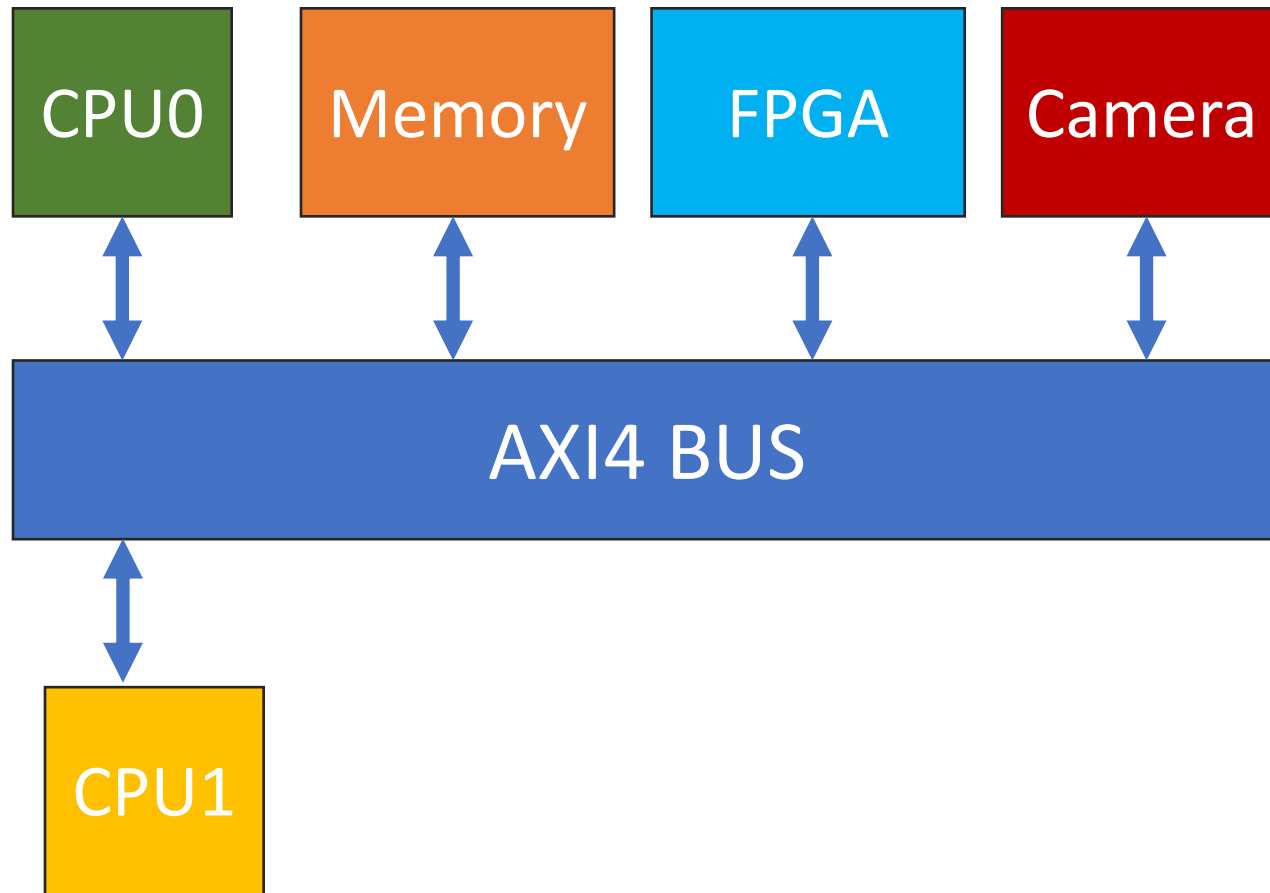
4K Video: 1697 Mbps* = 212 MB / second

~85% CPU utilization for Copy!

What about Ethernet?

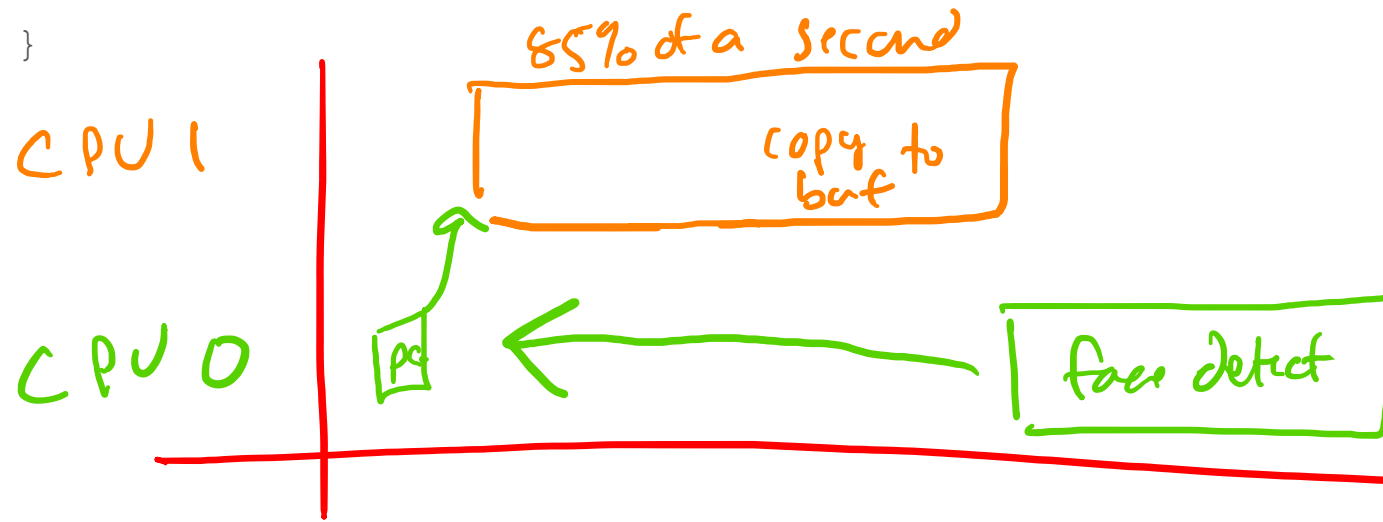
- CPU can do 1B instructions/second. (1GHz)
- 4 Instructions per copy loop
 - 1 load, 1 store, 1 increment, 1 branch
- 250M copies/second
- 1Gbps Ethernet:
- 1 Gbps Receive + 1Gbps Transmit = 2 Gbps
- 2Gbps = 250MB/second
- **Nothing. ~100% of CPU required?**

What if we do the copy on a new CPU, **CPU1**?



What if we do the copy on a new CPU, CPU1?

```
int main () {  
  
    while (true){  
  
        ask_cpu1_to_copy_image(camera, buf, BUF_SIZE);  
  
        CPU0 detect_face(buf);  
    }  
}
```



time

What if we do the copy on CPU1?

```
int [buf] buf0;  
int [buf] buf1;
```

```
int main () {
```

```
    while (true) {
```

```
        ask_cpu1_to_copy_image(camera, bufbuf0, BUF_SIZE);
```

```
        detect_face(buf);
```

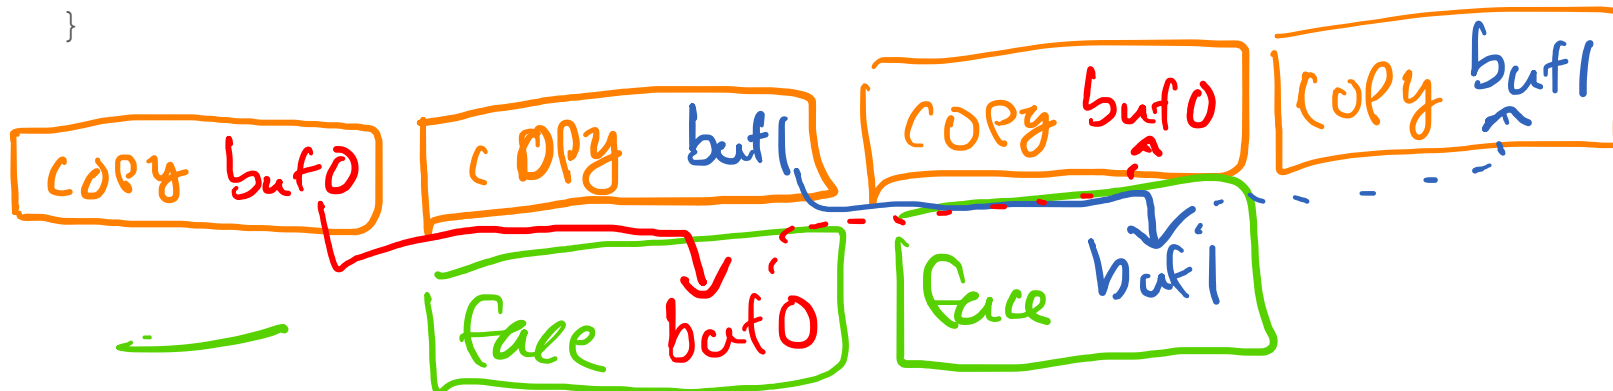
```
    }
```

```
}
```

buf0: Red
buf1: Blue
CPU0: Green
CPU1: Orange

CPU1:

CPU0:

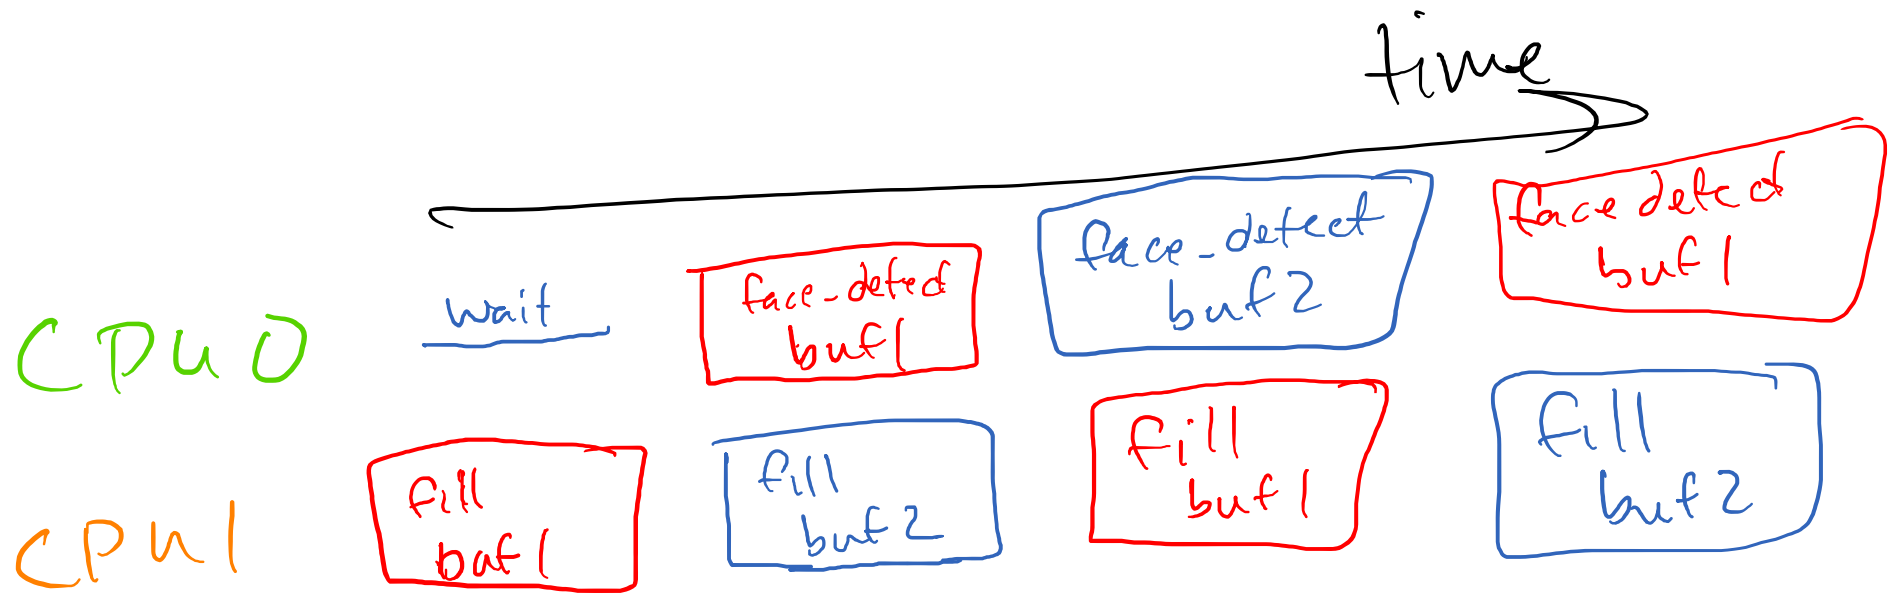


double buffering

Double Buffering explained.

Copy on CPU1, Version 2.

```
int main () {  
  
    ask_cpu1_to_copy_image(camera, buf1, BUF_SIZE);  
    wait_for_cpu1_done();  
  
    while (true){  
        ask_cpu1_to_copy_image(camera, buf2, BUF_SIZE);  
        detect_face(buf1);  
        wait_for_cpu1_done();  
  
        ask_cpu1_to_copy_image(camera, buf1, BUF_SIZE);  
        detect_face(buf2);  
        wait_for_cpu1_done();  
    }  
}
```

Why are we wasting an entire CPU for this?

```
void copy_image (uint32_t * from,
                 uint32_t * to,
                 uint32_t size)
{
    register uint32_t reg;

    for (int i = 0; i < size; ++i){

        reg = *from;

        to[i] = reg;

    }
}
```

DMA: Direct Memory Access

- **A mini-CPU that (only) does copy for you:**

```
void copy (uint32_t * from,
           uint32_t * to,
           uint32_t size)
{
    register uint32_t reg;

    for (int i = 0; i < size; ++i){

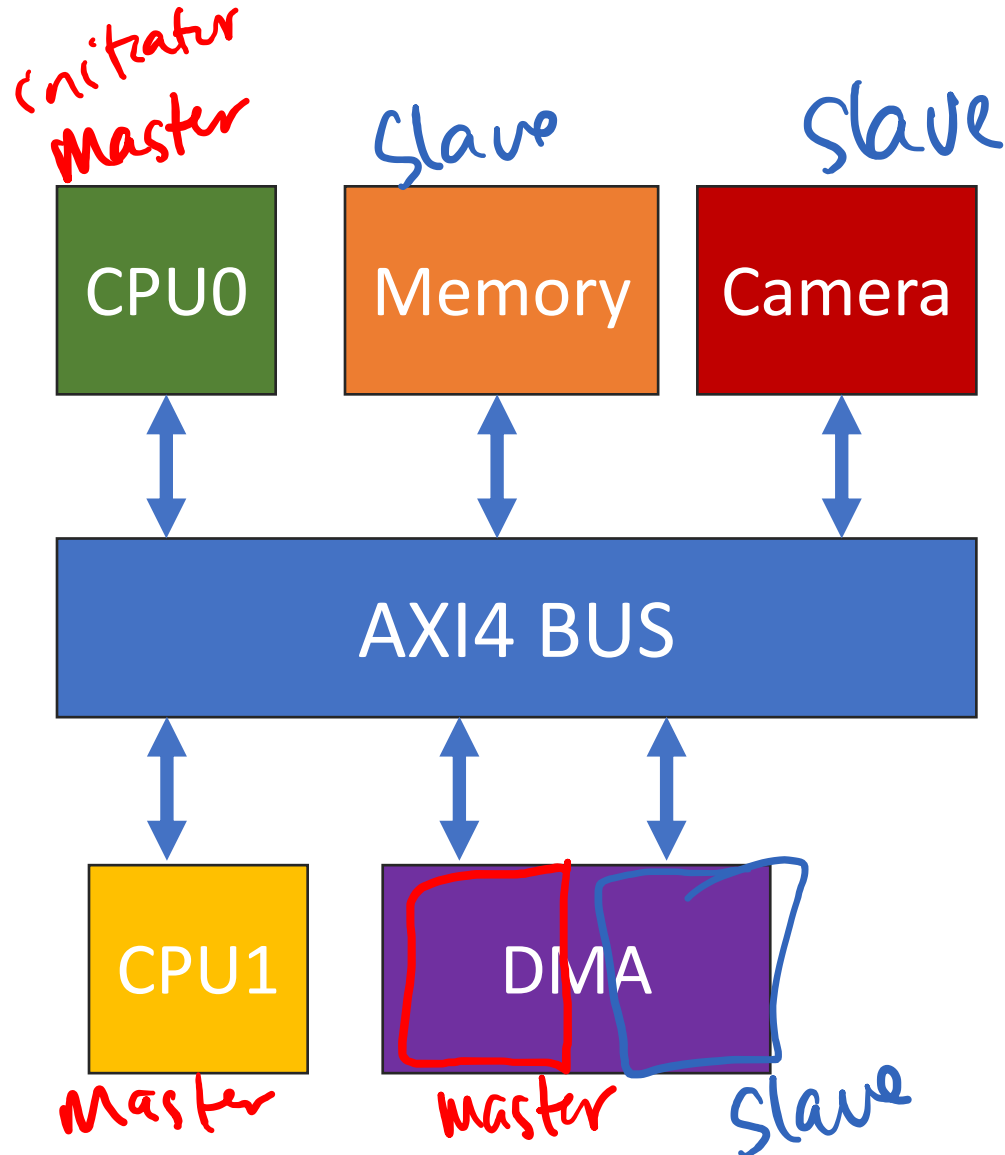
        reg = *from;

        to[i] = reg;
    }
}
```

Copy with DMA

```
int main () {  
  
    dma_start_copy (camera, buf1, BUF_SIZE);  
    dma_wait_for_done();  
  
    while (true){  
        dma_start_copy (camera, buf2, BUF_SIZE);  
        detect_face(buf1);  
        dma_wait_for_done();  
  
        dma_start_copy (camera, buf1, BUF_SIZE);  
        detect_face(buf2);  
        dma_wait_for_done();  
    }  
}
```

DMA has 2 interfaces



- Interface 1: Copy Memory
 - Data-Intensive Interface
 - AXI4 Master
 - Initiates Loads / Stores
- Interface 2: Tell DMA what to copy
 - Control Interface
 - AXI4 Slave
 - Responds to Loads/Stores


What's needed to do this in Hardware?

```
void dma_copy (uint32_t * from,  
               uint32_t * to,  
               uint32_t size)  
{  
    register uint32_t reg;  
  
    for (int i = 0; i < size; ++i){  
  
        reg = *from;  
  
        to[i] = reg;  
    }  
}
```

start
→

done
←


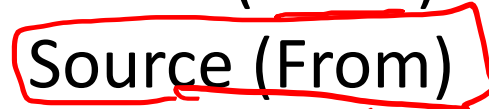
Hardware Needs:



```
void dma_copy (uint32_t * from,
               uint32_t * to,
               uint32_t size)
{
    register uint32_t reg;

    for (int i = 0; i < size; ++i) {
        reg = *from; //load

        to[i] = reg; //store
    }
}
```

- AXI4 Master Interface
 - Actual Loads + Stores
- AXI4 Slave Interface
- 5 MMIO registers
 - Control (Start) 
 - Status (Done)
 - Source (From) 
 - Destination (To)
 - Size (in Bytes)

MyDMA MMIO Interface

- 0x0400: Control Register
- 0x0404: Status Register
- 0x0408: Source Address
- 0x040C: Destination Address
- 0x0410: Transfer Size in Bytes

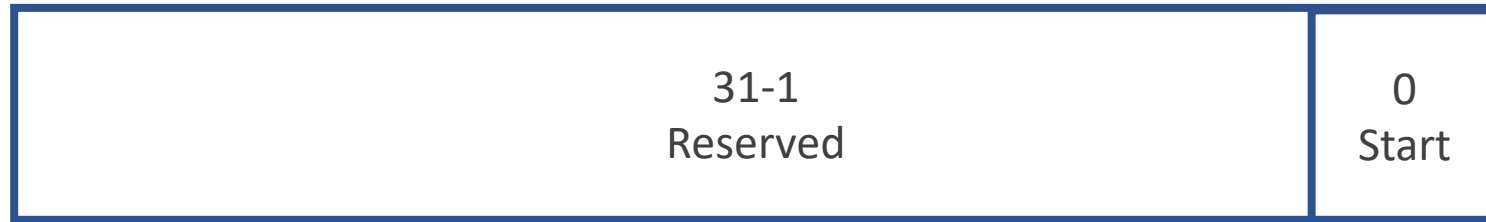
MMIO Control Register

0x4000



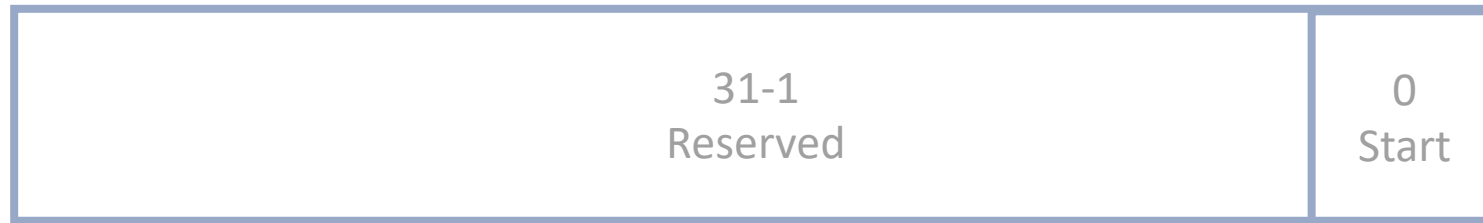
MMIO Control Register

Control - 0x0400

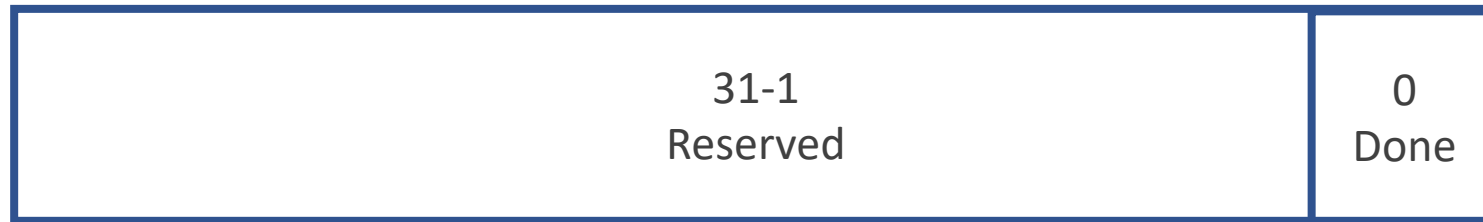


MMIO Status Register

Control - 0x0400
Read/Writer



Status - 0x0404
Read - Only

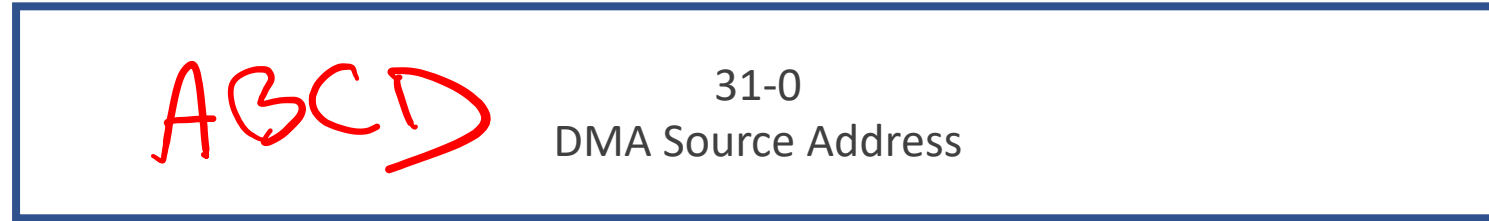


= 1 DMA Done
= 0 DMA not Done

MMIO Data Registers

1st

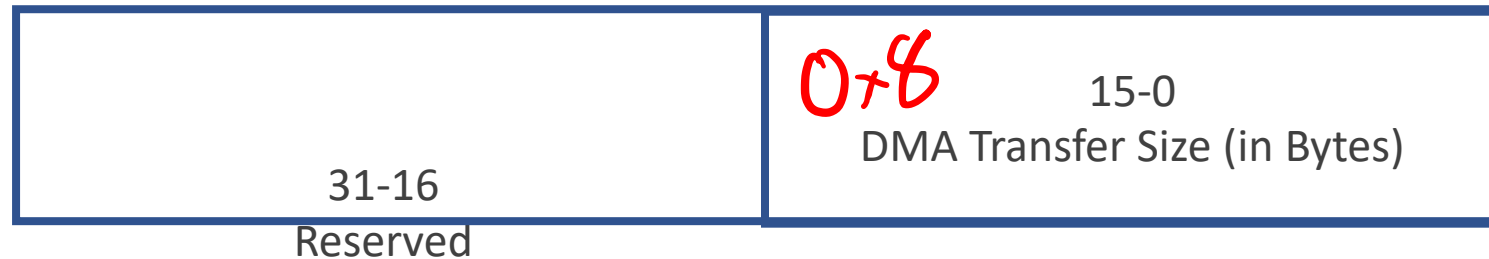
Store ABCD
to 0x0408
Source - 0x0408



Store 1234
to 0x040C
Destination - 0x040C



Store 0x8
Size - 0x0410

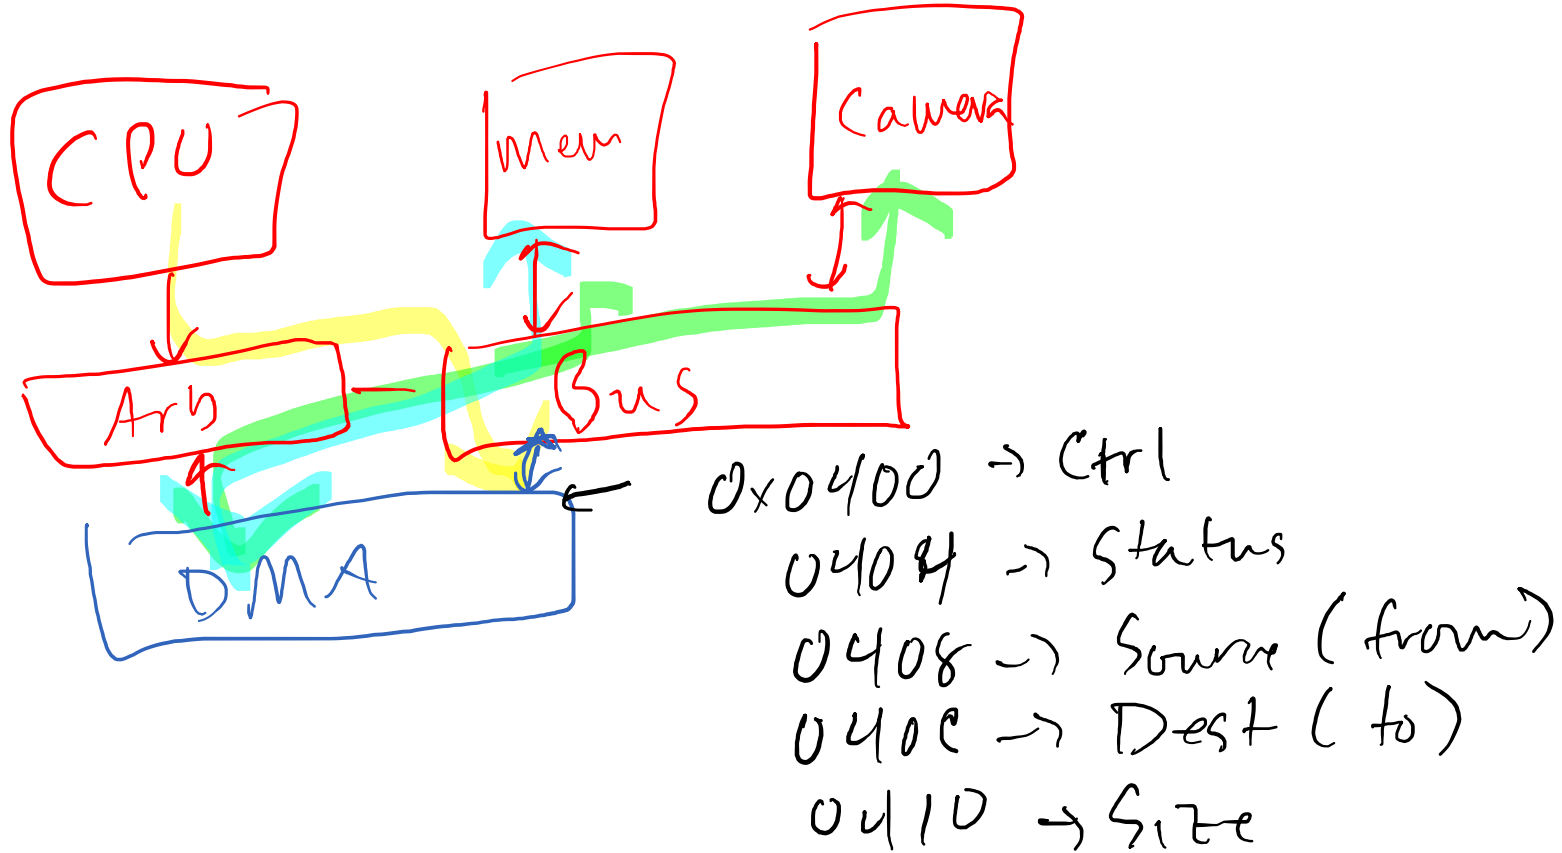


Stop

All MMIO Registers

Control - 0x0400	31-1 Reserved	0 Start
Status - 0x0404	31-1 Reserved	0 Done
Source - 0x0408	31-0 DMA Source Address	
Destination - 0x040C	31-1 DMA Destination Address	
Size - 0x0410	31-16 Reserved	15-0 DMA Transfer Size (in Bytes)

MyDMA Interface



MyDMA Internals

MyDMA Internals

- IDLE: Status[Done]=1, wait for Control[Start]
- START: Status[Done] = 0, i = 0;
- LOAD: tmp = [Source+i]
- STORE: Dest+i = tmp

Does the AXI4 Full Interface have an address?

Does the AXI4 Full Interface have an MMIO Address?

- Is pretending to be memory, or a CPU?
- Does a CPU have a memory address?
- **No.**
- MMIO is for SLAVE interfaces.

Using DMA from the CPU:

0x0400: Control Register
0x0404: Status Register
0x0408: Source Address
0x040C: Destination Address
0x0410: Transfer Size in Bytes

```
void dma_copy ( uint32_t * source,  
                uint32_t * dest,  
                uint32_t size) {  
  
    register uint32_t reg;  
  
    for (int i = 0; i < size; ++i){  
  
        reg = *source; //load  
        dest[i] = reg; //store  
    }  
  
    //code me!  
}
```

Using DMA from the CPU:

0x0400: Control Register
0x0404: Status Register
0x0408: Source Address
0x040C: Destination Address
0x0410: Transfer Size in Bytes

```
void dma_copy ( uint32_t * source,
                uint32_t * dest,
                uint32_t size) {

    *((volatile uint32_t *) (0x0408)) = source;
    *((volatile uint32_t *) (0x040C)) = dest;
    *((volatile uint32_t *) (0x0410)) = size;
    *((volatile uint32_t *) (0x0400)) = 0x1; //start

    //spin until copy done
    while( *((volatile uint32_t *) (0x0404)) != 0x1) {;}
}
```

Using DMA from the CPU:

0x0400: Control Register
0x0404: Status Register
0x0408: Source Address
0x040C: Destination Address
0x0410: Transfer Size in Bytes

```
void dma_start_copy (    uint32_t * source,
                        uint32_t * dest,
                        uint32_t size){

    *((volatile uint32_t *) (0x0408))= source;
    *((volatile uint32_t *) (0x040C))= dest;
    *((volatile uint32_t *) (0x0410))= size;
    *((volatile uint32_t *) (0x0400))= 0x1; //start
}

void dma_wait_for_done(){
    //spin until copy done?
    while( *((uint32_t) (0x0404)) != 0x1){;}
}
```

volatile keyword omitted!

Using DMA from C:

```
int main () {  
  
    dma_start_copy (camera, buf1, BUF_SIZE);  
    dma_wait_for_done();  
  
    while (true){  
        dma_start_copy (camera, buf2, BUF_SIZE);  
        detect_face(buf1);  
        dma_wait_for_done();  
  
        dma_start_copy (camera, buf1, BUF_SIZE);  
        detect_face(buf2);  
        dma_wait_for_done();  
    }  
}
```

11: Multi-Master Buses

Engr 315: Hardware / Software Codesign
Andrew Lukefahr
Indiana University

