

# 17: Pipelining II

ENGR 315: Hardware/Software CoDesign


Andrew Lukefahr

Indiana University

Some material taken from:

[https://github.com/trekhleb/homemade-machine-learning/tree/master/homemade/neural\\_network](https://github.com/trekhleb/homemade-machine-learning/tree/master/homemade/neural_network)

<http://cs231n.github.io/neural-networks-1/>




# Announcements

- P6 out
- No Class on Wednesday (Oct 26<sup>th</sup>)
- New SD Cards coming for P7...
- Exam Nov 7th

# P6: Adds DMA + AXI-Stream to Popcount

- DMA
  - Add DMA engine to move data via AXI4-Full to AXI-Stream interface
- Popcount.sv:
  - Add AXI-Stream Interface
  - Keep AXI4-Lite Interface to read result



## P7 – DMA from C

1. Start the MM2S channel running by setting the run/stop bit to 1 (MM2S\_DMACR.RS = 1). The halted bit (DMASR.Halted) should deassert indicating the MM2S channel is running.

### 2. Skip

3. Write a valid source address to the MM2S\_SA register.
4. Write the number of bytes to transfer in the MM2S\_LENGTH register.  
The MM2S\_LENGTH register must be written last.
5. Wait until MM2S\_DMASR.Idle==1 for completion


# P8+ Accelerate Machine Learning

- Goal: Accelerate reference neural network
- Harder, more open-ended projects

# Simple Neural Network

=====  
Index: 0

Image:




→ ML Classification Result: 2 ←

Real Value: 2

Correct Result: True


- Takes in image of number
- Returns integer value
- How? artificial neural network

# Python Neuron



```
class Neuron(object):  
    # ...  
    def forward(self, inputs):  
        """ assume inputs and weights are 1-D numpy arrays and bias is a number """  
        cell_body_sum = np.sum(inputs * self.weights) + self.bias  
        firing_rate = 1.0 / (1.0 + math.exp(-cell_body_sum)) # sigmoid activation function  
        return firing_rate
```

# Why Dot Product?



```
# forward-pass of a 3-layer neural network:  
f = lambda x: 1.0/(1.0 + np.exp(-x)) # activation function (use sigmoid)  
x = np.random.randn(3, 1) # random input vector of three numbers (3x1)  
h1 = f(np.dot(W1, x) + b1) # calculate first hidden layer activations (4x1)  
h2 = f(np.dot(W2, h1) + b2) # calculate second hidden layer activations (4x1)  
out = np.dot(W3, h2) + b3 # output neuron (1x1)
```



# Matrix Multiplication (Dot Product)

$$\begin{bmatrix} i_0 & i_1 \end{bmatrix} \times \begin{bmatrix} w_{00} & w_{10} & w_{20} \\ w_{01} & w_{11} & w_{21} \end{bmatrix} = \begin{bmatrix} o_0 & o_1 & o_2 \end{bmatrix}$$

$$o_0 = i_0 \cdot w_{00} + i_1 \cdot w_{01}$$

$$o_1 = i_0 \cdot w_{10} + i_1 \cdot w_{11}$$

$$o_2 = i_0 \cdot w_{20} + i_1 \cdot w_{21}$$

# Matrix Multiplication (Dot Product)

$$\begin{array}{c} \text{inputs} \\ \underline{[0.1 \quad 0.2]} \end{array} \times \begin{array}{c} \text{weights} \\ \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \end{array} =$$

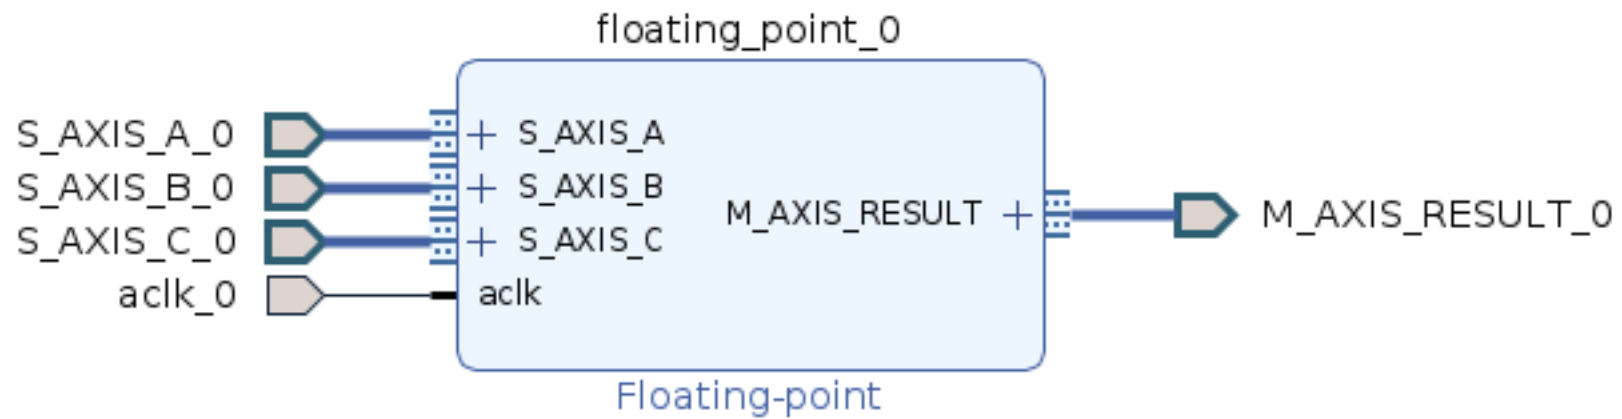
$$= \begin{bmatrix} (0.1 \times 1 + 0.2 \times 4) & (0.1 \times 2 + 0.2 \times 5) & (0.1 \times 3 + 0.2 \times 6) \end{bmatrix}$$

$$= \begin{array}{c} \text{(Answer)} \\ \underline{[0.9 \quad 1.2 \quad 1.5]} \end{array}$$

# Floating-Point Multiply-Accumulate (FMAC)

- Math:  $a * b + c$

# Floating-Point Multiply in Hardware




- $\text{result} = a * b + c$

# Floating-Point math takes 8 cycles.

- Floating-Point is complicated.
- 8 cycles of complicated.

# Demo Time



# Floating-Point math takes 8 cycles.

- Floating-Point is complicated.
- 8 cycles of complicated.
- How do we work around an 8 cycle latency?
- Pipelining!

# Floating-Point math takes 8 cycles.

- Floating-Point is complicated.
- How do we work around an 8 cycle latency?
- Pipelining!

$A \cdot B + C$   
~~~~~  
MAC



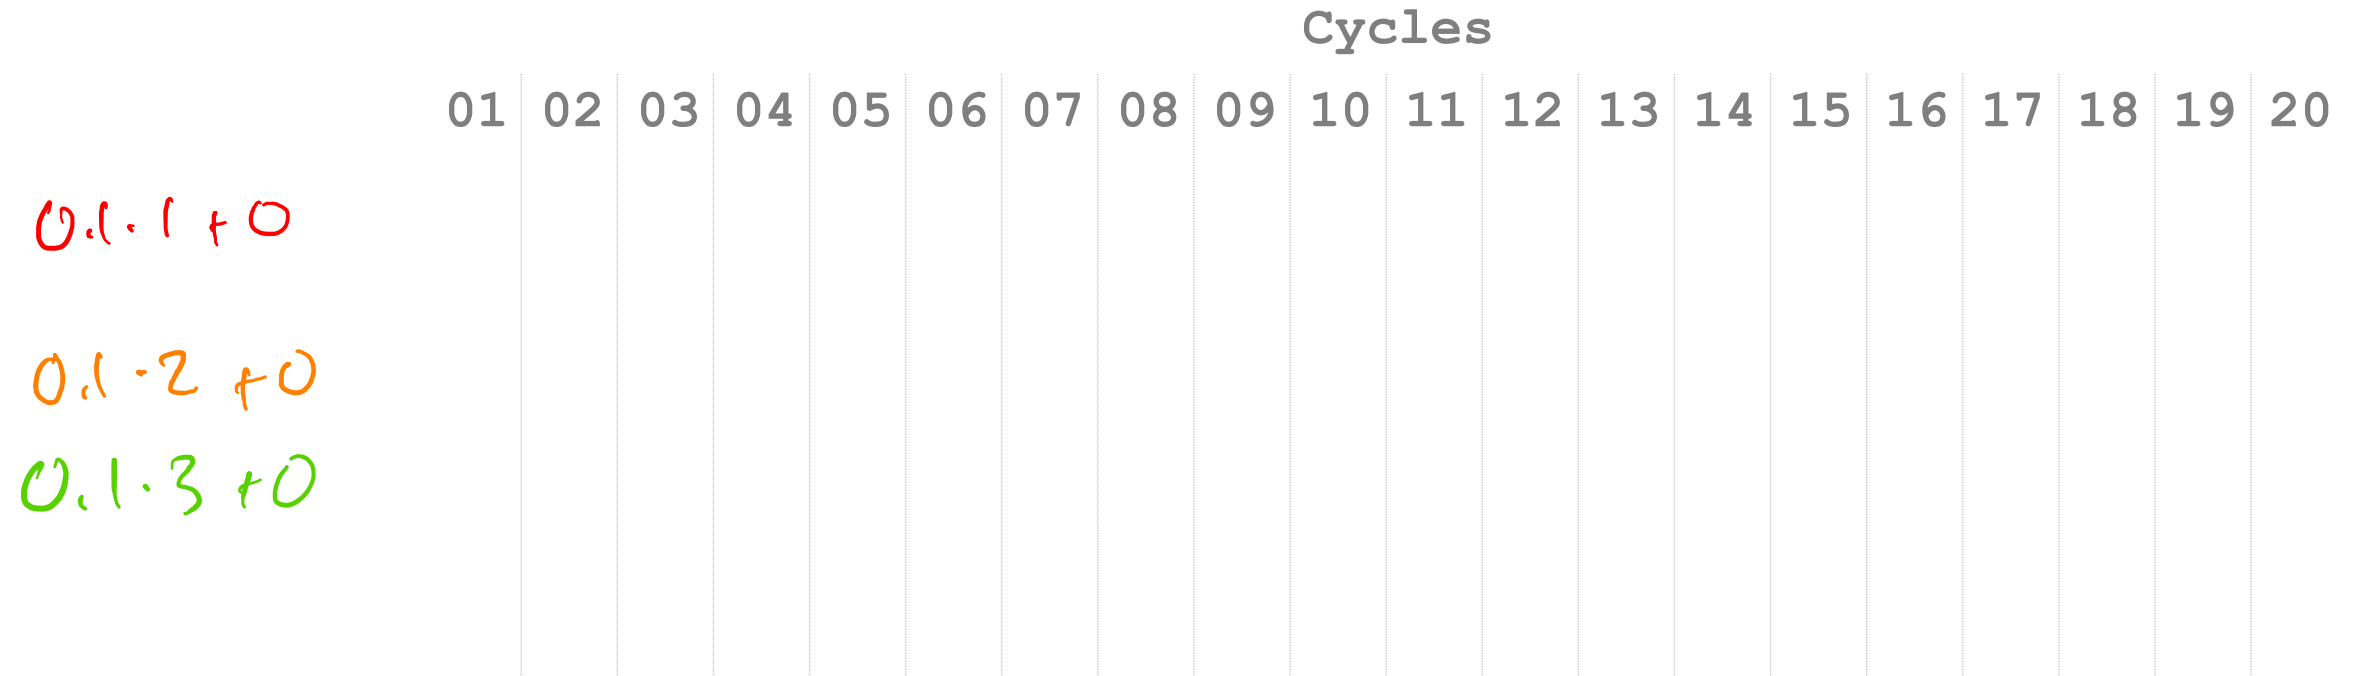
# Pipelining in hardware

$$\begin{array}{r} 2 \\ \times 3 \\ \hline 6 \end{array} \Rightarrow \begin{array}{r} 10 \\ \times 11 \\ \hline 110 \end{array} \Rightarrow \begin{array}{r} 10 \\ \times 1 \\ \hline 10 \end{array} + \begin{array}{r} 10 \\ \times 10 \\ \hline 100 \end{array} = 110$$

# Pipelining in hardware

$$\begin{array}{r} 2 \\ \times 3 \\ \hline 6 \end{array} \Rightarrow \begin{array}{r} 10 \\ \times 11 \\ \hline 110 \end{array} \Rightarrow \begin{array}{r} 10 \\ \times 1 \\ \hline 10 \end{array} + \begin{array}{r} 10 \\ \times 10 \\ \hline 100 \end{array} = 110$$

# FMAC Pipelining



# FMAC Pipelining

## Cycles

01 02 03 04 05 06 07 08 09 10 11 12 13 14 15 16 17 18 19 20

# Latency vs. Throughput

- **Latency:** How long does an individual operation take to complete?
- **Throughput:** How many operations can you complete per second (or per cycle)?

# Pipelining

- FMAC takes 8 cycles for 1 value
  - But can accept a new value every cycle.
- 
- What is Latency:
  - What is Throughput:

# Recall: Matrix Multiplication (Dot Product)

$$\begin{array}{c} \text{inputs} \\ \underline{[0.1 \quad 0.2]} \end{array} \times \begin{array}{c} \text{weights} \\ \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \end{array} =$$

$$= \begin{bmatrix} (0.1 \times 1 + 0.2 \times 4) & (0.1 \times 2 + 0.2 \times 5) & (0.1 \times 3 + 0.2 \times 6) \end{bmatrix}$$

$$= \begin{array}{c} \text{(Answer)} \\ \underline{[0.9 \quad 1.2 \quad 1.5]} \end{array}$$

# Alternative Dot Computations

$$\begin{bmatrix} 0.1 & 0.2 \end{bmatrix} \times \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} = \begin{bmatrix} 0.9 & 1.2 & 1.5 \end{bmatrix}$$

partial  
result

$$\begin{bmatrix} 0.1 \cdot 1 & 0.1 \cdot 2 & 0.1 \cdot 3 \end{bmatrix} = \begin{bmatrix} 0.1 & 0.2 & 0.3 \end{bmatrix}$$

$$\begin{bmatrix} 0.2 \cdot 4 & 0.2 \cdot 5 & 0.2 \cdot 6 \end{bmatrix} = \begin{bmatrix} 0.8 & 1.0 & 1.2 \end{bmatrix}$$

---

$$\begin{bmatrix} 0.9 & 1.2 & 1.5 \end{bmatrix}$$



# Multiply-Accumulate Dot Computations

$$\begin{bmatrix} 0.1 & 0.2 \end{bmatrix} \times \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} = \begin{bmatrix} 0.9 & 1.2 & 1.5 \end{bmatrix}$$

# Python Time

$$\begin{array}{c} \text{inputs} \\ [0.1 \quad 0.2 \quad 0.3] \end{array} \times \begin{array}{c} \text{weights} \\ \begin{bmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \end{bmatrix} \end{array} =$$

```
weights = np.array( [[1,2,3,4],[5,6,7,8],[9,10,11,12]], dtype=np.float32)
inputs = np.array([[0.1,0.2,0.3]], dtype=np.float32)
outputs = np.dot(inputs, weights)
```

| Input         |   | Weights           |   | Output         |
|---------------|---|-------------------|---|----------------|
| [0.1 0.2 0.3] | . | [1. 2. 3. 4.]     | = | [3.8000002 4.4 |
|               |   | [5. 6. 7. 8.]     |   | 5. 5.6000004]  |
|               |   | [ 9. 10. 11. 12.] |   |                |

| Input           |   | Weights           |   | Output         |            |
|-----------------|---|-------------------|---|----------------|------------|
| [[0.1 0.2 0.3]] | . | [1. 2. 3. 4.]     | = | [3.8000002 4.4 | 5.         |
|                 |   | [5. 6. 7. 8.]     |   |                | 5.6000004] |
|                 |   | [ 9. 10. 11. 12.] |   |                |            |

# Mult-Accum Dot

*# how its done in dot.sv*

```
def pydot(inputs, weights):
    inputs = inputs[0] # remove outer nesting
    outs = np.zeros(weights.shape[1], dtype=np.float32)
    for i in range(weights.shape[0]): # input length
        for j in range(weights.shape[1]): # output length
            outs[j] = outs[j] + weights[i][j] * inputs[i]
    return outs
```

Inputs (Shape):  
(1, 3)  
Output (Shape):  
(1, 4)  
Weights (Shape):  
(3, 4)

# Dependencies

| Input           |   | Weights           |   | Output                       |
|-----------------|---|-------------------|---|------------------------------|
| [[0.1 0.2 0.3]] | . | [1. 2. 3. 4.]     | = | [3.8000002 4.4 5. 5.6000004] |
|                 |   | [5. 6. 7. 8.]     |   |                              |
|                 |   | [ 9. 10. 11. 12.] |   |                              |

## Cycles

| 01 | 02 | 03 | 04 | 05 | 06 | 07 | 08 | 09 | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 | 18 | 19 | 20 |
|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|
|    |    |    |    |    |    |    |    |    |    |    |    |    |    |    |    |    |    |    |    |

Next Time: More on  
Dependencies

# Latency on Pipelined FMAC

- Solution: Stall at the end of a row.
- Drain the pipeline.

# Hardware Parallelism

- CPU: 1 Floating-Point Unit
- FPGA? 10 Floating-Point Units?  
20 ?  
100 ?

# Finding Parallelism

- Some some computation that doesn't depend on other computation's results
- Shared Inputs are OK.



Next Time: Can we use 2+ FMACs?

$$\begin{bmatrix} 0.1 & 0.2 & 0.3 \end{bmatrix} \begin{bmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \end{bmatrix} = \begin{bmatrix} 3.8 & 4.4 & 5 & 5.6 \end{bmatrix}$$

Option 1

Stopped here!

| <u>Cycle</u> | fmac comp         |
|--------------|-------------------|
| 0            | $0.1 \cdot 1 + 0$ |
| 1            | $0.1 \cdot 2 + 0$ |
| 2            | $0.1 \cdot 3 + 0$ |
| 3            | $0.1 \cdot 4 + 0$ |

| <u>Cycle</u> | fmac 1            |
|--------------|-------------------|
| 0            | $0.1 \cdot 1 + 0$ |
| 1            | $0.1 \cdot 3 + 0$ |

| <u>fmac 2</u>     |
|-------------------|
| $0.1 \cdot 2 + 0$ |
| $0.1 \cdot 3 + 0$ |

Option 2

→  $0.1 * 1 + 0$   
 $0.1 * 2 + 0$   
 $0.1 * 3 + 0$   
 $0.1 * 4 + 0$

$0.2 * 5 + 0$   
 $0.2 * 6 + 0$   
 $0.2 * 7 + 0$   
 $0.2 * 8 + 0$

$$\begin{bmatrix} 0.1 & 0.2 & 0.3 \end{bmatrix} \begin{bmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \end{bmatrix} = \begin{bmatrix} 3.8 & 4.4 & 5 & 5.6 \end{bmatrix}$$

$$\begin{bmatrix} 0 & 0 & 0 & 0 \end{bmatrix}$$

$$0.1 \cdot \begin{bmatrix} 1 & 2 & 3 & 4 \end{bmatrix} + \begin{bmatrix} 0 & 0 & 0 & 0 \end{bmatrix} = \begin{bmatrix} 0.1 & 0.2 & 0.3 & 0.4 \end{bmatrix}$$

$$0.2 \cdot \begin{bmatrix} 5 & 6 & 7 & 8 \end{bmatrix} + \begin{bmatrix} 0.1 & 0.2 & 0.3 & 0.4 \end{bmatrix} = \begin{bmatrix} 1.1 & 1.4 & 1.7 & 2.0 \end{bmatrix}$$

$$0.3 \cdot \begin{bmatrix} 9 & 10 & 11 & 12 \end{bmatrix} + \begin{bmatrix} 1.1 & 1.4 & 1.7 & 2.0 \end{bmatrix} = \begin{bmatrix} 3.8 & 4.4 & 5 & 5.6 \end{bmatrix}$$

# Parallize Alternative Dot Computations?

$$\begin{bmatrix} 0.1 & 0.2 \end{bmatrix} \times \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} = \begin{bmatrix} 0.9 & 1.2 & 1.5 \end{bmatrix}$$

$$\begin{bmatrix} 0 & 0 & 0 \end{bmatrix} \leftarrow \text{result}$$

$$0.1 \cdot \begin{bmatrix} 1 & 2 & 3 \end{bmatrix} + \begin{bmatrix} 0 & 0 & 0 \end{bmatrix} =$$

$$\begin{bmatrix} 0.1 & 0.2 & 0.3 \end{bmatrix} + \begin{bmatrix} 0 & 0 & 0 \end{bmatrix} = \begin{bmatrix} 0.1 & 0.2 & 0.3 \end{bmatrix} \leftarrow \text{temp result}$$

$$0.2 \cdot \begin{bmatrix} 4 & 5 & 6 \end{bmatrix} + \begin{bmatrix} 0.1 & 0.2 & 0.3 \end{bmatrix} =$$

$$\begin{bmatrix} 0.8 & 1.0 & 1.2 \end{bmatrix} + \begin{bmatrix} 0.1 & 0.2 & 0.3 \end{bmatrix} = \begin{bmatrix} 0.9 & 1.2 & 1.5 \end{bmatrix}$$

# Can we parallelize Dot?

```
# how its done in dot.sv
def pydot(inputs,weights):
    inputs = inputs[0] # remove outer nesting
    outs = np.zeros(weights.shape[1], dtype=np.float32)
    for i in range(weights.shape[0]): # input length
        for j in range(weights.shape[1]): # output length
            outs[j] = outs[j] + weights[i][j] * inputs[i]
    return outs
```

# Can we parallelize Dot?

```
# how its done in dot.sv
def pydot(inputs, weights):
    inputs = inputs[0] # remove outer nesting
    outs = np.zeros(weights.shape[1], dtype=np.float32)
    for i in range(weights.shape[0]): # input length
        for j in range(weights.shape[1]): # output length
            outs[j] = outs[j] + weights[i][j] * inputs[i]
    return outs
```

```
def par_pydot(inputs, weights):
    par_inputs = [inputs[:, ::2], inputs[:, 1::2]]
    par_weights = [weights[:, ::2, :], weights[:, 1::2, :]]

    par_outputs = [pydot(par_inputs[0], par_weights[0]),
                   pydot(par_inputs[1], par_weights[1])]

    outputs = par_outputs[0] + par_outputs[1]
    return outputs
```

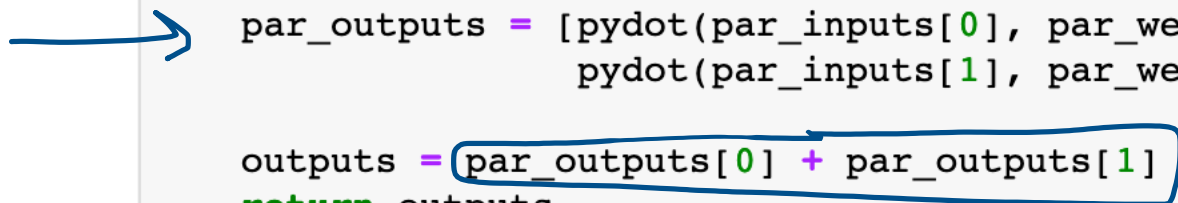
# Can we parallelize Dot?

```
# how its done in dot.sv
def pydot(inputs, weights):
    inputs = inputs[0] # remove outer nesting
    outs = np.zeros(weights.shape[1], dtype=np.float32)
    for i in range(weights.shape[0]): # input length
        for j in range(weights.shape[1]): # output length
            outs[j] = outs[j] + weights[i][j] * inputs[i]
    return outs
```

```
def par_pydot(inputs, weights):
    par_inputs = [inputs[:, ::2], inputs[:, 1::2]]
    par_weights = [weights[:, ::2, :], weights[1::2, :, :]]

    par_outputs = [pydot(par_inputs[0], par_weights[0]),
                   pydot(par_inputs[1], par_weights[1])]

    outputs = par_outputs[0] + par_outputs[1]
    return outputs
```



# 19: Hardware Acceleration III

Engr 315: Hardware / Software Codesign  
Andrew Lukefahr  
*Indiana University*

